

# Fundamentals of CUDA 1

Prof. Seokin Hong

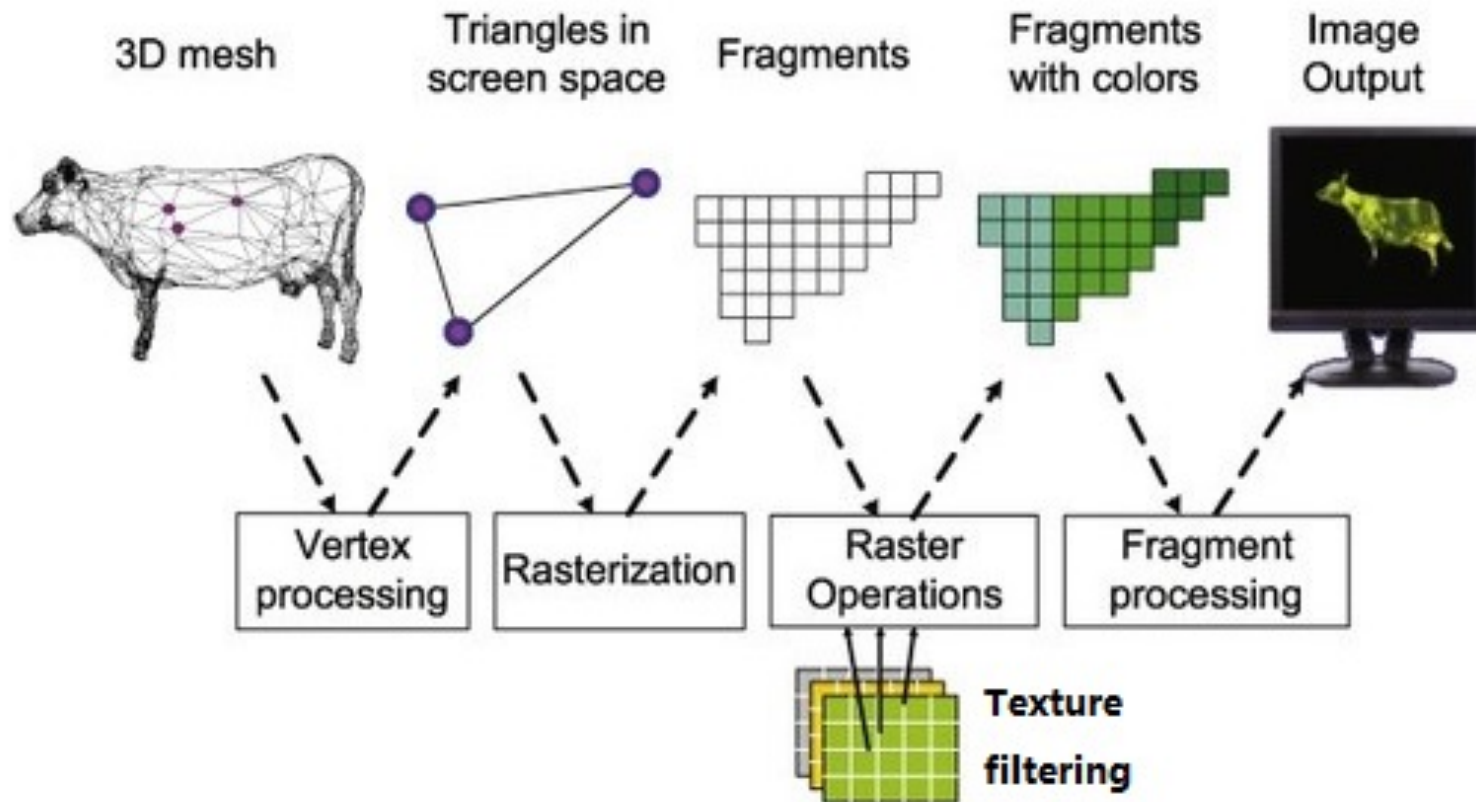
# Agenda

---

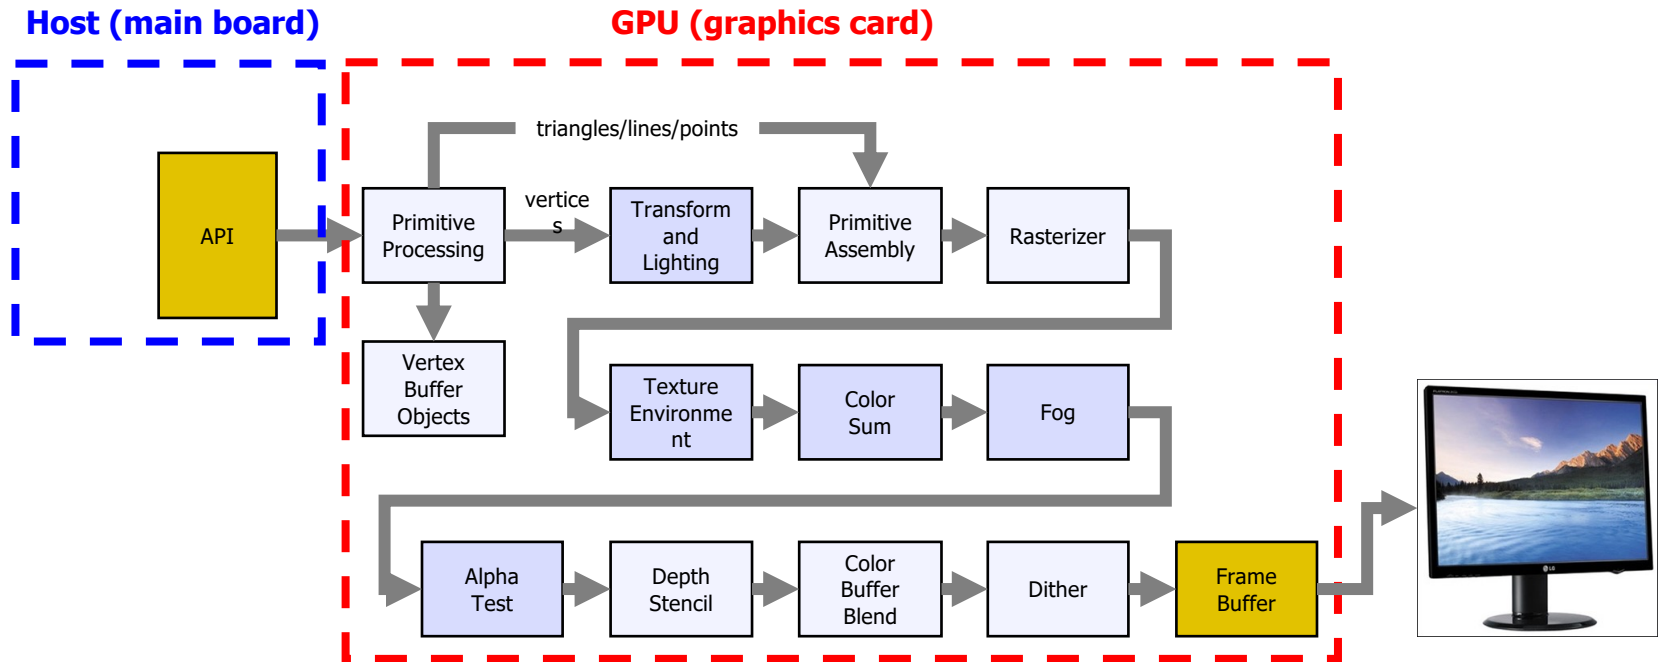
- **History**
- **What is CUDA?**
- **Device Global Memory and Data Transfer**
- **Error Checking**
- A Vector Addition Kernel
- Kernel Functions and Threading
- Kernel Launch

# History

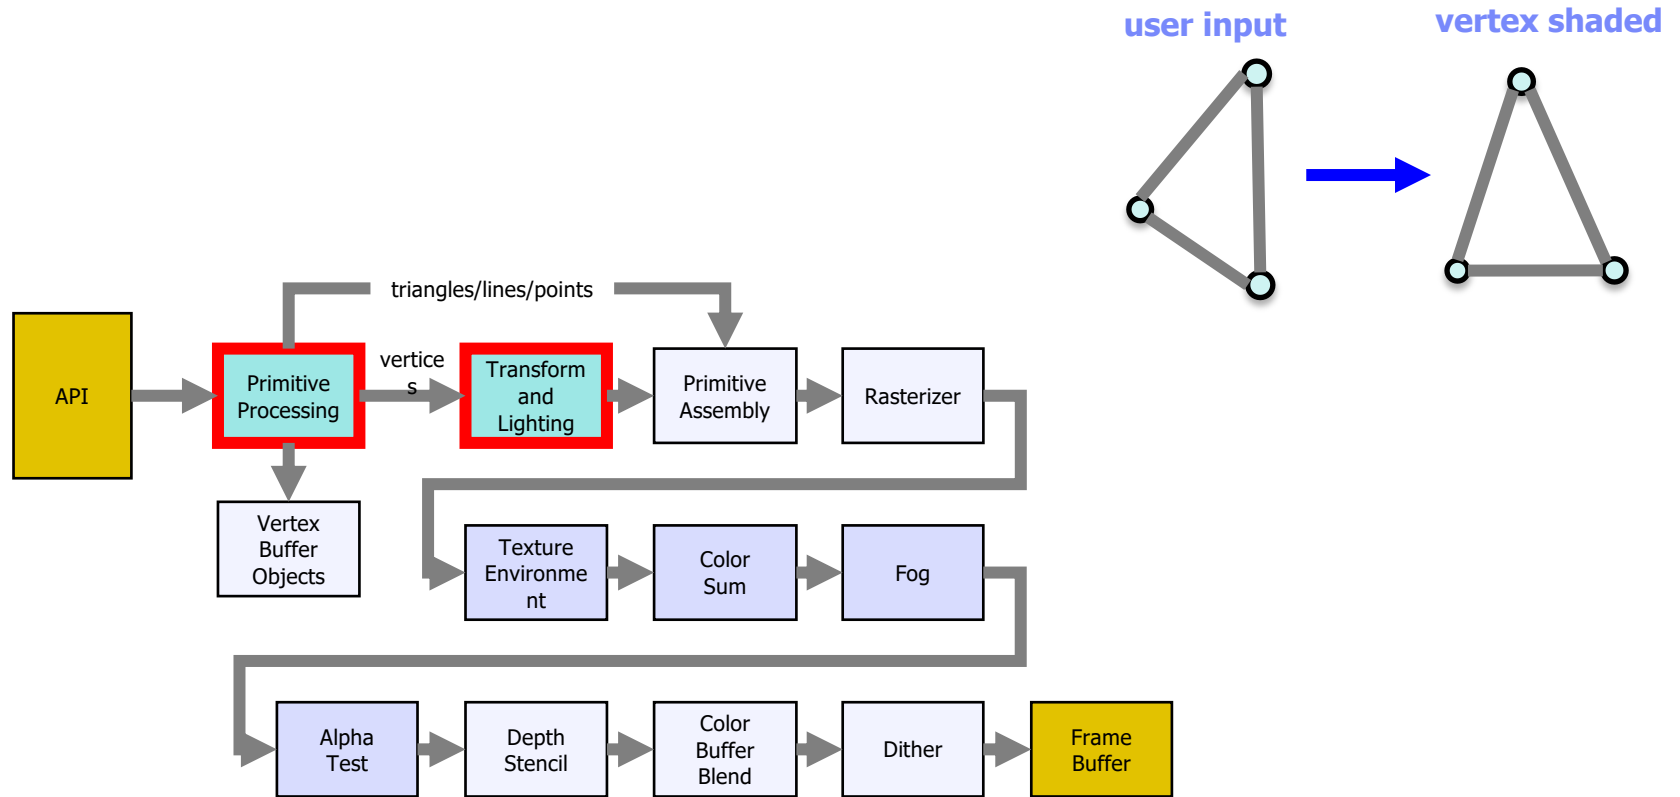
# 3D Graphics Pipeline



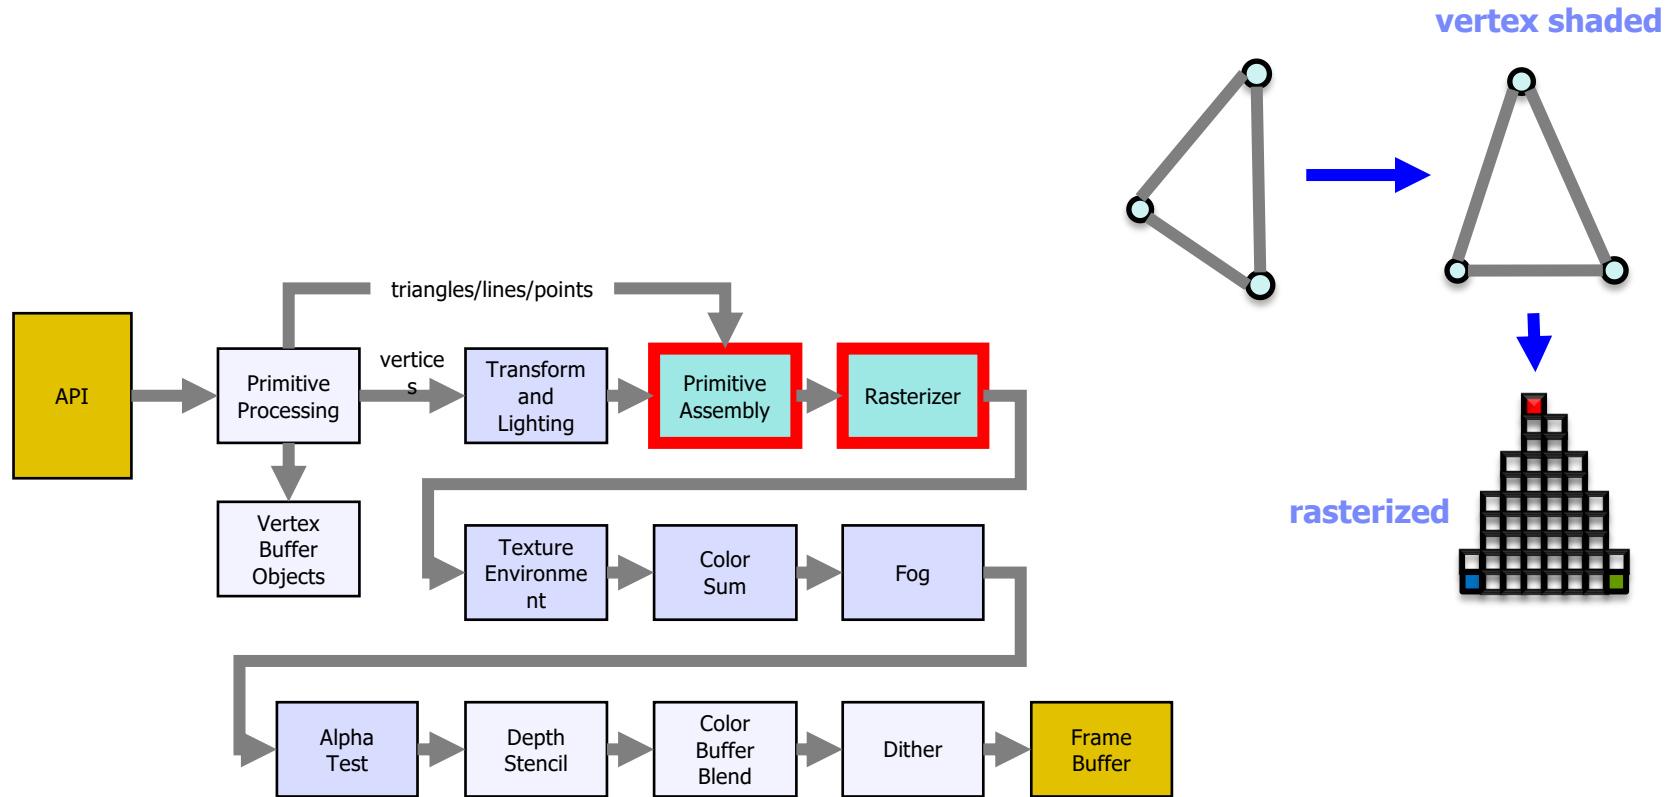
# 3D Graphics Pipeline



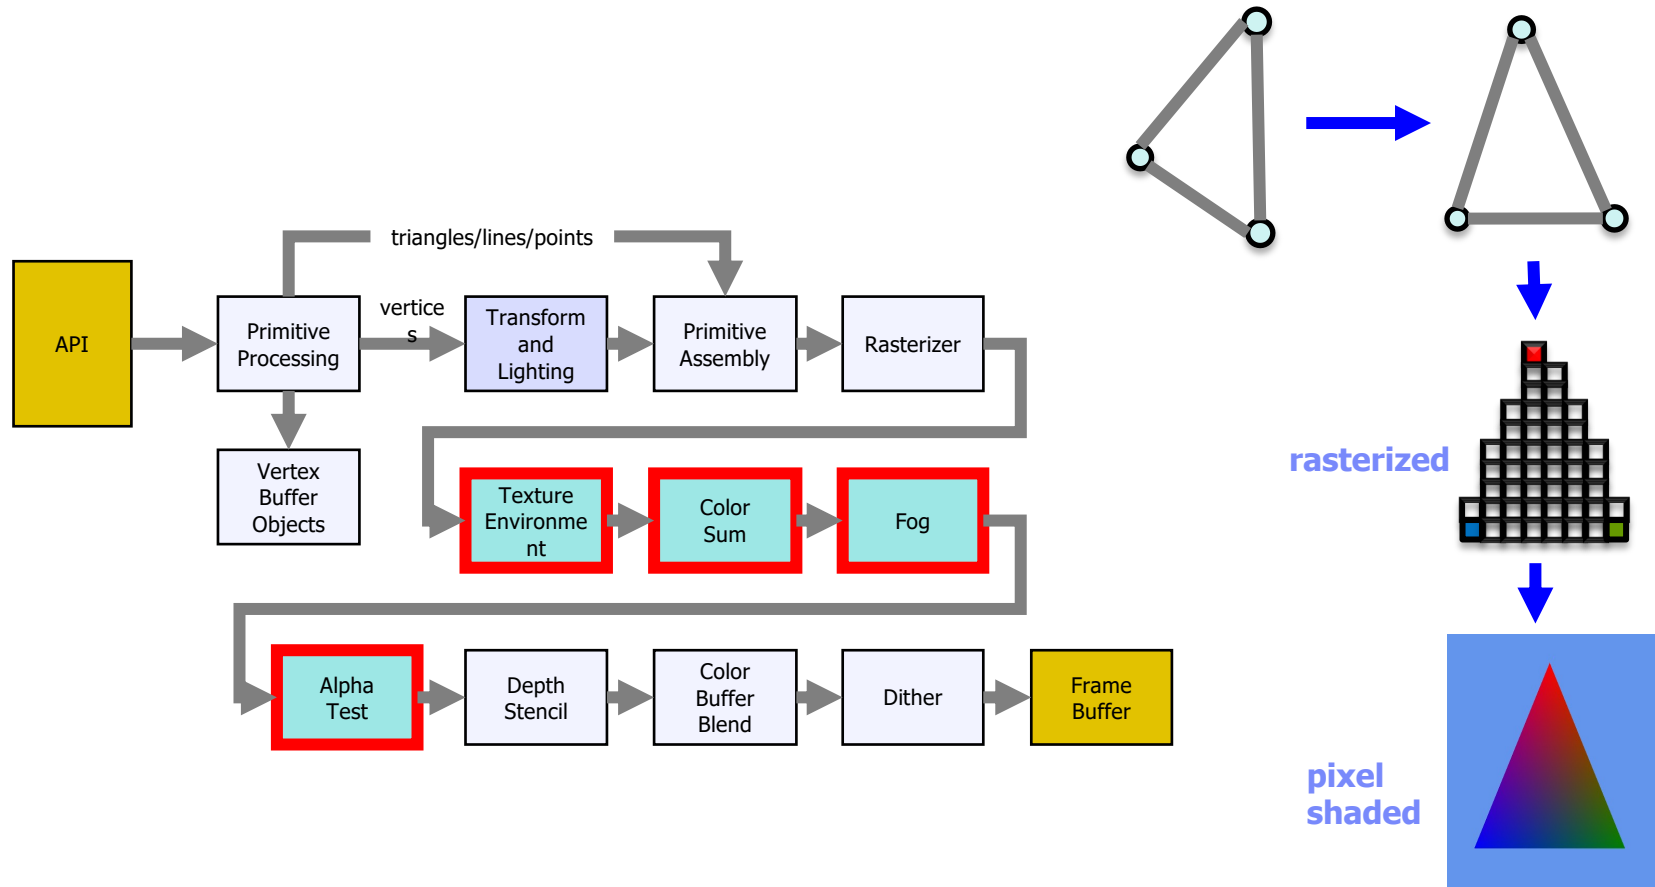
# 3D Graphics Pipeline



# 3D Graphics Pipeline



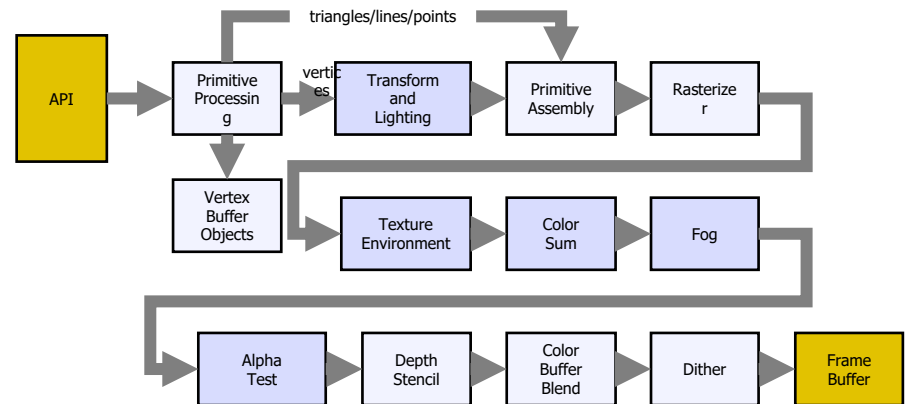

# 3D Graphics Pipeline





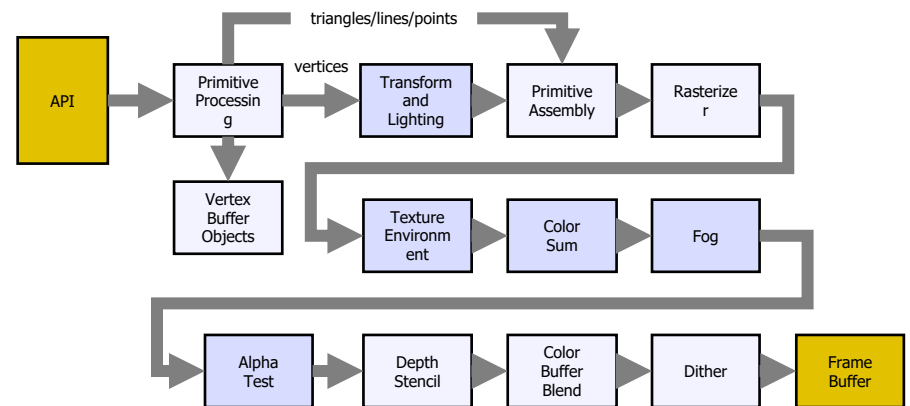

# The Graphics Pipeline – 1st Gen.

- One chip/board per stage
- Fixed data flow through pipeline



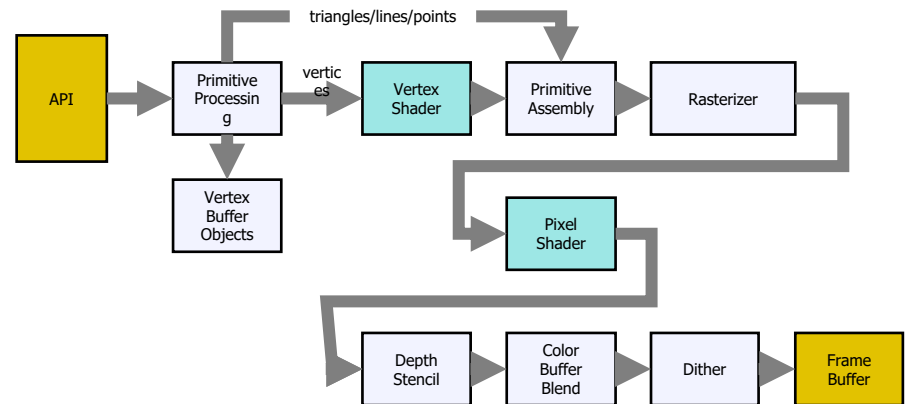
# The Graphics Pipeline – 2nd Gen.

- Everything fixed function,  
with a certain number of modes
- Number of modes for each stage grew over time
- Hard to optimize HW
- Developers always wanted more flexibility



# The Graphics Pipeline – 3rd Gen.

- Vertex & pixel processing became programmable
- GPU architecture increasingly centers around 'shader' execution



# Before CUDA

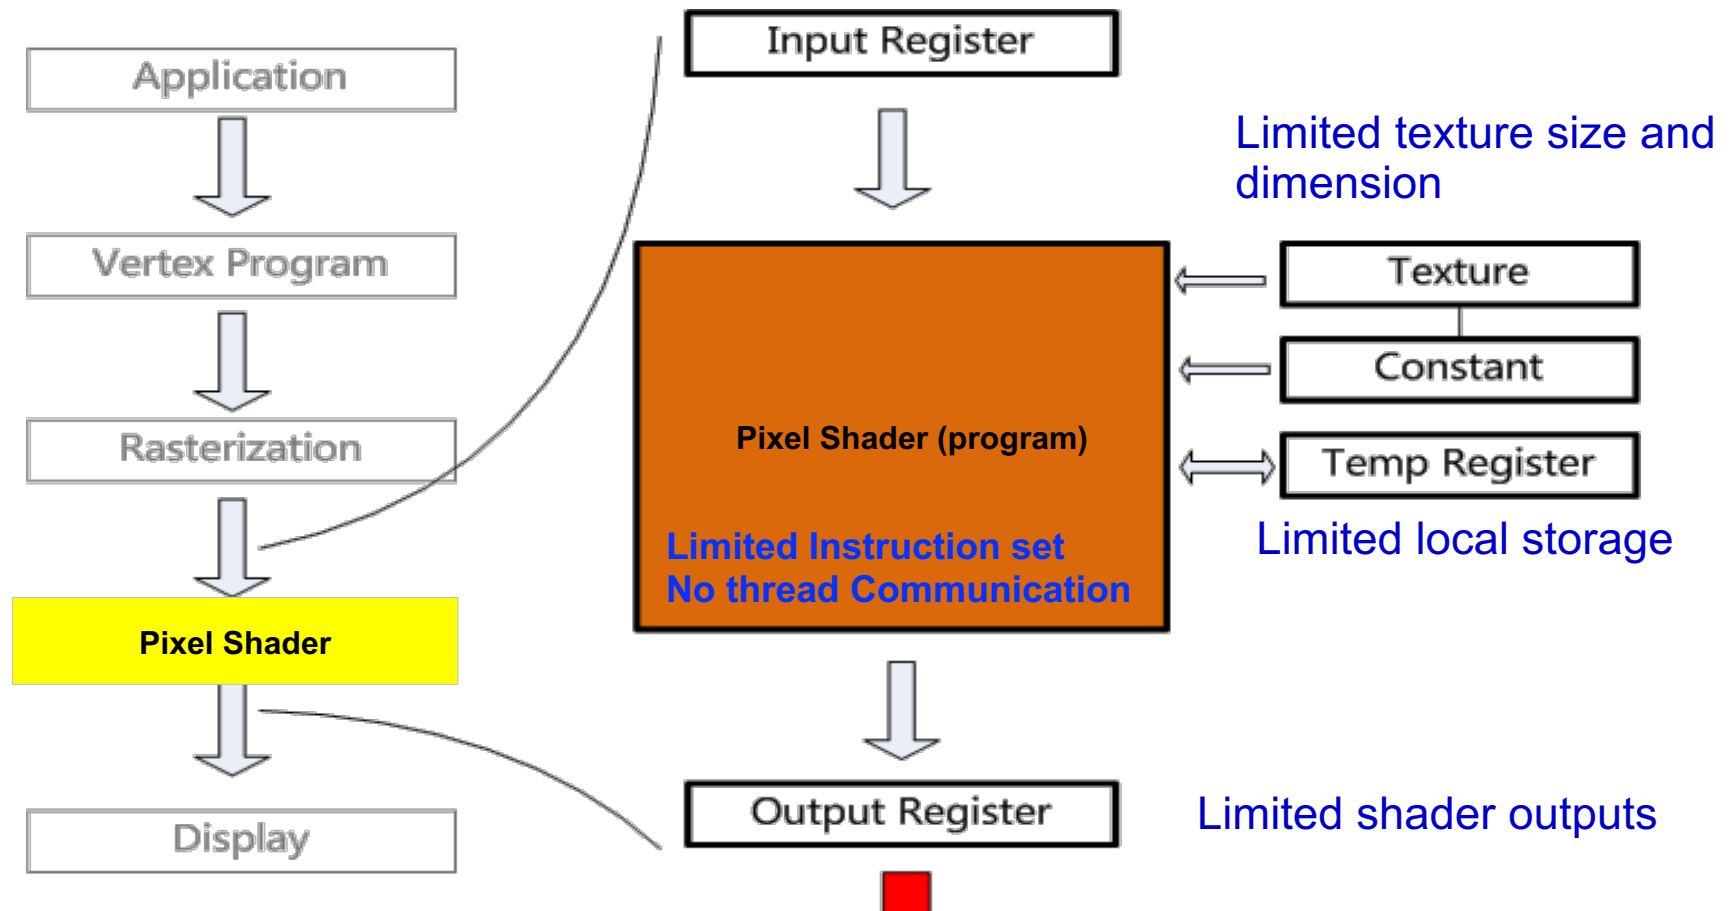
---

- Use the GPU for general-purpose
- Computing by casting problem as graphics
  - Turn data into images ("texture maps")
  - Turn algorithms into image synthesis ("rendering passes")
- **Drawback:**
  - Tough learning curve
  - Potentially high overhead of graphics API
  - Highly constrained memory layout & access model

# Before CUDA

- What's wrong with the old GPGPU programming model 1

APIs are specific to Graphics




# Before CUDA

- What's wrong with the old GPGPU programming model 2




Vertex Shader




Pixel Shader



Vertex Shader



Pixel Shader



**What is CUDA?**

# What is CUDA?

## ■ What is CUDA?: **Compute Unified Device Architecture**

- Parallel computing platform and application programming interface (API) model
- A powerful parallel programming model for issuing and managing computations on the GPU without mapping them to a graphics API

## ■ Targeted Software stack

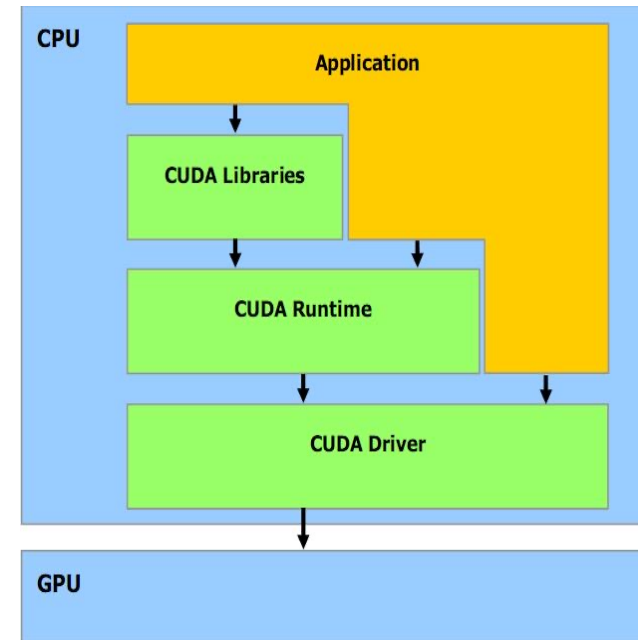
- Library, Runtime, Driver

## ■ Advantages

- SW: program the GPU in C
  - Scalable data parallel execution/memory model
  - C with minimal yet powerful extensions
- HW: fully general data-parallel architecture

## ■ Features

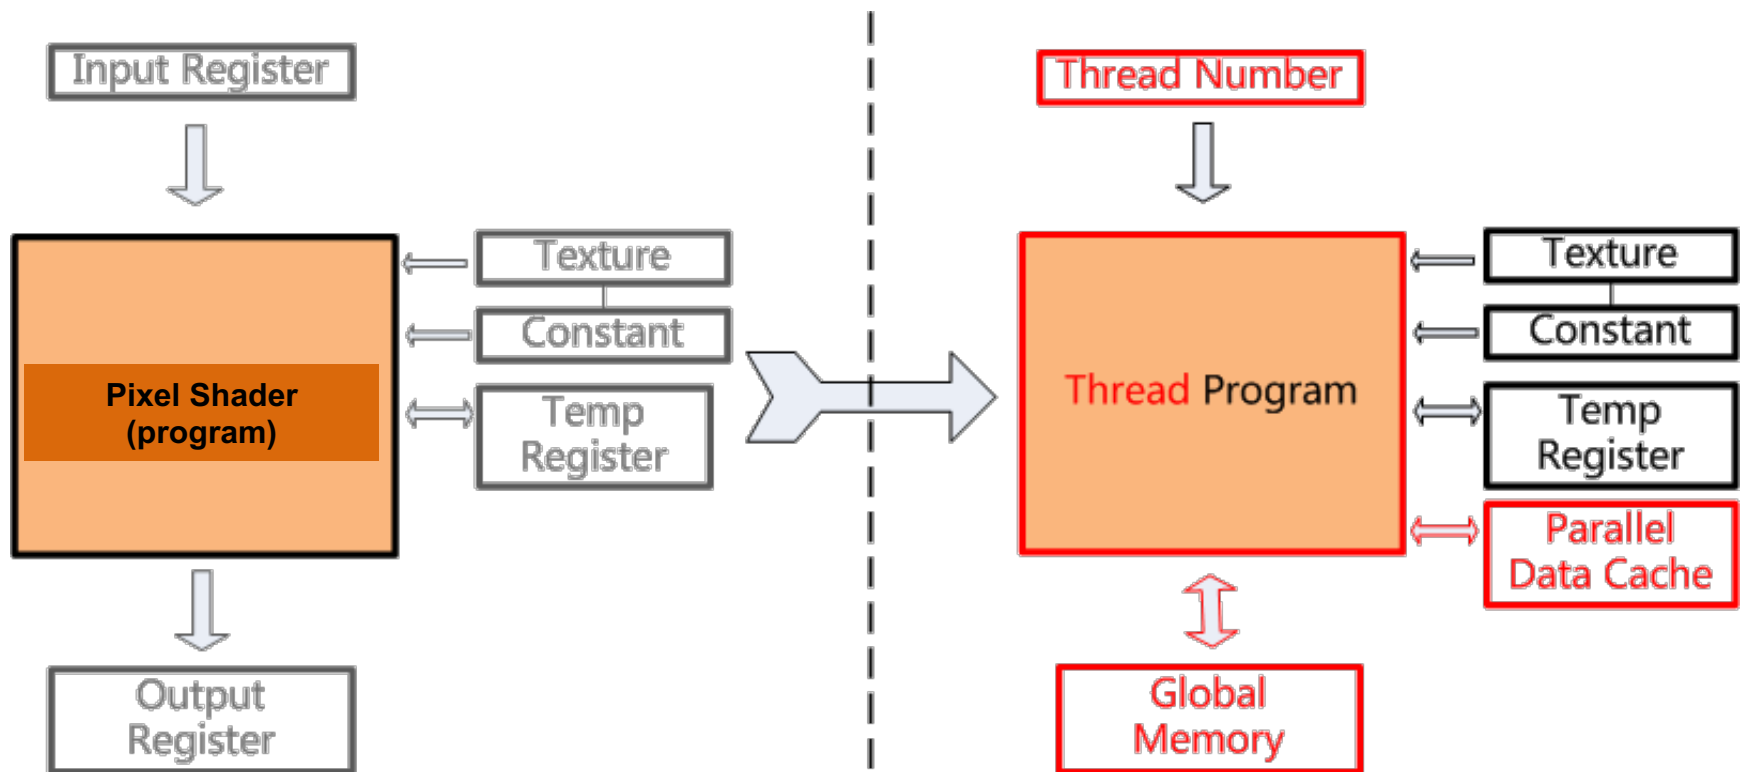
- **Heterogenous** - mixed serial-parallel programming
- **Scalable** - hierarchical thread execution model
- **Accessible** - minimal but expressive changes to C






# What is CUDA?

---



# Review : Heterogeneous Computing


- Use more than one kind of processor or cores
  - CPUs for sequential parts
  - GPUs for parallel parts



# Simple CUDA Model

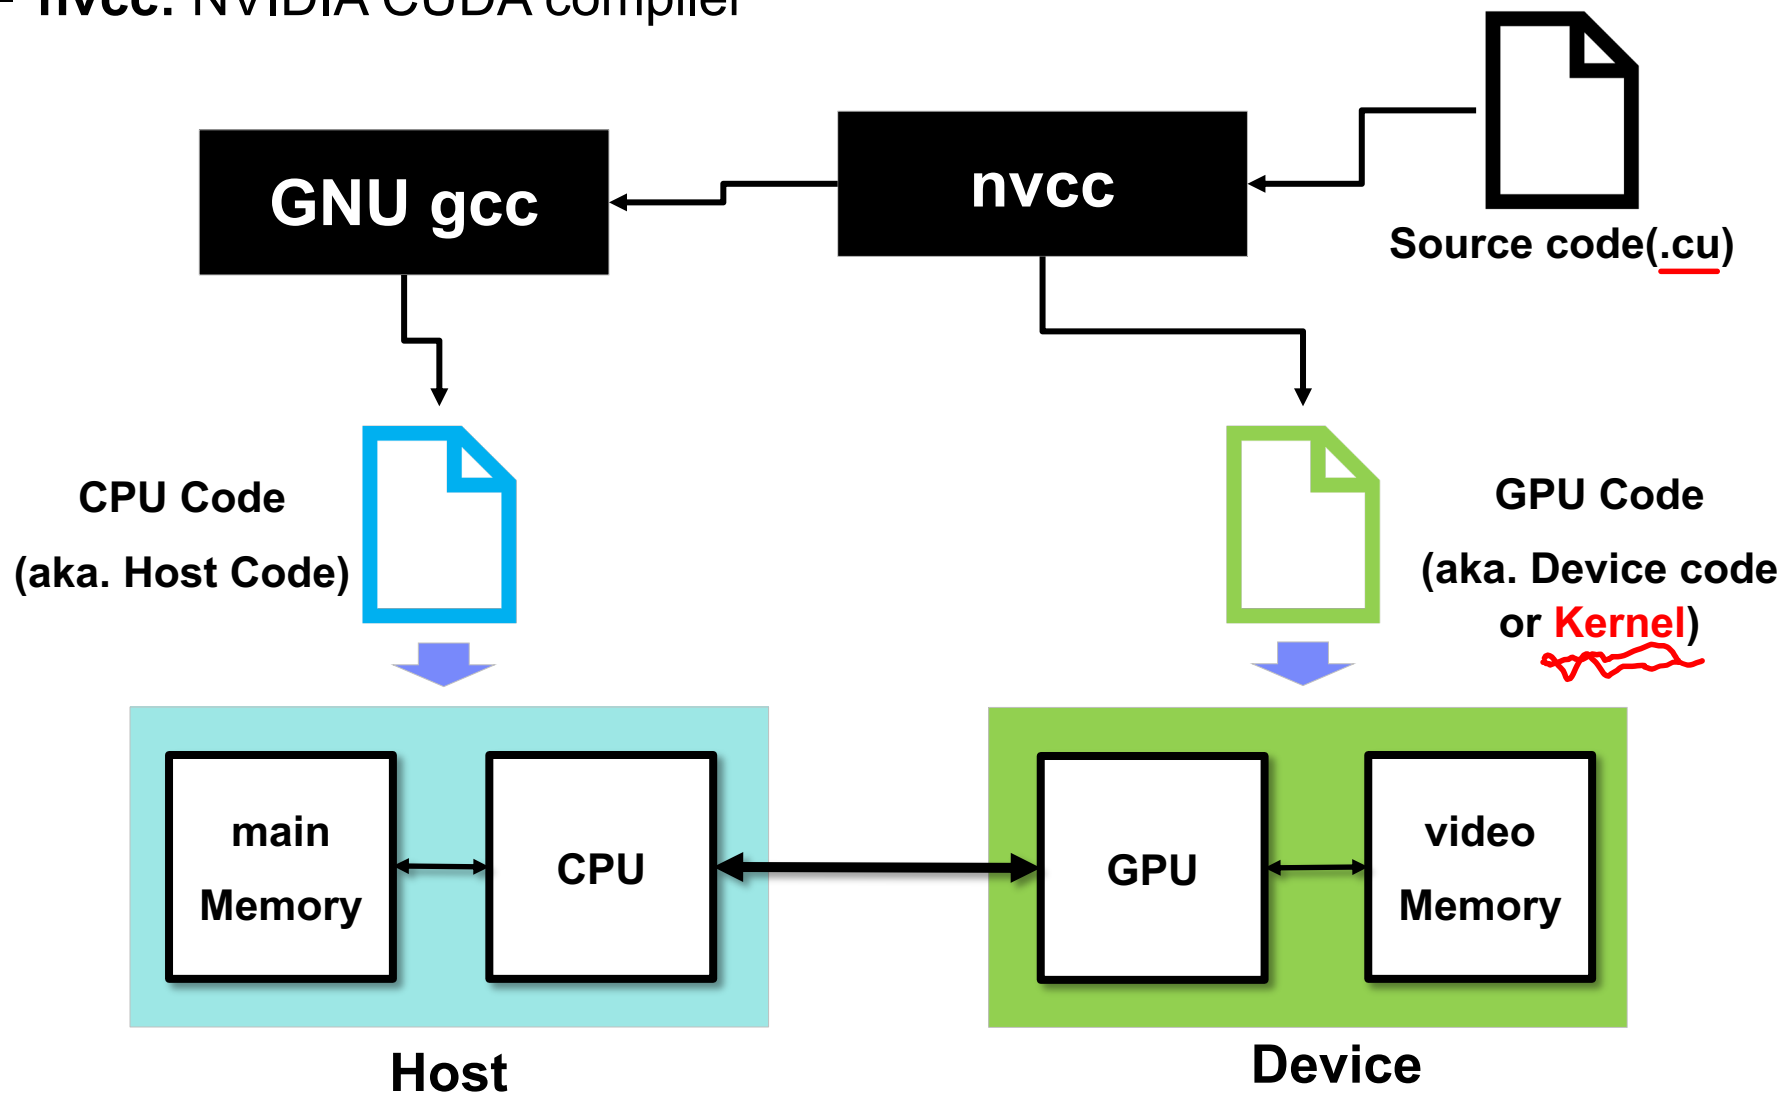
---

- **Host** : CPU + main memory (host memory)
- **Device** : GPU + video memory (device memory)



# Simple CUDA Model

- **GNU gcc** : linux c compiler
- **nvcc**: NVIDIA CUDA compiler


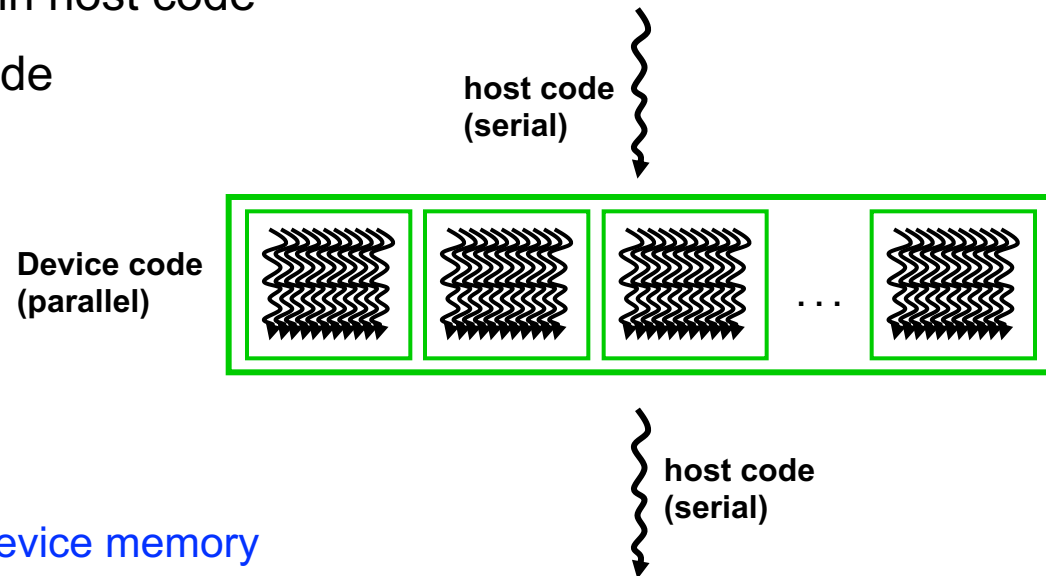


# CUDA Program Execution Scenario

- Integrated “host+device” app C program
  - Serial or modestly parallel parts in host code
  - Highly parallel parts in device code

- **Execution Scenario**


- **Step 1: host code**
  - Serial execution: read data
  - Prepare parallel execution
  - Copy data from host memory to device memory
- **Step 2: device code (kernel)**
  - Parallel processing
  - Read/write data from device memory to device memory
- **Step 3: host code**
  - Copy data from device memory to host memory
  - Serial execution: print data




# **Device Global Memory and Data Transfer**

# CUDA program uses CUDA memory

- GPU cores share the “global memory” (device memory)
  - DRAM (e.g., GDDR, HBM) is used as global memory
- To execute a kernel on a device,
  - **allocate** global memory on the device
  - **transfer** data from the host memory to allocated device memory
  - **transfer** result data from the device memory back to the host memory
  - **release** global memory



Red lines are global memory



# Memory Spaces

---

- **CPU and GPU have separate memory spaces**
  - Data is moved across data bus
  - Use functions to allocate/set/copy memory on GPU
  - Very similar to corresponding C functions
  
- **Pointers are just addresses**
  - Use pointer to access CPU and GPU memory
  - Can't tell from the pointer value whether the address is on CPU or GPU
  - **Dereferencing CPU pointer on GPU will likely crash**
  - Same for vice versa



# CPU Memory Allocation / Release

---

- **Host (CPU) manages host (CPU) memory:**
  - `void* malloc (size_t nbytes)`
  - `void* memset (void* pointer, int value, size_t count)`
  - `void free (void* pointer)`

```
int n = 1024;  
int nbytes = 1024*sizeof(int);  
int* ptr = 0;  
ptr = malloc( nbytes );  
memset( ptr, 0, nbytes);  
free( ptr );
```

# GPU Memory Allocation / Release

---

- **Host (CPU) manages device (GPU) memory:**
  - **cudaMalloc** (void\*\* pointer, size\_t nbytes)
  - **cudaMemset** (void\* pointer, int value, size\_t count)
  - **cudaFree** (void\* pointer)

```
int n = 1024;
```

```
int nbytes = 1024*sizeof(int);
```

```
int* dev_a = 0;
```

```
cudaMalloc( (void**)&dev_a, nbytes );
```

```
cudaMemset( dev_a, 0, nbytes);
```

```
cudaFree(dev_a);
```

# CUDA function rules

---

- Every library function starts with “`cuda`”
  
- Most of them returns error code (or `cudaSuccess`).
  - `cudaError_t cudaMalloc(void** devPtr, size_t size);`
  - `cudaError_t cudaFree(void* devPtr);`
  - `cudaError_t cudaMemcpy(void* dst, const void* src, size_t size);`
  
- Example:
  - `if (cudaMalloc(&devPtr, SIZE) != cudaSuccess) {  
 exit(1);  
}`

# CUDA Malloc

---

- `cudaError_t cudaMalloc( void** devPtr, size_t nbytes );`
  - allocates *nbytes* bytes of linear memory on the device
  - The start address is stored into “*devPtr*”
  - The memory is not cleared.
  - returns **cudaSuccess** or **cudaErrorMemoryAllocation**
  
- `cudaError_t cudaFree( void* devPtr );`
  - frees the memory space pointed by *devPtr*
  - if *devPtr* == 0, no operation
  - returns **cudaSuccess** or **cudaErrorInvalidDevicePointer**

# CUDA mem set

---

- `cudaError_t cudaMemset( void* devPtr, int value, size_t nbytes );`
  - fills the first ***nbytes*** byte of the memory area pointed by ***devPtr*** with the *value*
  - returns `cudaSuccess`, `cudaErrorInvalidValue`, `cudaErrorInvalidDevicePointer`

# Data Copy


---

- `cudaError_t cudaMemcpy( void* dst,  
                            void* src,  
                            size_t nbytes,  
                            enum cudaMemcpyKind direction);`
  - returns after the copy is complete
  - **blocks** CPU thread until all bytes have been copied
  - **doesn't start** copying until previous CUDA calls complete
  
- **enum cudaMemcpyKind**
  - `cudaMemcpyHostToDevice`
  - `cudaMemcpyDeviceToHost`
  - `cudaMemcpyDeviceToDevice`
  - `cudaMemcpyHostToHost`

# Data Copy

---

- host → host : `memcpy` (in C/C++)
- host → device, device → device, device → host : `cudaMemcpy` (CUDA)



# Example: Host-Device Mem copy

---

## ▪ step 1.


- make a block of data
- print out the source data

## ▪ step 2.

- copy from host memory to device memory
- copy from device memory to device memory
- copy from device memory to host memory


## ▪ step 3.

- print out the result





# Code: cuda memcpy (1/4)



```
#include <iostream>
```

```
int main(void) {
```

```
    // host-side data
```

```
    const int SIZE = 5;
```

```
    const int a[SIZE] = { 1, 2, 3, 4, 5 }; // source data
```

```
    int b[SIZE] = { 0, 0, 0, 0, 0 }; // final destination
```

```
    // print source
```

```
    printf("a = {%d,%d,%d,%d,%d}\n", a[0], a[1], a[2], a[3], a[4]);
```

# Code: cuda memcpy (2/4)

```
// device-side data
```

```
int* dev_a = 0;
```

```
int* dev_b = 0;
```


```
// allocate device memory
```

```
cudaMalloc((void**)&dev_a, SIZE * sizeof(int));
```


```
cudaMalloc((void**)&dev_b, SIZE * sizeof(int));
```

```
// copy from host to device
```

```
cudaMemcpy(dev_a, a, SIZE * sizeof(int), cudaMemcpyHostToDevice);
```



# Code: cuda memcpy (3/4)




// copy from device to device

```
cudaMemcpy(dev_b, dev_a, SIZE * sizeof(int),  
           cudaMemcpyDeviceToDevice);
```

// copy from device to host

```
cudaMemcpy(b, dev_b, SIZE * sizeof(int), cudaMemcpyDeviceToHost);
```

# Code: cuda memcpy (4/4)



```
// free device memory
```

```
cudaFree(dev_a);
```

```
cudaFree(dev_b);
```

```
// print the result
```

```
printf("b = {%d,%d,%d,%d,%d}\n", b[0], b[1], b[2], b[3], b[4]);
```

```
// done
```

```
return 0;
```

```
}
```

# Execution Result

---

- Compile the source code
  - `nvcc memcpy.cu -o ./memcpy`
  
- executing `./memcpy`
  - `a = {1,2,3,4,5}`
  - `b = {1,2,3,4,5}`

# Error Checking

# Error Checking and Handling in CUDA

---

- It is important for a program to check and handle errors
- CUDA API functions return flags that indicate whether an error has occurred
- Most of them returns error code (or `cudaSuccess`).
  - `cudaError_t` `cudaMalloc(void** devPtr, size_t size);`
  - `cudaError_t` `cudaFree(void* devPtr);`
  - `cudaError_t` `cudaMemcpy(void* dst, const void* src, size_t size);`
- Example:
  - ```
if (cudaMalloc(&devPtr, SIZE) != cudaSuccess) {  
    exit(1);  
}
```

# cudaError\_t : data type

---

- typedef enum cudaError cudaError\_t
- possible values:
  - cudaSuccess, cudaErrorMissingConfiguration, cudaErrorMemoryAllocation, cudaErrorInitializationError, cudaErrorLaunchFailure, cudaErrorLaunchTimeout, cudaErrorLaunchOutOfResources, cudaErrorInvalidDeviceFunction, cudaErrorInvalidConfiguration, cudaErrorInvalidDevice, cudaErrorInvalidValue, cudaErrorInvalidPitchValue, cudaErrorInvalidSymbol, cudaErrorUnmapBufferObjectFailed, cudaErrorInvalidHostPointer, cudaErrorInvalidDevicePointer, cudaErrorInvalidTexture, cudaErrorInvalidTextureBinding, cudaErrorInvalidChannelDescriptor, cudaErrorInvalidMemcpyDirection, cudaErrorInvalidFilterSetting, cudaErrorInvalidNormSetting, cudaErrorUnknown, cudaErrorNotYetImplemented, cudaErrorInvalidResourceHandle, cudaErrorInsufficientDriver, cudaErrorSetOnActiveProcess, cudaErrorStartupFailure, cudaErrorApiFailureBase



# cudaGetErrorName( *err* )

---

- **const char\* cudaGetErrorName( cudaError\_t err )**
  - *err* : error code to convert to string
  - returns:
    - char\* to a NULL-terminated string
    - NULL if the error code is not valid
  
- cout << cudaGetErrorName( cudaErrorMemoryAllocation )  
    << endl;
  
- cout << cudaGetErrorName( cudaErrorInvalidValue )  
    << endl;
- shows:
  - cudaErrorMemoryAllocation**
  - cudaErrorInvalidValue**

# cudaGetErrorString( *err* )

---

- **const char\* cudaGetErrorString( cudaError\_t err )**
  - *err* : error code to convert to string
  - returns:
    - char\* to a NULL-terminated string
    - NULL if the error code is not valid
  
- `cout << cudaGetErrorString( cudaErrorMemoryAllocation )`  
`<< endl;`
  
- `cout << cudaGetErrorString( cudaErrorInvalidValue )`  
`<< endl;`
  - shows:
    - out of memory**
    - invalid argument**

# cudaGetLastError( void )

---

- `cudaError_t cudaGetLastError( void )`

- returns the last error due to CUDA runtime calls in the same host thread
- and **resets** it to **cudaSuccess**
- So, if no CUDA error since the last call, it returns **cudaSuccess**
- For multiple errors, it contains the last error only.

- `cudaError_t cudaPeekAtLastError( void )`

- returns the last error due to CUDA runtime calls in the same host thread
- Note that this call does **NOT reset**
- So, the last error code is still available

# A simple CUDA error check code

---

```
cudaMemcpy( ... );  
cudaError_t e = cudaGetLastError();  
if (e != cudaSuccess) {  
    printf("cuda failure %s:%d: '%s'\n",  
        __FILE__, __LINE__,  
        cudaGetErrorString(e) );  
    exit(0);  
}
```

# A simple CUDA error check macro

```
#define  cudaCheckError( )      do { \
    cudaError_t e = cudaGetLastError(); \
    if (e != cudaSuccess) { \
        printf("cuda failure %s:%d: '%s'\n", \
            __FILE__, __LINE__, \
            cudaGetErrorString(e) ); \
        exit(0); \
    } \
} while (0)
```

# Example

---

- code segment

// allocate device memory

```
cudaMalloc((void**)&dev_a, sizeof(int));
```

```
cudaMalloc((void**)&dev_b, sizeof(int));
```

```
cudaMalloc((void**)&dev_c, sizeof(int));
```

```
cudaCheckError( );
```

# More advanced macro

```
#ifdef DEBUG // debug mode

#define CUDA_CHECK(x)      do {\
    (x); \
    cudaError_t e = cudaGetLastError(); \
    if (cudaSuccess != e) {\
        printf("cuda failure %s at %s:%d\n", \
            cudaGetErrorString(e), \
            __FILE__, __LINE__); \
        exit(1); \
    } \
} while (0)

#else

#define CUDA_CHECK(x)      (x)          // release mode

#endif
```

# error\_check.cu

```
#include <iostream>

#ifdef DEBUG
#define CUDA_CHECK(x) do {\
    (x); \
    cudaError_t e = cudaGetLastError(); \
    if (cudaSuccess != e) { \
        printf("cuda failure \"%s\" at %s:%d\n", \
            cudaGetErrorString(e), \
            __FILE__, __LINE__); \
        exit(1); \
    } \
} while (0)
#else
#define CUDA_CHECK(x) (x)
#endif
```



# error\_check.cu

---

```
// main program for the CPU
int main(void) {
    // host-side data
    const int SIZE = 5;
    const int a[SIZE] = { 1, 2, 3, 4, 5 };
    int b[SIZE] = { 0, 0, 0, 0, 0 };
    // print source
    printf("a = {%d,%d,%d,%d,%d}\n", a[0], a[1], a[2], a[3], a[4]);
    // device-side data
    int *dev_a = 0;
    int *dev_b = 0;
    // allocate device memory
    CUDA_CHECK( cudaMalloc((void**)&dev_a, SIZE * sizeof(int)) );
    CUDA_CHECK( cudaMalloc((void**)&dev_b, SIZE * sizeof(int)) );
    // copy from host to device
    CUDA_CHECK( cudaMemcpy(dev_a, a, SIZE * sizeof(int), cudaMemcpyDeviceToDevice) ); // BOMB here !
    // copy from device to device
    CUDA_CHECK( cudaMemcpy(dev_b, dev_a, SIZE * sizeof(int), cudaMemcpyDeviceToDevice) );
    // copy from device to host
    CUDA_CHECK( cudaMemcpy(b, dev_b, SIZE * sizeof(int), cudaMemcpyDeviceToHost) );
    // free device memory
    CUDA_CHECK( cudaFree(dev_a) );
    CUDA_CHECK( cudaFree(dev_b) );
    // print the result
    printf("b = {%d,%d,%d,%d,%d}\n", b[0], b[1], b[2], b[3], b[4]);
    // done
    return 0;
}
```

# Execution Result

---

- Compile the source code
  - `nvcc error_check.cu -DDEBUG -o ./error_check`
  - `nvcc error_check.cu -o ./error_check`
  
- executing `./error_check`
  - `a = {1,2,3,4,5}`
  - `b = {.....}`

# Agenda

---

- What is CUDA?
- Device Global Memory and Data Transfer
- Error Checking
- **A Vector Addition Kernel**
- **Kernel Functions and Threading**
- **Kernel Launch**

# CUDA Resources

---

## ■ CUDA API reference:

- <http://docs.nvidia.com/cuda/index.html>
- <http://docs.nvidia.com/cuda/cuda-runtime-api/index.html>

## ■ CUDA course:

- <https://developer.nvidia.com/cuda-education-training>
- <https://developer.nvidia.com/cuda-training>

