

CUDA Thread 1

Prof. Seokin Hong

Agenda

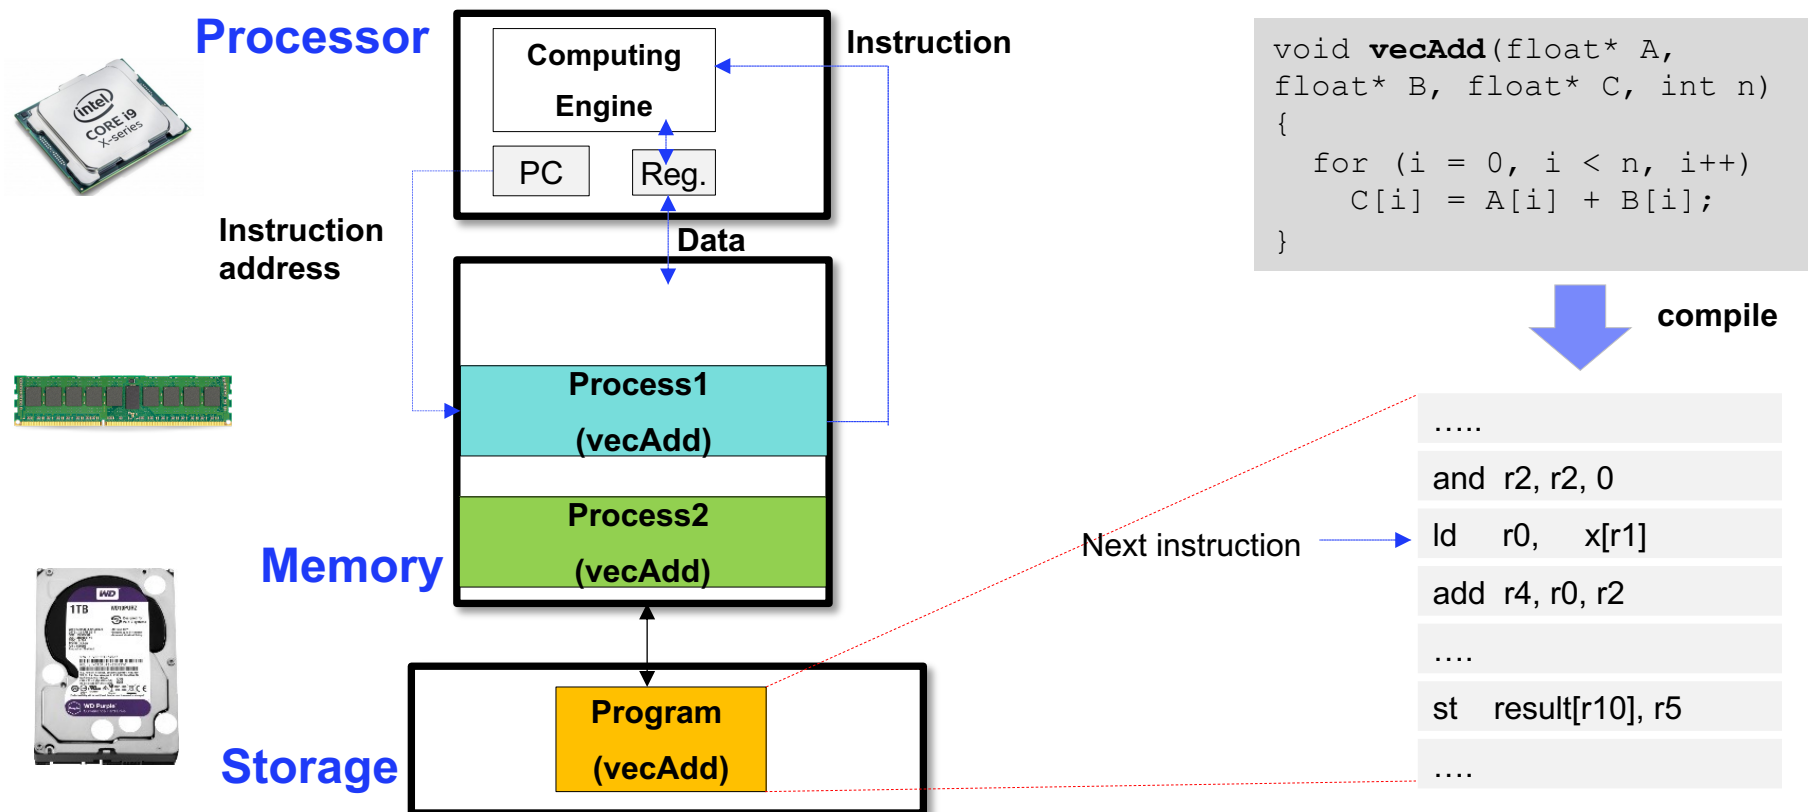
- What is Thread?
- CUDA Thread Organization
- Extended VecAdd.cu
- Matrix Addition Kernel
- Mapping Threads to Multidimensional Data
- Synchronization and Transparent Scalability
- Resource Assignment
- Querying Device Properties
- Thread Scheduling and Latency Tolerance

What is Thread?

Process and Thread

■ Process

- An instance of a computer program that is being executed
- Program code + Execution state (address of the next instruction, register state, memory contents)



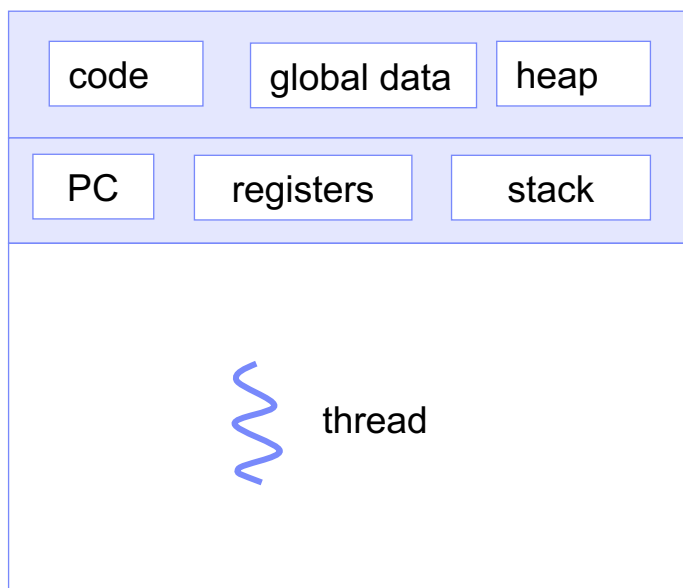
Process and Thread (Cont'd)

■ Process

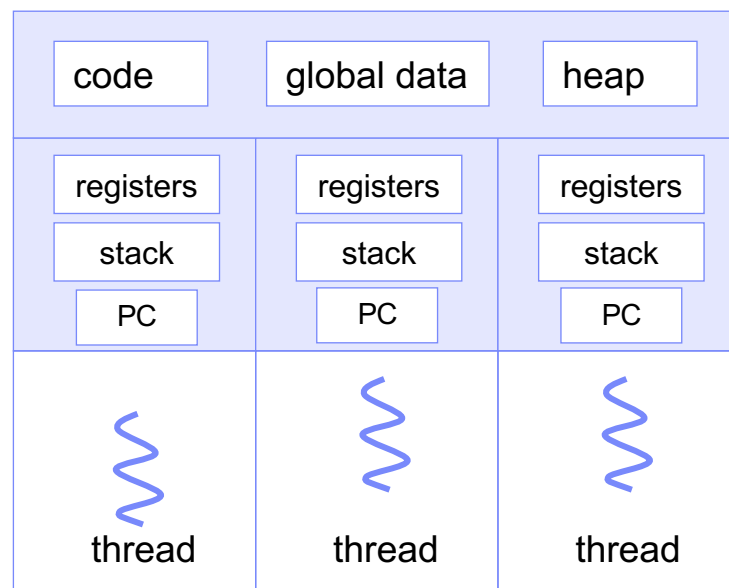
- **An instance of a computer program that is being executed**
- Program code + Execution state (address of the next instruction, register state, memory contents)

■ Thread

- **A flow of execution within a process**
- Execution state (address of the next instruction, register state, stack)

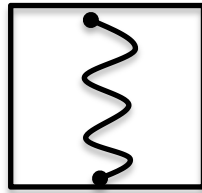


single-threaded process

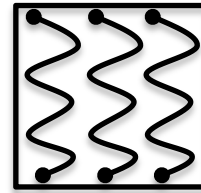


multithreaded process

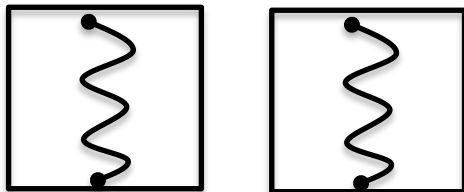
Process and Thread (Cont'd)



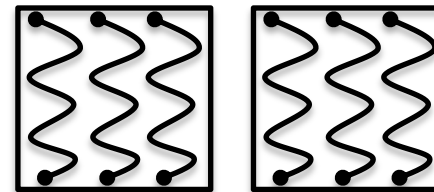
**one process,
one thread**



**one process,
multiple threads**



**multiple processes,
one thread per process**



**multiple processes,
multiple threads per process**

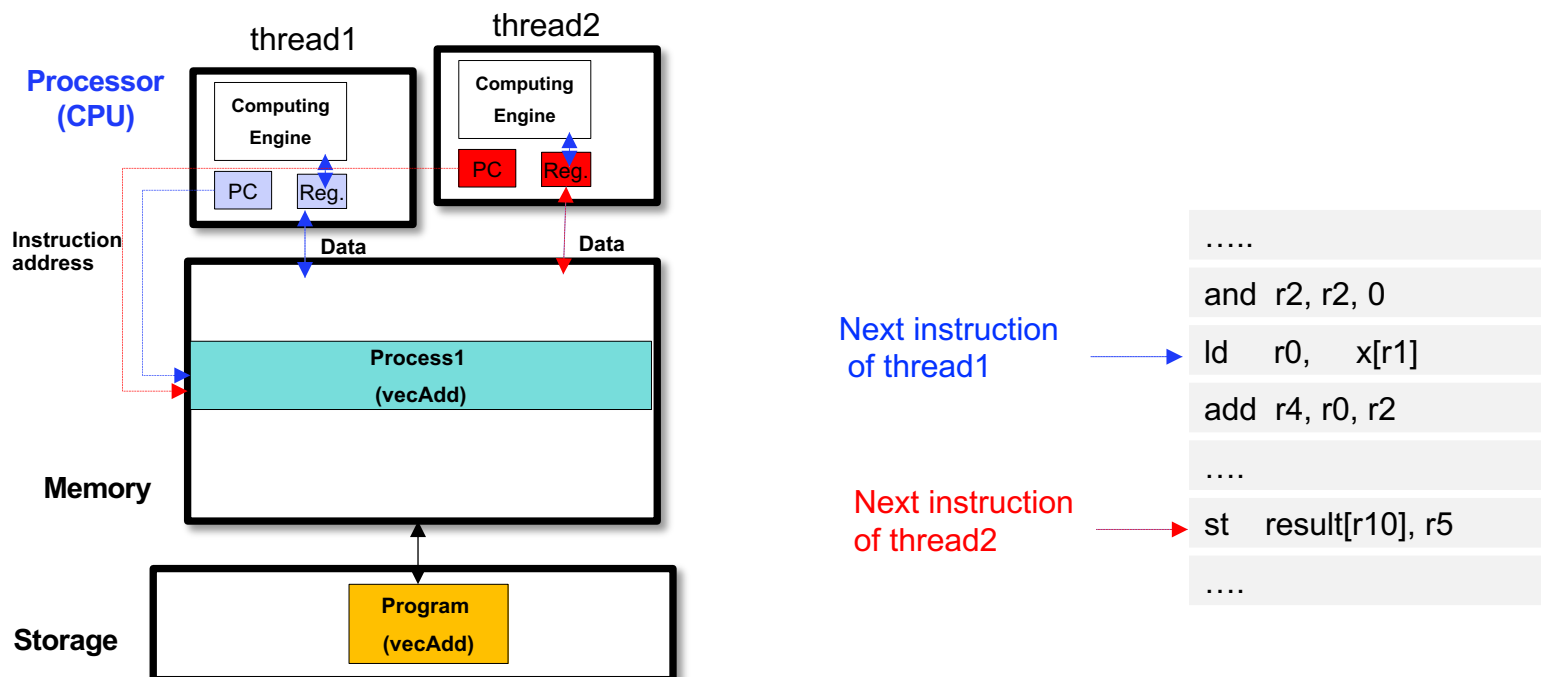
Process and Thread (Cont'd)

■ Process

- An instance of a computer program that is being executed
- Program code + Execution state (address of the next instruction, register state, memory contents)

■ Thread

- A flow of execution within a process
- Execution state (address of the next instruction, register state)



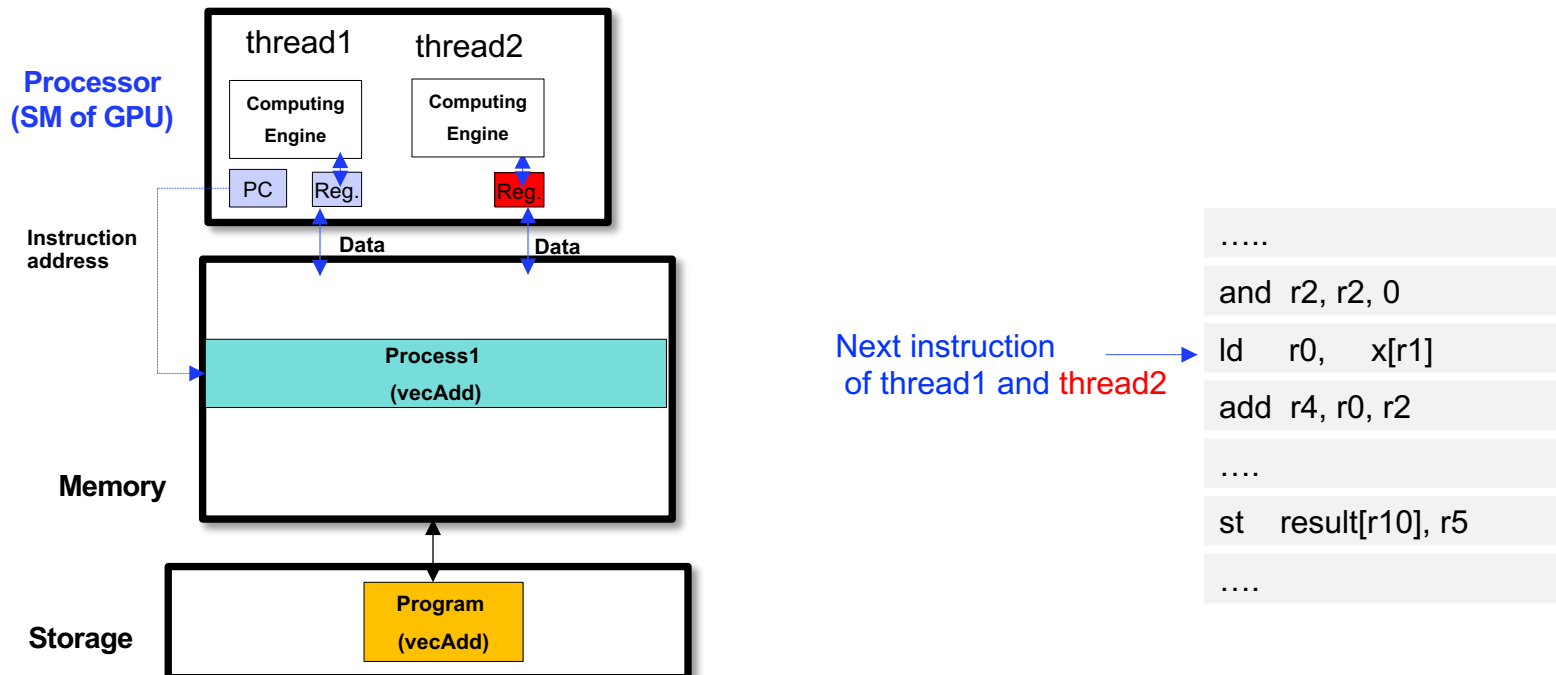
Process and Thread (Cont'd)

■ Process

- An instance of a computer program that is being executed
- Program code + Execution state (address of the next instruction, register state, memory contents)

■ Thread

- A flow of execution within a process
- Execution state (address of the next instruction, register state)



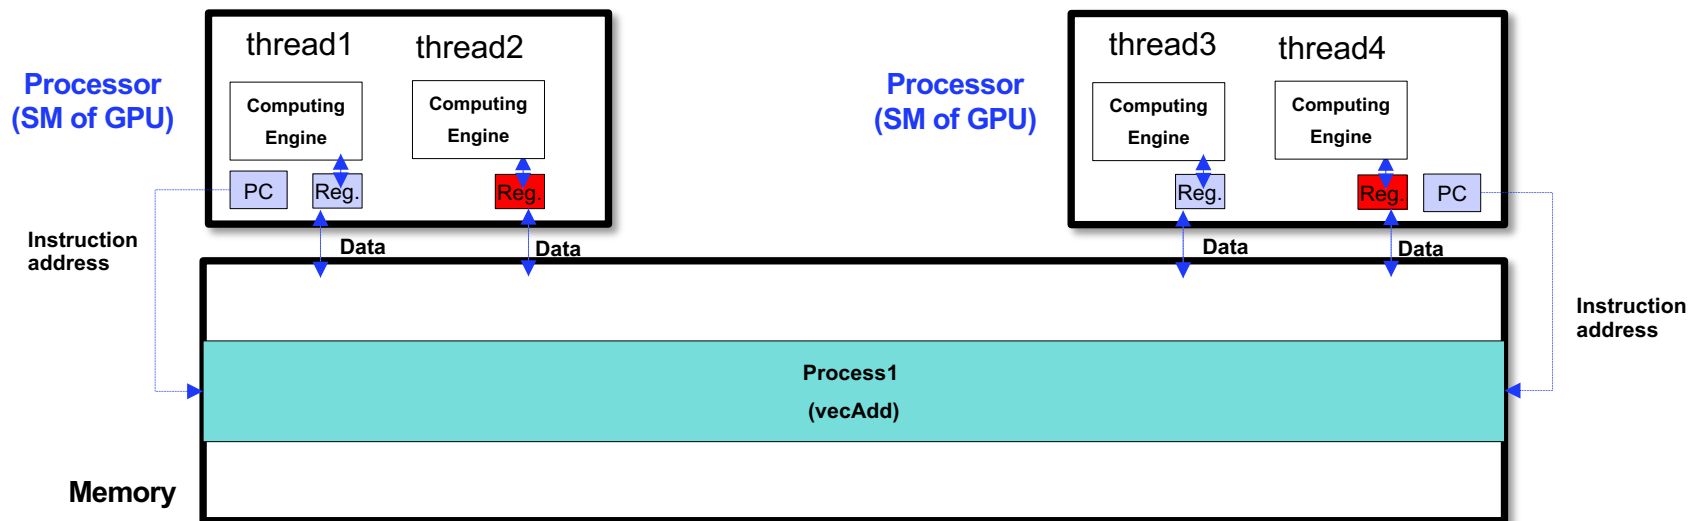
Process and Thread (Cont'd)

■ Process

- An instance of a computer program that is being executed
- Program code + Execution state (address of the next instruction, register state, memory contents)

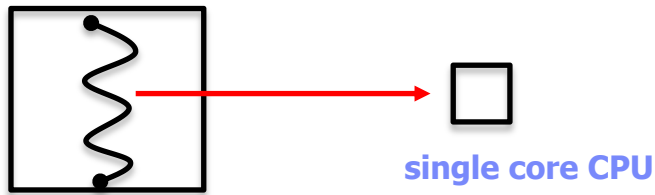
■ Thread

- A flow of execution within a process
- Execution state (address of the next instruction, register state)



Thread Execution

- On Single Core Processors
 - single thread → Serial processing



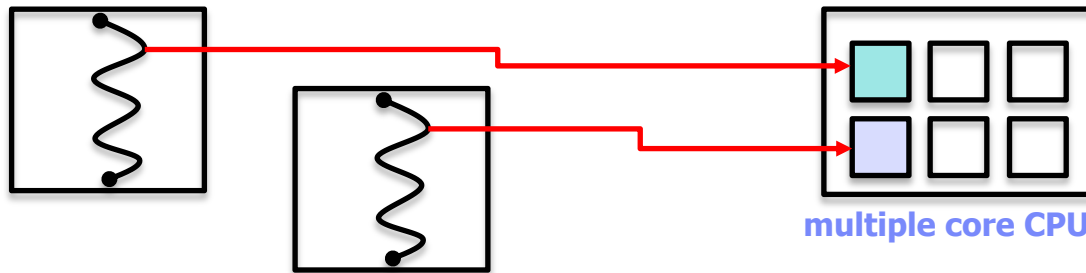
- multiple thread → time sharing (concurrent processing)



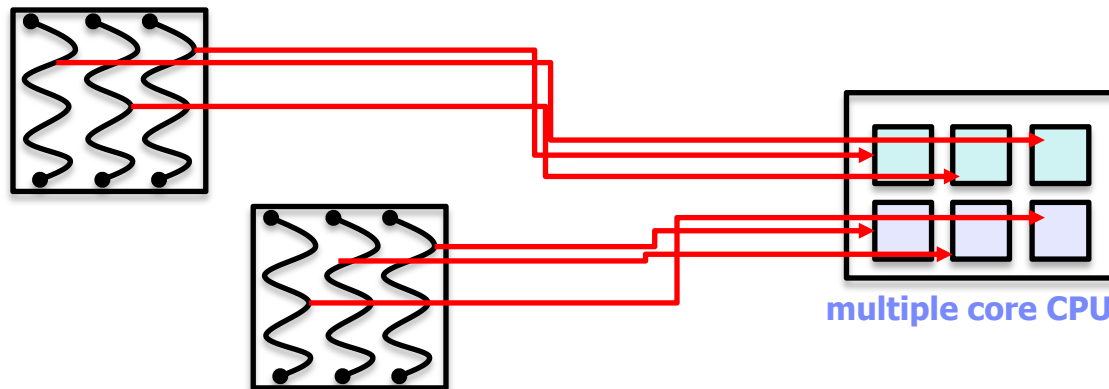
Thread Execution on Different Types of Processor (Cont'd)

■ On Multicore Processors

- single thread, multiple process → parallel processing



- multiple thread → parallel processing



CUDA Thread Organization

Hierarchy of Threads

- Kernel is composed of many threads
 - All threads execute same sequential program
 - Use parallel threads rather than sequential loop

- Threads are grouped into thread blocks

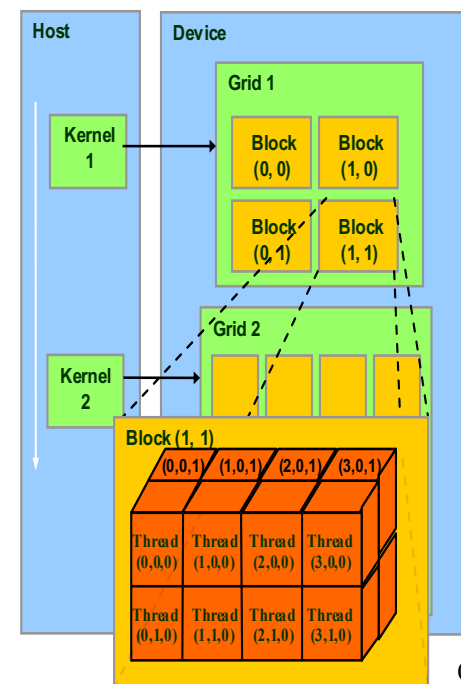
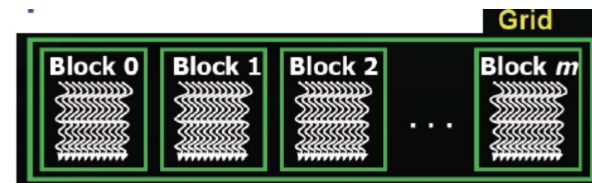
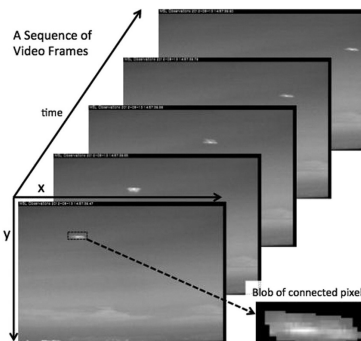
- Threads in block can share data with a shared memory

- Blocks are grouped into grids

- Why hierarchy of threads?

- Real-world systems(data) are organized hierarchically

- Video file
 - › Pixel
 - › Image : multiple pixels
 - › Video : A sequence of frames



Courtesy: NDVIA

IDs and Dimensions

- Threads and blocks have unique IDs

- Threads:**

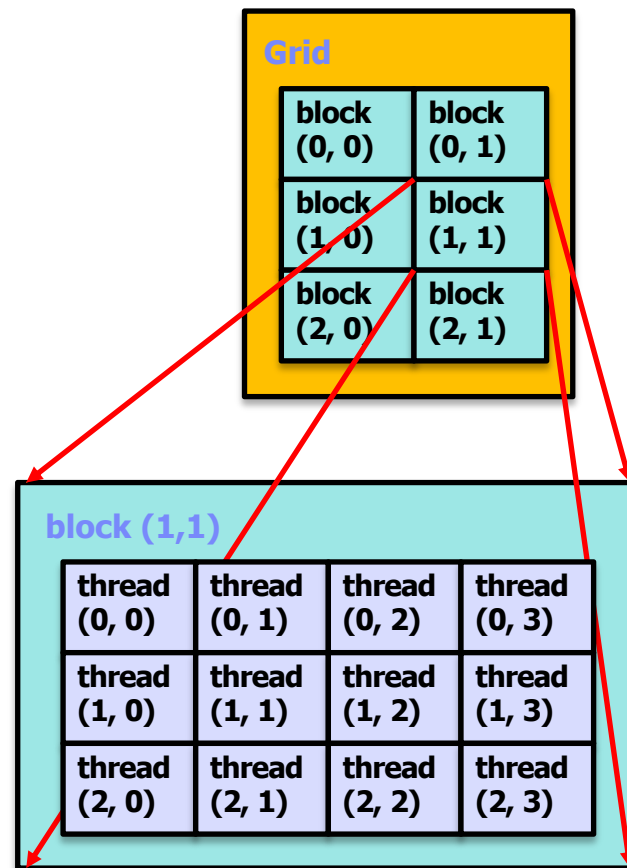
- `threadIdx`: 1D, 2D, or 3D
- unique within a block

- Blocks:**

- `blockIdx`: 1D, 2D, or 3D
- unique within a grid

- Dimensions set at launch**

- Can be unique for each grid
- `blockDim`: dimension of block
- `gridDim`: dimension of grid



Each block and thread is labelled with (y, x)

CUDA pre-defined variables

■ Pre-defined variables

- dim3 **gridDim** dimensions of grid
- dim3 **blockDim** dimensions of block
- uint3 **blockIdx** block index within grid
- uint3 **threadIdx** thread index within block
- int **warpSize** number of threads in warp

■ dim3 can take 1, 2, or 3 arguments (x, y, z):

- dim3 blocks1D(5);
- dim3 blocks2D(5, 5);
- dim3 blocks3D(5, 5, 5);

Configuring Thread Organization

- A kernel function must be called with an execution configuration:

- **Example1**

```
KernelFunc<<< blocks, threads >>>(...);
```

Or

- **Example2**

- Launch the KernelFunc() kernel function, and generate a 1D grid that consists of 32 blocks, each of which consists of 128 threads

```
dim3 DimGrid(32, 1, 1);  
dim3 DimBlock(128, 1, 1);  
KernelFunc<<< DimGrid, DimBlock >>>(...);
```

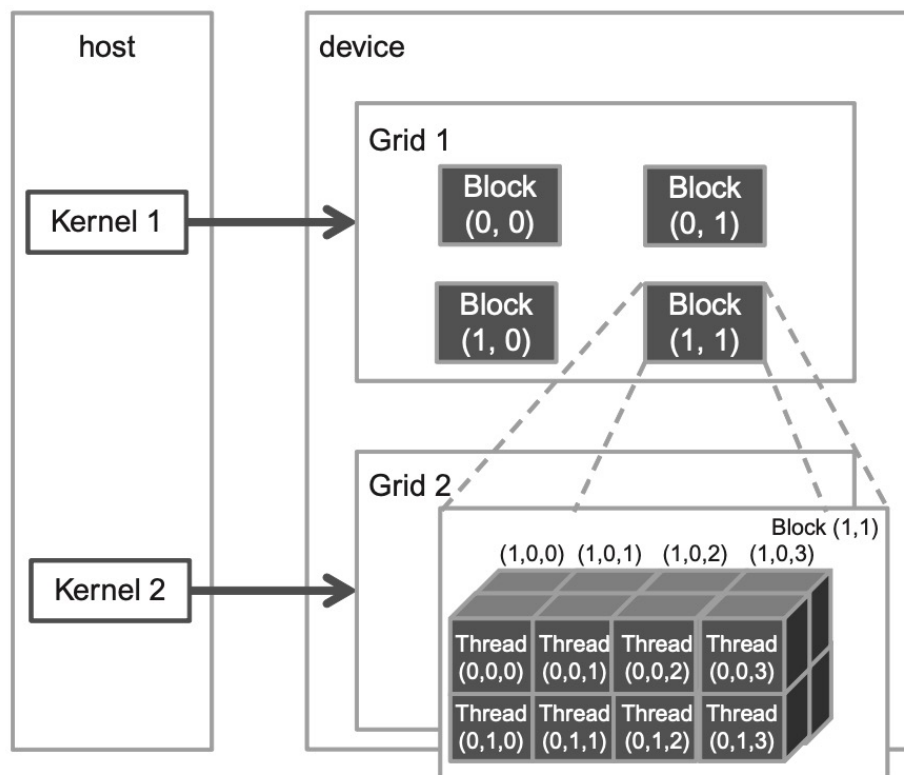
- **dim3** : a C struct with three unsigned integer fields: x, y, z
- Total number of threads in the grid is $128 \times 32 = 4096$

Configuring Thread Organization

- A kernel function must be called with an execution configuration:

- **Example3**

```
dim3 DimGrid(2, 2, 1);  
dim3 DimBlock(4, 2, 2);  
KernelFunc<<< DimGrid, DimBlock >>>(...);
```



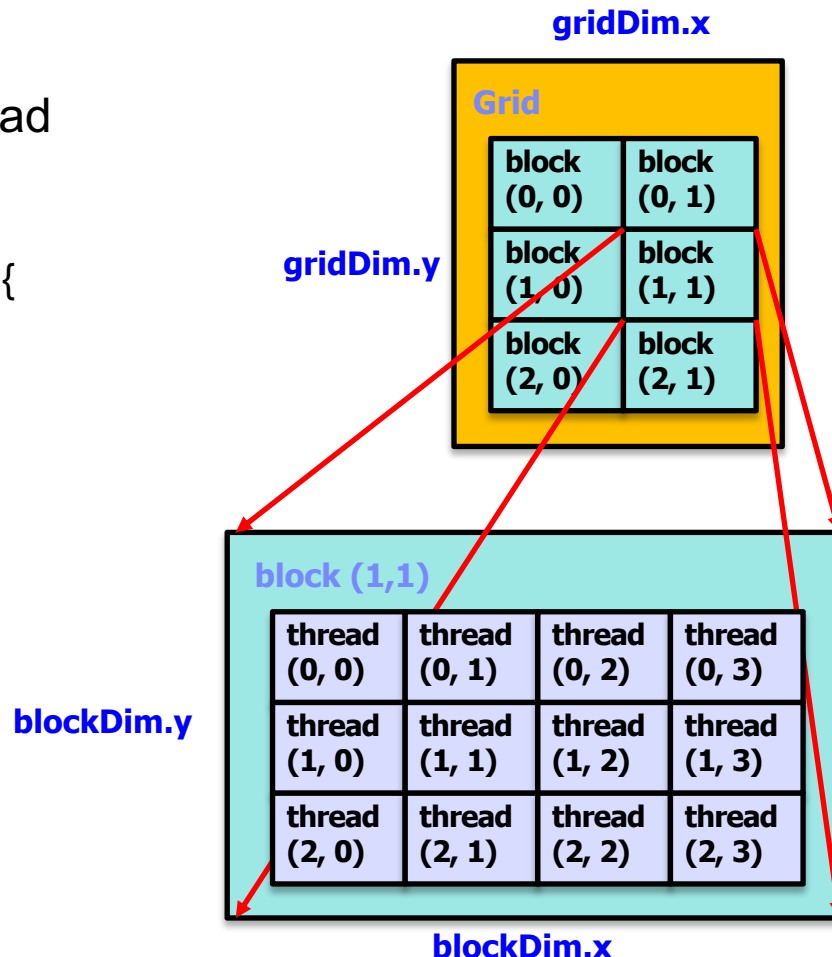
Kernel with 2D Indexing

dim3 DimGrid(100, 50); → **gridDim** in kernel

dim3 DimBlock(4, 8, 8); → **blockDim** in kernel

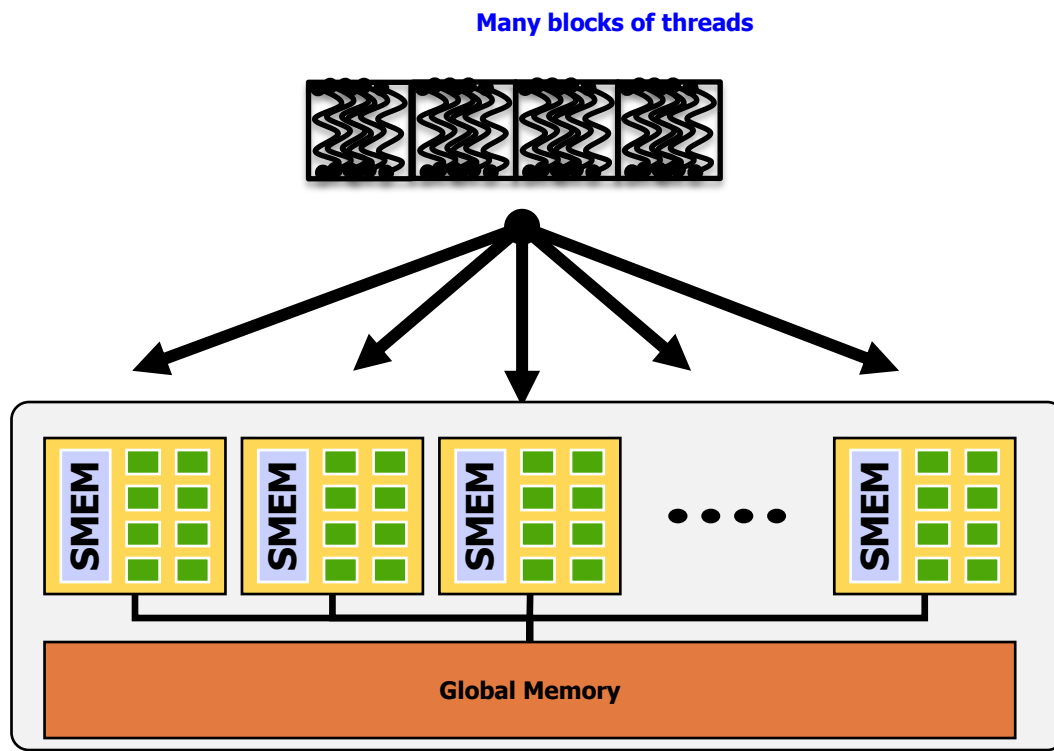
blockIdx, **threadIdx** : unique for each thread

```
__global__ void kernel( int *a, int dimx, int dimy ) {  
    int ix  = blockIdx.x*blockDim.x + threadIdx.x;  
    int iy  = blockIdx.y*blockDim.y + threadIdx.y;  
    int idx = iy*dimx + ix;  
    a[idx] = a[idx]+1;  
}
```

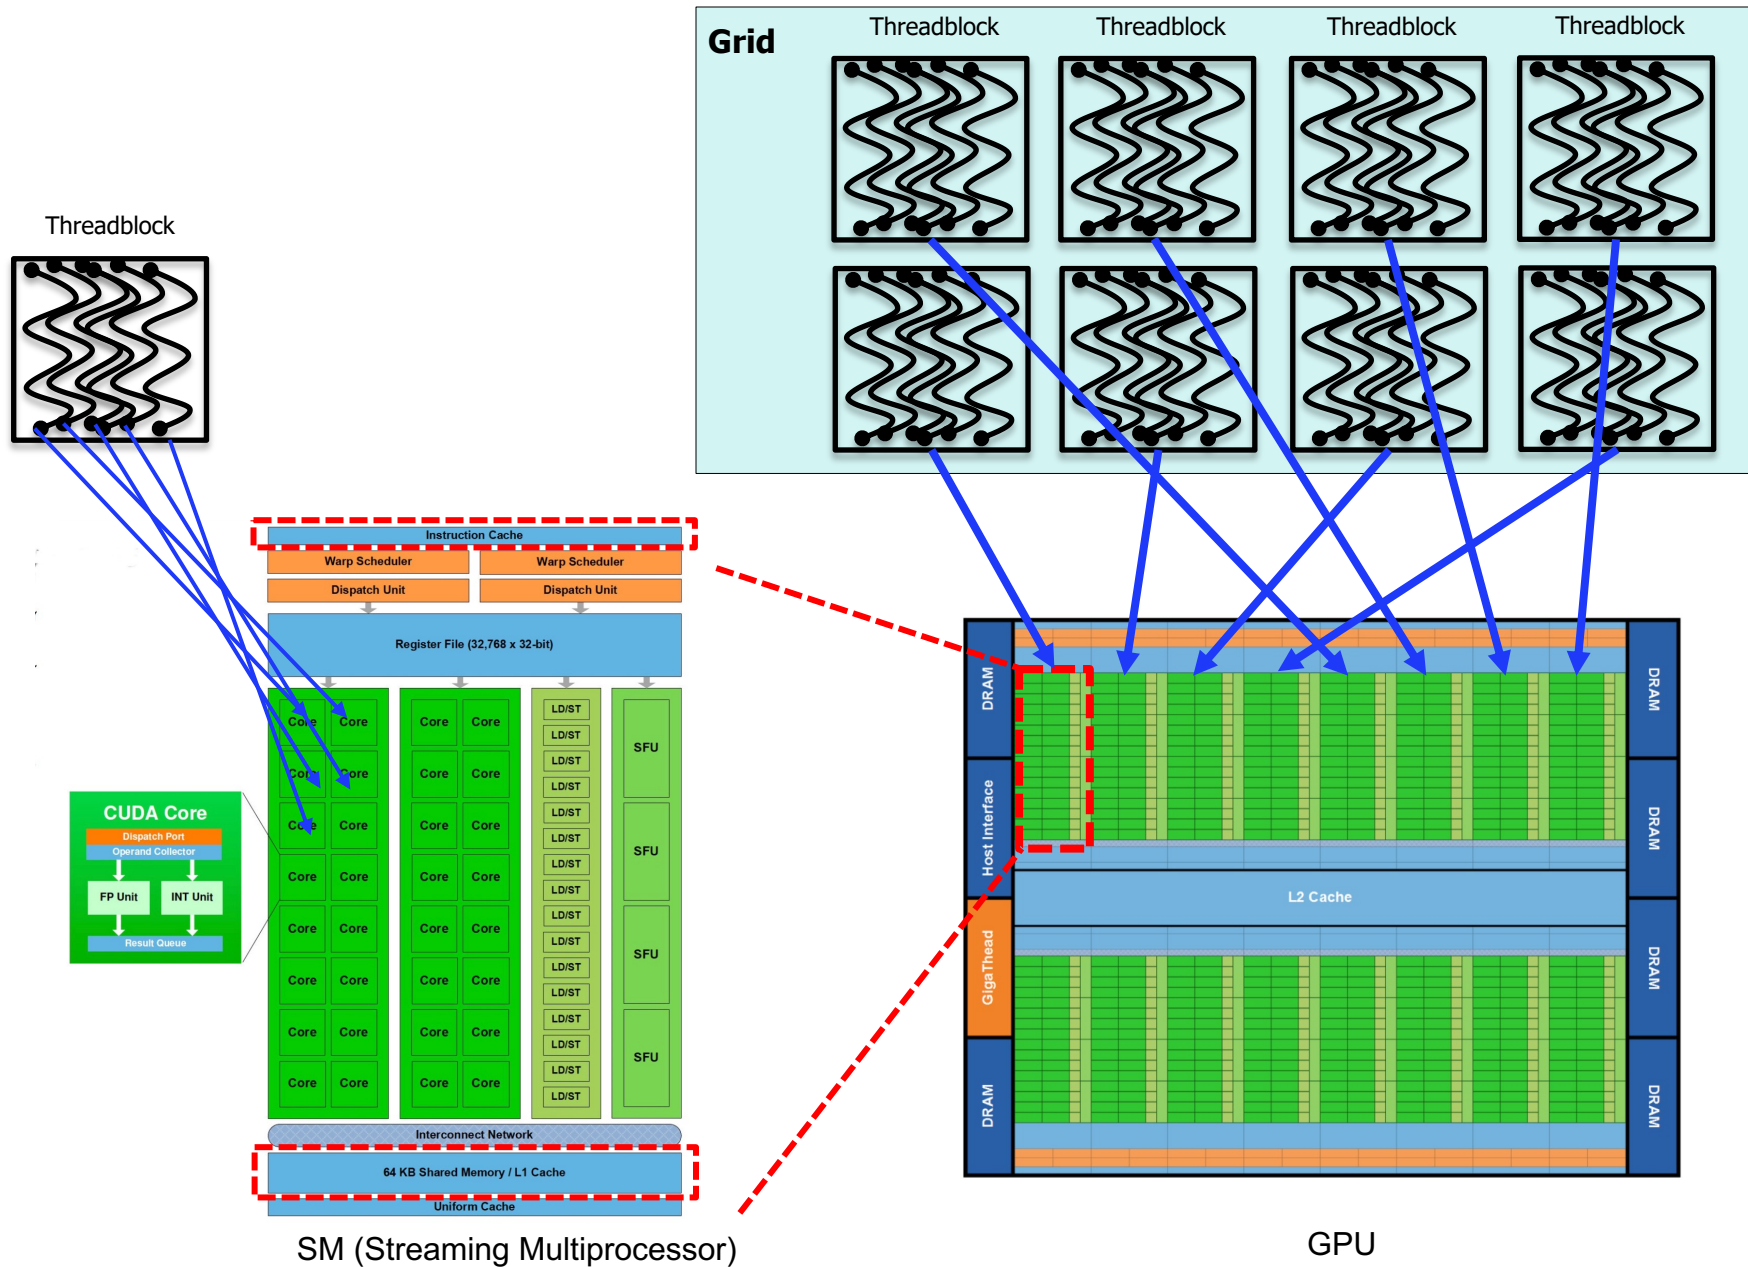


Executing threads on GPU

- Each thread block goes to a SM
 - Threads running on the same SM can share data through shared memory (SMEM)
 - Shared memory is much faster than the global memory



Executing threads on GPU



Extending Vector Addition Program for scalable parallel execution

VecAdd_ext.cu

Host code

```
void vecAdd(int* A, int* B, int* C, int n)
{
    int size = n * sizeof(int);
    int* A_d=0;
    int* B_d=0;
    int* C_d=0;
    // Allocate device memory
    cudaMalloc((void **) &A_d, size);
    cudaMalloc((void **) &B_d, size);
    cudaMalloc((void **) &C_d, size);

    // Transfer A and B to device memory
    cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);

    addKernel<<<1, n>>>(A_d, B_d, C_d);

    // Transfer C from device to host
    cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(A_d); cudaFree(B_d); cudaFree(C_d);
}
```



```
void vecAdd(int* A, int* B, int* C, int n)
{
    int size = n * sizeof(int);
    int* A_d=0;
    int* B_d=0;
    int* C_d=0;
    // Allocate device memory
    cudaMalloc((void **) &A_d, size);
    cudaMalloc((void **) &B_d, size);
    cudaMalloc((void **) &C_d, size);

    // Transfer A and B to device memory
    cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);

    addKernel<<<ceil(n/256.0), 256>>>(A_d, B_d, C_d,n);

    // Transfer C from device to host
    cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(A_d); cudaFree(B_d); cudaFree(C_d);
}
```

- $\text{ceil}(x)$: maps x to the least integer greater than or equal to x
ex) $\text{ceil}(4.5) = 5$

Extended Vector Addition Kernel

```
__global__  
void addKernel(int* A_d, int* B_d, int* C_d)  
{  
    // each thread knows its own index  
    int i = threadIdx.x;  
    C_d[i] = A_d[i] + B_d[i];  
}
```



```
__global__  
void addKernel(int* A_d, int* B_d, int* C_d, int n)  
{  
    int i = blockDim.x*blockIdx.x + threadIdx.x;  
    if(i<n) C_d[i] = A_d[i] + B_d[i];  
}
```

Host code

```
int main(void) {
    const int SIZE = 2048;
    int a[SIZE];
    int b[SIZE];
    int c[SIZE];

    for(int i=0;i<SIZE;i++)
    {
        a[i]=i;
        b[i]=i;
    }
    vecAdd(a,b,c, SIZE );

    // print the result
    for(int i=0;i<SIZE;i++)
        printf("%d\n", c[i]);

    // done
    return 0;
}
```


Scalable Parallel Execution

- **addKernel**<<<ceil(n/256.0), 256>>>(A_d, B_d, C_d,n);
 - The number of threads per block: 256
 - The number of blocks can be different depending on the total number of threads
 - If n is 750, 3 blocks will be used
 - If n is 4000, 16 thread blocks will be used
 - If n is 2000000, 7813 blocks will be used
 - All the thread blocks operate on different parts of the vectors
 - The thread block can be executed in any arbitrary order
 - Programmers must not make any assumptions regarding execution order
 - **CUDA supports Scalable Parallel Execution**
 - A small GPU with a small amount of SM may execute only one or two of the thread blocks in parallel
 - A larger GPU may execute 64 or 128 blocks in parallel
- So, same code runs at lower speed on small GPU and higher speed on larger GPU

Matrix Addition Kernel

Matrix Addition

- We want to add matrix A and matrix B

$$\mathbf{C} = \mathbf{A} + \mathbf{B} = \begin{bmatrix} a_{00} & a_{01} & a_{02} & a_{03} & a_{04} \\ a_{10} & a_{11} & a_{12} & a_{13} & a_{14} \\ a_{20} & a_{21} & a_{22} & a_{23} & a_{24} \\ a_{30} & a_{31} & a_{32} & a_{33} & a_{34} \\ a_{40} & a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} + \begin{bmatrix} b_{00} & b_{01} & b_{02} & b_{03} & b_{04} \\ b_{10} & b_{11} & b_{12} & b_{13} & b_{14} \\ b_{20} & b_{21} & b_{22} & b_{23} & b_{24} \\ b_{30} & b_{31} & b_{32} & b_{33} & b_{34} \\ b_{40} & b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix}$$

- We need a two-dimensional kernel !

- $c_{ij} = a_{ij} + b_{ij}$

Memory Layout of Matrix

- row-major matrix storage

- logical layout:

$$\begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{bmatrix}$$

- physical layout: 1D array

$$[a_{00} \quad a_{01} \quad a_{02} \quad a_{10} \quad a_{11} \quad a_{12} \quad a_{20} \quad a_{21} \quad a_{22}]$$

- re-interpret:

$$[a_0 \quad a_1 \quad a_2 \quad a_3 \quad a_4 \quad a_5 \quad a_6 \quad a_7 \quad a_8]$$

- index change:

- $\text{idx} = y * \text{WIDTH} + x$

```
#include <iostream>
```

```
int main(void) {
```

```
    // host-side data
```

```
    const int WIDTH = 5;
```

```
    int a[WIDTH][WIDTH];
```

```
    int b[WIDTH][WIDTH];
```

```
    int c[WIDTH][WIDTH] = { 0 };
```

```
    // make a, b matrices
```

```
    for (int y = 0; y < WIDTH; ++y) {
```

```
        for (int x = 0; x < WIDTH; ++x) {
```

```
            a[y][x] = y * 10 + x;
```

```
            b[y][x] = (y * 10 + x) * 100;
```

```
        }
```

```
    }
```

```
}
```

A 5x5 matrix A is shown. The rows are indexed by y (vertical axis, 0 to 4) and the columns by x (horizontal axis, 0 to 4). The matrix elements are:

| | $x=0$ | $x=1$ | $x=2$ | $x=3$ | $x=4$ |
|-------|-------|-------|-------|-------|-------|
| $y=0$ | 0 | 1 | 2 | 3 | 4 |
| $y=1$ | 10 | 11 | 12 | 13 | 14 |
| $y=2$ | 20 | 21 | 22 | 23 | 24 |
| $y=3$ | 30 | 31 | 32 | 33 | 34 |
| $y=4$ | 40 | 41 | 42 | 43 | 44 |

matadd-host.cu (cont'd)

```
// calculate
```

```
for (int y = 0; y < WIDTH; ++y) {  
    for (int x = 0; x < WIDTH; ++x) {  
        c[y][x] = a[y][x] + b[y][x];  
    }  
}
```

$$\begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 10 & 11 & 12 & 13 & 14 \\ 20 & 21 & 22 & 23 & 24 \\ 30 & 31 & 32 & 33 & 34 \\ 40 & 41 & 42 & 43 & 44 \end{bmatrix} + \begin{bmatrix} 0 & 100 & 200 & 300 & 400 \\ 1000 & 1100 & 1200 & 1300 & 1400 \\ 2000 & 2100 & 2200 & 2300 & 2400 \\ 3000 & 3100 & 3200 & 3300 & 3400 \\ 4000 & 4100 & 4200 & 4300 & 4400 \end{bmatrix}$$

matadd-host.cu (cont'd)

```
// print the result
```

```
for (int y = 0; y < WIDTH; ++y) {  
    for (int x = 0; x < WIDTH; ++x) {  
        printf("%5d", c[y][x]);  
    }  
    printf("\n");  
}
```

```
// done
```

```
return 0;
```

```
}
```

matadd-host.cu: a Complete Version

A Complete Version

```
#include <iostream>
```

```
int main(void) {
```

```
    // host-side data
```

```
    const int WIDTH = 5;
```

```
    int a[WIDTH][WIDTH];
```

```
    int b[WIDTH][WIDTH];
```

```
    int c[WIDTH][WIDTH] = { 0 };
```

```
    // make a, b matrices
```

```
    for (int y = 0; y < WIDTH; ++y) {
```

```
        for (int x = 0; x < WIDTH; ++x) {
```

```
            a[y][x] = y * 10 + x;
```

```
            b[y][x] = (y * 10 + x) * 100;
```

```
        }
```

```
    }
```

```
    // calculate
```

```
    for (int y = 0; y < WIDTH; ++y) {
```

```
        for (int x = 0; x < WIDTH; ++x)
```

```
            c[y][x] = a[y][x] + b[y][x];
```

```
    }
```

```
    // print the result
```

```
    for (int y = 0; y < WIDTH; ++y) {
```

```
        for (int x = 0; x < WIDTH; ++x)
```

```
            printf("%5d", c[y][x]);
```

```
        printf("\n");
```

```
    }
```

```
    // done
```

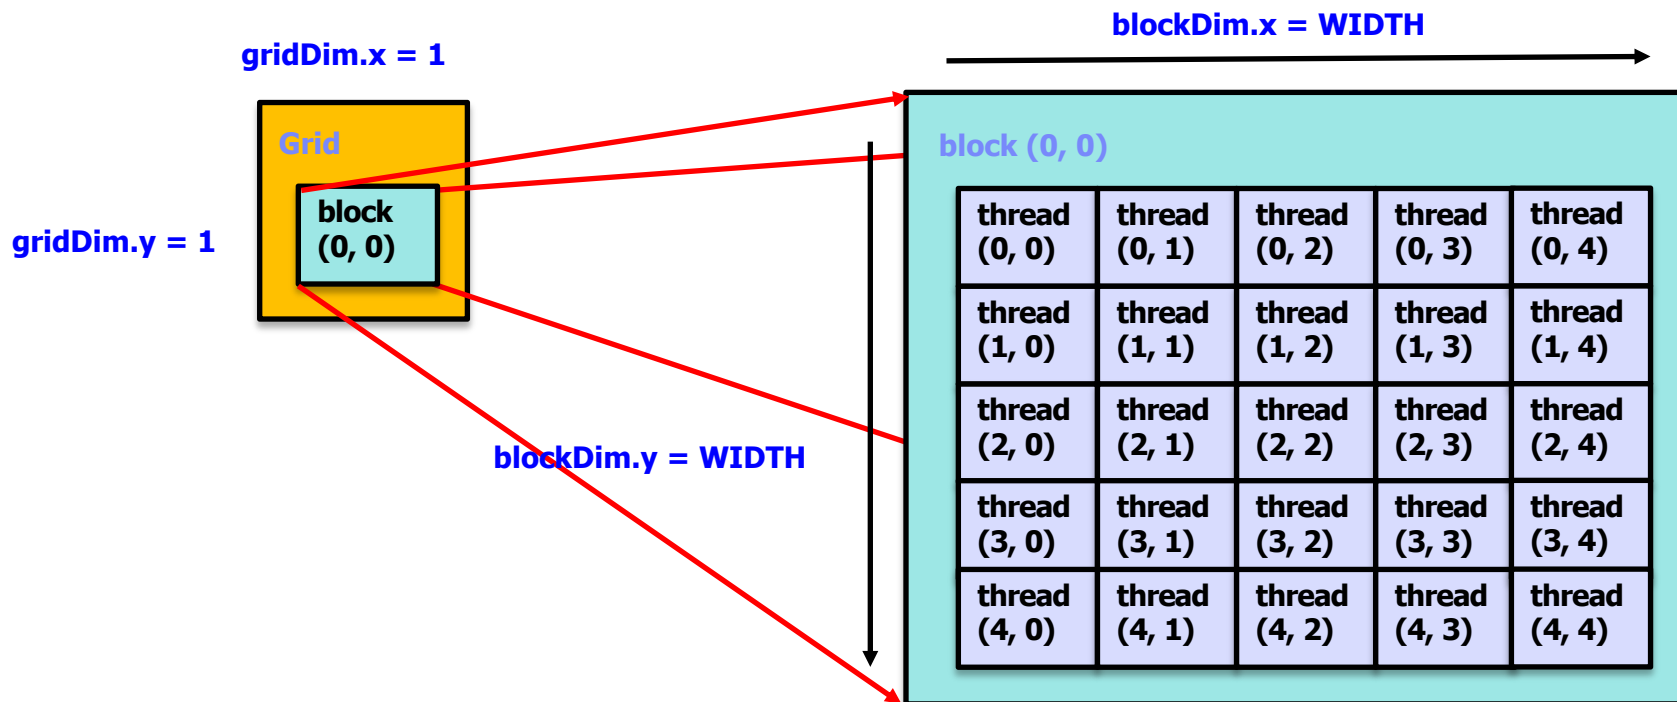
```
    return 0;
```

```
}
```

```
$ nvcc matadd-host.cu -o matadd-host
```

```
$ ./matadd-host
```


- Let's use GPU!
- Thread Organization
 - Matrix \rightarrow 2D
 - Small size matrix \rightarrow a single block



matadd-dev.cu (cont'd)

■ CUDA kernel

```
#include <iostream>

// kernel program for the device (GPU): compiled by NVCC
__global__ void addKernel(int* c, const int* a, const int* b) {
    int x = threadIdx.x;
    int y = threadIdx.y;
    int i = y * (blockDim.x) + x;    // [y][x] = y * WIDTH + x;
    c[i] = a[i] + b[i];
}
```

matadd-dev.cu (cont'd)

■ Host code

```
int main(void) {  
    // host-side data  
    const int WIDTH = 5;  
    int a[WIDTH][WIDTH];  
    int b[WIDTH][WIDTH];  
    int c[WIDTH][WIDTH] = { 0 };  
    // make a, b matrices  
    for (int y = 0; y < WIDTH; ++y) {  
        for (int x = 0; x < WIDTH; ++x) {  
            a[y][x] = y * 10 + x;  
            b[y][x] = (y * 10 + x) * 100;  
        }  
    }  
}
```

matadd-dev.cu (cont'd)

▪ Host code (cont'd)

```
// device-side data
```

```
int *dev_a = 0;
```

```
int *dev_b = 0;
```

```
int *dev_c = 0;
```

```
// allocate device memory
```

```
cudaMalloc((void**)&dev_a, WIDTH * WIDTH * sizeof(int));
```

```
cudaMalloc((void**)&dev_b, WIDTH * WIDTH * sizeof(int));
```

```
cudaMalloc((void**)&dev_c, WIDTH * WIDTH * sizeof(int));
```

```
// copy from host to device
```

```
cudaMemcpy(dev_a, a, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice);
```

```
cudaMemcpy(dev_b, b, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice);
```

matadd-dev.cu (cont'd)

▪ Host code (cont'd)

```
dim3 dimGrid(1, 1, 1);  
dim3 dimBlock(WIDTH, WIDTH, 1); // x, y, z  
addKernel<<<dimGrid, dimBlock>>>(dev_c, dev_a, dev_b);
```

```
// copy from device to host
```

```
cudaMemcpy(c, dev_c, WIDTH * WIDTH * sizeof(int),  
          cudaMemcpyDeviceToHost);
```

```
// free device memory
```

```
cudaFree(dev_c);  
cudaFree(dev_a);  
cudaFree(dev_b);
```

matadd-dev.cu (cont'd)

▪ Host code (cont'd)

```
// print the result
for (int y = 0; y < WIDTH; ++y) {
    for (int x = 0; x < WIDTH; ++x) {
        printf("%5d", c[y][x]);
    }
    printf("\n");
}

// done
return 0;
}
```

matadd-dev.cu: a Complete Version

```
#include <iostream>
#include "common.h"

// kernel program for the device (GPU): compiled by NVCC
__global__ void addKernel(int* c, const int* a, const int* b) {
    int x = threadIdx.x;
    int y = threadIdx.y;
    int i = y * (blockDim.x) + x;    // [y][x] = y * WIDTH + x;
    c[i] = a[i] + b[i];
}
```

```
// main program for the CPU: compiled by MS-VC++
int main(void) {
    // host-side data
    const int WIDTH = 5;
    int a[WIDTH][WIDTH];
    int b[WIDTH][WIDTH];
    int c[WIDTH][WIDTH] = { 0 };
    // make a, b matrices
    for (int y = 0; y < WIDTH; ++y) {
        for (int x = 0; x < WIDTH; ++x) {
            a[y][x] = y * 10 + x;
            b[y][x] = (y * 10 + x) * 100;
        }
    }
    // device-side data
    int* dev_a = 0;
    int* dev_b = 0;
    int* dev_c = 0;
```

```
// allocate device memory
    CUDA_CHECK( cudaMalloc((void**)&dev_a, WIDTH * WIDTH * sizeof(int)) );
    CUDA_CHECK( cudaMalloc((void**)&dev_b, WIDTH * WIDTH * sizeof(int)) );
    CUDA_CHECK( cudaMalloc((void**)&dev_c, WIDTH * WIDTH * sizeof(int)) );

    // copy from host to device
    CUDA_CHECK( cudaMemcpy(dev_a, a, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice) );
    CUDA_CHECK( cudaMemcpy(dev_b, b, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice) );

    // launch a kernel on the GPU with one thread for each element.
    dim3 dimBlock(WIDTH, WIDTH, 1); // x, y, z
    addKernel <<< 1, dimBlock>>>(dev_c, dev_a, dev_b);    // dev_c = dev_a + dev_b;

    CUDA_CHECK( cudaPeekAtLastError() );

    // copy from device to host
    CUDA_CHECK( cudaMemcpy(c, dev_c, WIDTH * WIDTH * sizeof(int), cudaMemcpyDeviceToHost) );

    // free device memory
    CUDA_CHECK( cudaFree(dev_c) );
    CUDA_CHECK( cudaFree(dev_a) );
    CUDA_CHECK( cudaFree(dev_b) );

    // print the result
    for (int y = 0; y < WIDTH; ++y) {
        for (int x = 0; x < WIDTH; ++x) {
            printf("%5d", c[y][x]);
        }
        printf("\n");
    }

    // done
    return 0;
}
```

matadd-dev.cu: a Complete Version

common.h

```
#ifdef DEBUG // debug mode
#define CUDA_CHECK(x)    do {\
    (x); \
    cudaError_t e = cudaGetLastError(); \
    if (cudaSuccess != e) { \
        printf("cuda failure %s at %s:%d\n", \
            cudaGetErrorString(e), \
            __FILE__, __LINE__); \
        exit(1); \
    } \
} while (0)
#else
#define CUDA_CHECK(x)    (x)           // release mode
#endif
```


Next?

- What is Thread?
- CUDA Thread Organization
- Extended VecAdd.cu
- Maxtrix Addition Kernel
- Mapping Threads to Multidimensional Data
- Synchronization and Transparent Scalability
- Resource Assignment
- Querying Device Properties
- Thread Scheduling and Latency Tolerance