

# **CUDA Memory Model 1**

Prof. Seokin Hong

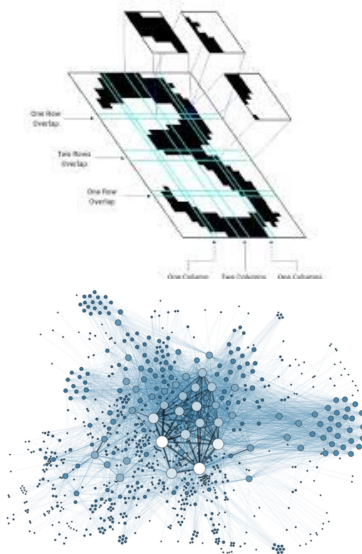
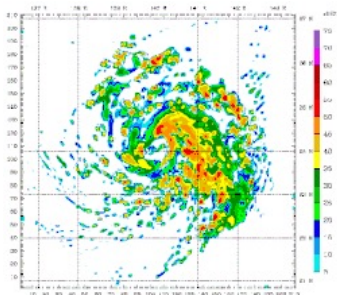
# Agenda

---

- **Matrix Multiplication**
  - **Basic Version**
  - **Tiled Version**
- Review: Memory Hierarchy
- Importance of Memory Access Efficiency
- GPU Memory Hierarchy
- Improving Tiled Matrix Multiplication
- Impact of Memory on Parallelism

# Matrix Multiplication is Fundamental in HPC

- Many algorithmic problems can be solved by means of matrix computation
  - Scientific computing
  - Pattern Recognition
  - Graph Analysis
- LINPACK relies heavily on the matrix multiplication
  - LINPACK is used for measuring the performance of Supercomputer



Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	<b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,299,072	415,530.0	513,854.7	28,335
2	<b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148,600.0	200,794.9	10,096
3	<b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	125,712.0	7,438
4	<b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371

# Basic Matrix Multiplication

# Matrix Multiplication

---

$$\mathbf{C} = \mathbf{A} \cdot \mathbf{B} = \begin{bmatrix} a_{00} & a_{01} & a_{02} & a_{03} & a_{04} \\ a_{10} & a_{11} & a_{12} & a_{13} & a_{14} \\ a_{20} & a_{21} & a_{22} & a_{23} & a_{24} \\ a_{30} & a_{31} & a_{32} & a_{33} & a_{34} \\ a_{40} & a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \cdot \begin{bmatrix} b_{00} & b_{01} & b_{02} & b_{03} & b_{04} \\ b_{10} & b_{11} & b_{12} & b_{13} & b_{14} \\ b_{20} & b_{21} & b_{22} & b_{23} & b_{24} \\ b_{30} & b_{31} & b_{32} & b_{33} & b_{34} \\ b_{40} & b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix}$$

# Matrix Multiplication

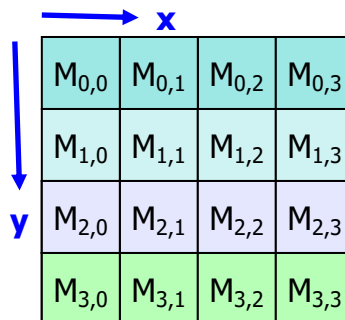
- $C_{ij}$  = dot product of  $A_{i\_}$  and  $B_{\_j}$

$$\mathbf{C} = \begin{bmatrix} c_{00} & c_{01} & c_{02} & c_{03} & c_{04} \\ c_{10} & c_{11} & c_{12} & c_{13} & c_{14} \\ c_{20} & c_{21} & c_{22} & c_{23} & c_{24} \\ c_{30} & \boxed{c_{31}} & c_{32} & c_{33} & c_{34} \\ c_{40} & c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} & a_{03} & a_{04} \\ a_{10} & a_{11} & a_{12} & a_{13} & a_{14} \\ a_{20} & a_{21} & a_{22} & a_{23} & a_{24} \\ \boxed{a_{30} \quad a_{31} \quad a_{32} \quad a_{33} \quad a_{34}} \\ a_{40} & a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \cdot \begin{bmatrix} b_{00} & \boxed{b_{01}} & b_{02} & b_{03} & b_{04} \\ b_{10} & \boxed{b_{11}} & b_{12} & b_{13} & b_{14} \\ b_{20} & \boxed{b_{21}} & b_{22} & b_{23} & b_{24} \\ b_{30} & \boxed{b_{31}} & b_{32} & b_{33} & b_{34} \\ b_{40} & \boxed{b_{41}} & b_{42} & b_{43} & b_{44} \end{bmatrix}$$

$$c_{31} = \begin{bmatrix} a_{30} & a_{31} & a_{32} & a_{33} & a_{34} \end{bmatrix} \cdot \begin{bmatrix} b_{01} \\ b_{11} \\ b_{21} \\ b_{31} \\ b_{41} \end{bmatrix}$$

# Row-major Matrix Layout in C/C++

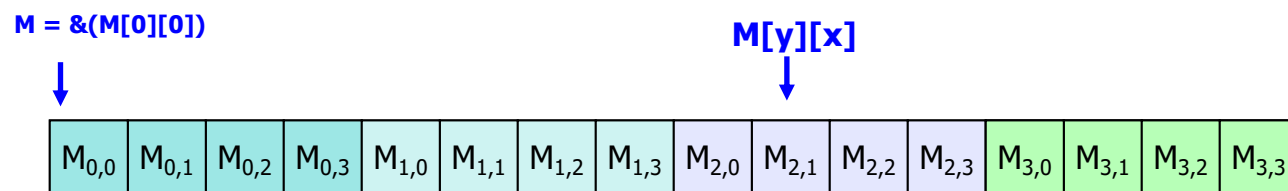
- logical layout:



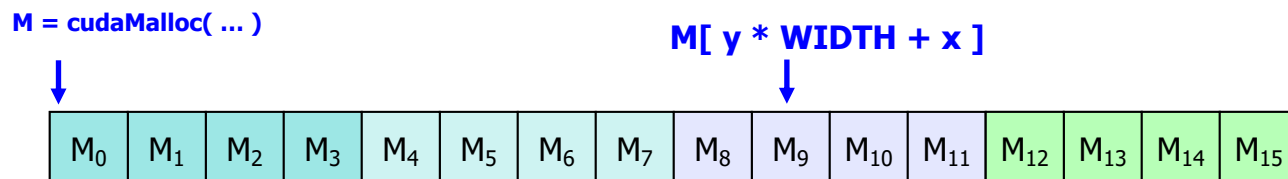
A 4x4 grid representing a matrix. The columns are labeled with an 'x' and an arrow pointing right, and the rows are labeled with a 'y' and an arrow pointing down. The elements are arranged in rows: the first row contains M<sub>0,0</sub>, M<sub>0,1</sub>, M<sub>0,2</sub>, M<sub>0,3</sub>; the second row contains M<sub>1,0</sub>, M<sub>1,1</sub>, M<sub>1,2</sub>, M<sub>1,3</sub>; the third row contains M<sub>2,0</sub>, M<sub>2,1</sub>, M<sub>2,2</sub>, M<sub>2,3</sub>; and the fourth row contains M<sub>3,0</sub>, M<sub>3,1</sub>, M<sub>3,2</sub>, M<sub>3,3</sub>. The colors of the cells are: first row (light blue), second row (light blue), third row (light purple), and fourth row (light green).

M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>

- physical layout: 1D array



- re-interpret:



# Matrix Multiplication

---

- $C_{ij}$  = dot product of  $A_{i\_}$  and  $B_{\_j}$

$$c_{31} = \begin{bmatrix} a_{30} & a_{31} & a_{32} & a_{33} & a_{34} \end{bmatrix} \cdot \begin{bmatrix} b_{01} \\ b_{11} \\ b_{21} \\ b_{31} \\ b_{41} \end{bmatrix}$$

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

```
int sum = 0;

for (int k = 0; k < WIDTH; ++k) {
    sum += a[i][k] * b[k][j];
}

c[i][j] = sum;
```



# matmul-host.cu (1/2)

- calculate matrix multiplication on CPU

```
//prepare data
```

```
const int WIDTH = 5;  
int a[WIDTH][WIDTH];  
int b[WIDTH][WIDTH];  
int c[WIDTH][WIDTH] = { 0 };
```

```
//make matrices A, B
```

```
for (int y = 0; y < WIDTH; ++y) {  
    for (int x = 0; x < WIDTH; ++x) {  
        a[y][x] = y + x;  
        b[y][x] = y + x;  
    }  
}
```

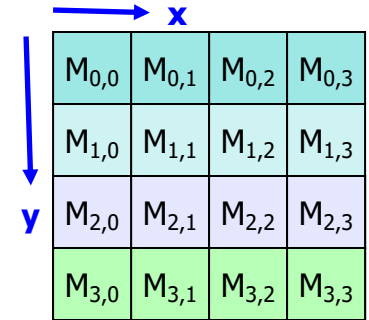
$$\begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 & 5 \\ 2 & 3 & 4 & 5 & 6 \\ 3 & 4 & 5 & 6 & 7 \\ 4 & 5 & 6 & 7 & 8 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 & 5 \\ 2 & 3 & 4 & 5 & 6 \\ 3 & 4 & 5 & 6 & 7 \\ 4 & 5 & 6 & 7 & 8 \end{bmatrix}$$

# matmul-host.cu (2/2)

- calculate matrix multiplication on CPU

//calculation code

```
for (int y = 0; y < WIDTH; ++y) {  
    for (int x = 0; x < WIDTH; ++x) {  
        int sum = 0;  
        for (int k = 0; k < WIDTH; ++k) {  
            sum += a[y][k] * b[k][x];  
        }  
        c[y][x] = sum;  
    }  
}
```



A 4x4 matrix diagram with rows and columns indexed from 0 to 3. The columns are labeled with 'x' and the rows with 'y'. The matrix elements are labeled M<sub>y,x</sub>. The matrix is divided into four colored quadrants: top-left (cyan), top-right (light blue), bottom-left (light purple), and bottom-right (light green).

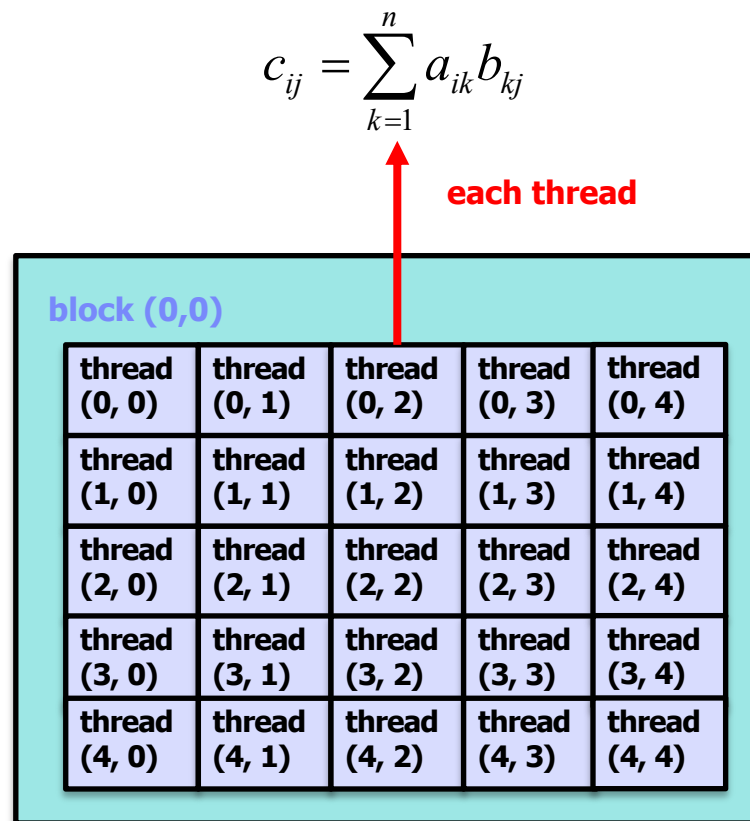
	x=0	x=1	x=2	x=3
y=0	M <sub>0,0</sub>	M <sub>0,1</sub>	M <sub>0,2</sub>	M <sub>0,3</sub>
y=1	M <sub>1,0</sub>	M <sub>1,1</sub>	M <sub>1,2</sub>	M <sub>1,3</sub>
y=2	M <sub>2,0</sub>	M <sub>2,1</sub>	M <sub>2,2</sub>	M <sub>2,3</sub>
y=3	M <sub>3,0</sub>	M <sub>3,1</sub>	M <sub>3,2</sub>	M <sub>3,3</sub>

$$\begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 & 5 \\ 2 & 3 & 4 & 5 & 6 \\ 3 & 4 & 5 & 6 & 7 \\ 4 & 5 & 6 & 7 & 8 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 & 5 \\ 2 & 3 & 4 & 5 & 6 \\ 3 & 4 & 5 & 6 & 7 \\ 4 & 5 & 6 & 7 & 8 \end{bmatrix}$$

# CUDA Matrix Multiplication

- $C = A * B$ 
  - Size: WIDTH x WIDTH
- Thread organization (layout)
  - One thread handles one element of C

$$C = \begin{bmatrix} c_{00} & c_{01} & c_{02} & c_{03} & c_{04} \\ c_{10} & c_{11} & c_{12} & c_{13} & c_{14} \\ c_{20} & c_{21} & c_{22} & c_{23} & c_{24} \\ c_{30} & c_{31} & c_{32} & c_{33} & c_{34} \\ c_{40} & c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$



# CUDA Matrix Multiplication (1/2)

## ■ $C = A * B$

- Size: WIDTH x WIDTH

$$C = \begin{bmatrix} c_{00} & c_{01} & c_{02} & c_{03} & c_{04} \\ c_{10} & c_{11} & c_{12} & c_{13} & c_{14} \\ c_{20} & c_{21} & c_{22} & c_{23} & c_{24} \\ c_{30} & \boxed{c_{31}} & c_{32} & c_{33} & c_{34} \\ c_{40} & c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} & a_{03} & a_{04} \\ a_{10} & a_{11} & a_{12} & a_{13} & a_{14} \\ a_{20} & a_{21} & a_{22} & a_{23} & a_{24} \\ \boxed{a_{30}} & \boxed{a_{31}} & \boxed{a_{32}} & \boxed{a_{33}} & \boxed{a_{34}} \\ a_{40} & a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \cdot \begin{bmatrix} b_{00} & \boxed{b_{01}} & b_{02} & b_{03} & b_{04} \\ b_{10} & b_{11} & b_{12} & b_{13} & b_{14} \\ b_{20} & b_{21} & b_{22} & b_{23} & b_{24} \\ b_{30} & b_{31} & b_{32} & b_{33} & b_{34} \\ \boxed{b_{40}} & \boxed{b_{41}} & b_{42} & b_{43} & b_{44} \end{bmatrix}$$

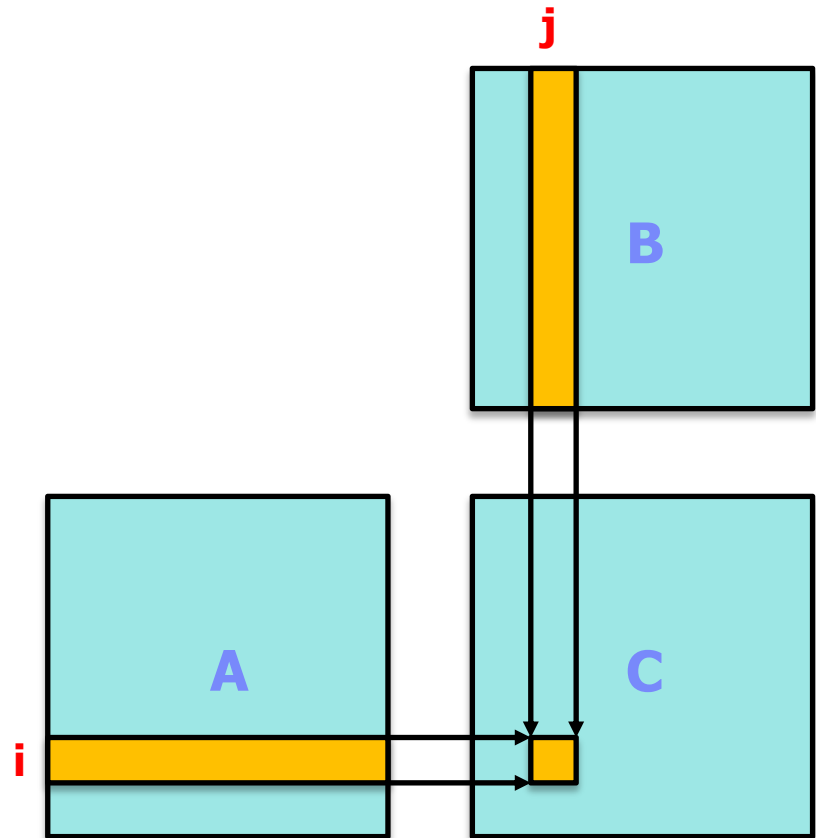
## ■ Memory usage view:

- read from A and B
- to calculate:

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

- using:

- $a[i * \text{WIDTH} + k]$
- $b[k * \text{WIDTH} + j];$



# CUDA Matrix Multiplication (2/2)

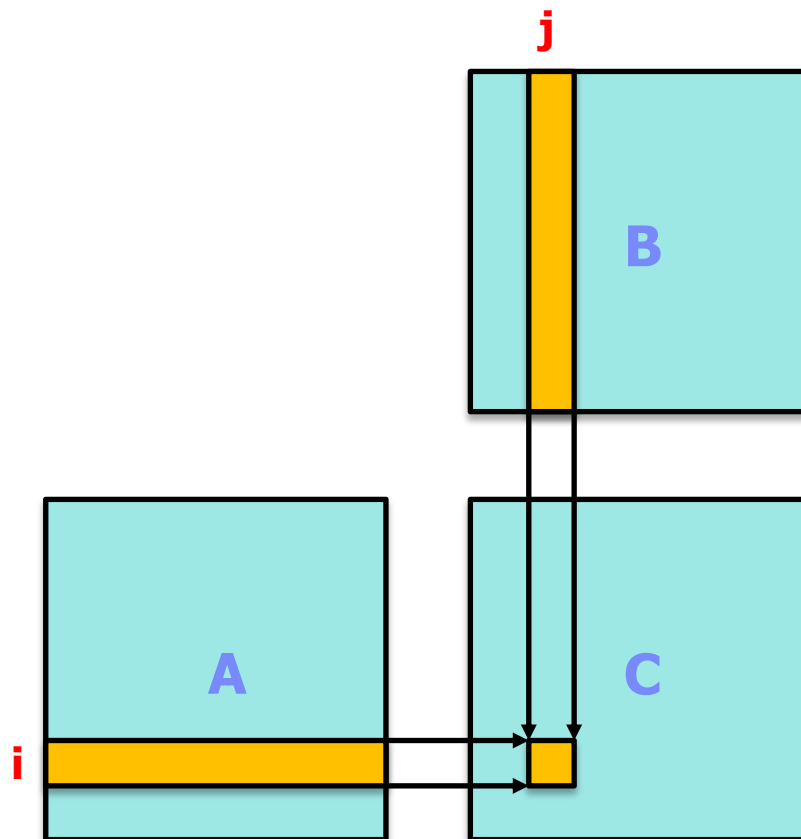
- **One Block of threads compute matrix C**

- Each thread computes one element of C

- **Each thread**

- loads a row of matrix A
- loads a column of matrix B
- Perform one multiply and addition for each pair of A and B elements

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$



# matmul-dev.cu (1/3) : Kernel code

---

- use  $\text{WIDTH} * \text{WIDTH}$  threads

```
__global__ void mulKernel(int* c, const int* a, const int* b, const int WIDTH) {  
    int x = threadIdx.x;  
    int y = threadIdx.y;  
  
    int i = y * WIDTH + x;                                // [y][x] = y * WIDTH + x;  
  
    int sum = 0;  
    for (int k = 0; k < WIDTH; ++k) {  
        sum += a[y * WIDTH + k] * b[k * WIDTH + x];  
    }  
    c[i] = sum;  
}
```

# matmul-dev.cu (2/3) : Host Code

```
int main(void) {  
    // host-side data  
    const int WIDTH = 5;  
    int a[WIDTH][WIDTH];  
    int b[WIDTH][WIDTH];  
    int c[WIDTH][WIDTH] = { 0 };  
    // make a, b matrices  
    for (int y = 0; y < WIDTH; ++y) {  
        for (int x = 0; x < WIDTH; ++x) {  
            a[y][x] = y * 10 + x;  
            b[y][x] = (y * 10 + x) * 100;  
        }  
    }  
  
    // device-side data  
    int *dev_a = 0;  
    int *dev_b = 0;  
    int *dev_c = 0;  
    //allocate device memory  
    cudaMalloc((void**)&dev_a, WIDTH * WIDTH * sizeof(int));  
    cudaMalloc((void**)&dev_b, WIDTH * WIDTH * sizeof(int));  
    cudaMalloc((void**)&dev_c, WIDTH * WIDTH * sizeof(int));  
  
    // copy from host to device  
    cudaMemcpy(dev_a, a, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice);  
    cudaMemcpy(dev_b, b, WIDTH * WIDTH * sizeof(int), cudaMemcpyHostToDevice);  
}
```

# matmul-dev.cu (3/3) : Host Code

// launch a kernel on the GPU

```
dim3 dimBlock(WIDTH, WIDTH, 1); // x, y, z
mulKernel<<<1, dimBlock>>>(dev_c, dev_a, dev_b, WIDTH);
CUDA_CHECK( cudaPeekAtLastError() );
```

// copy from device to host

```
cudaMemcpy(c, dev_c, WIDTH * WIDTH * sizeof(int),
           cudaMemcpyDeviceToHost);
```


// free device memory

```
cudaFree(dev_c);
cudaFree(dev_a);
cudaFree(dev_b);
```

// print the result

```
for (int y = 0; y < WIDTH; ++y) {
    for (int x = 0; x < WIDTH; ++x) {
        printf("%5d", c[y][x]);
    }
    printf("\n");
}
return 0;
```

```
}
```



0,0	0,1	0,2	0,3	0,4
1,0	1,1	1,2	1,3	1,4
2,0	2,1	2,2	2,3	2,4
3,0	3,1	3,2	3,3	3,4
4,0	4,1	4,2	4,3	4,4



# Tiled Matrix Multiplication

# Thread Organization in the Simple Matrix Multiplication

- matrix  $\rightarrow$  2D layout
- small size matrix  $\rightarrow$  a single block !

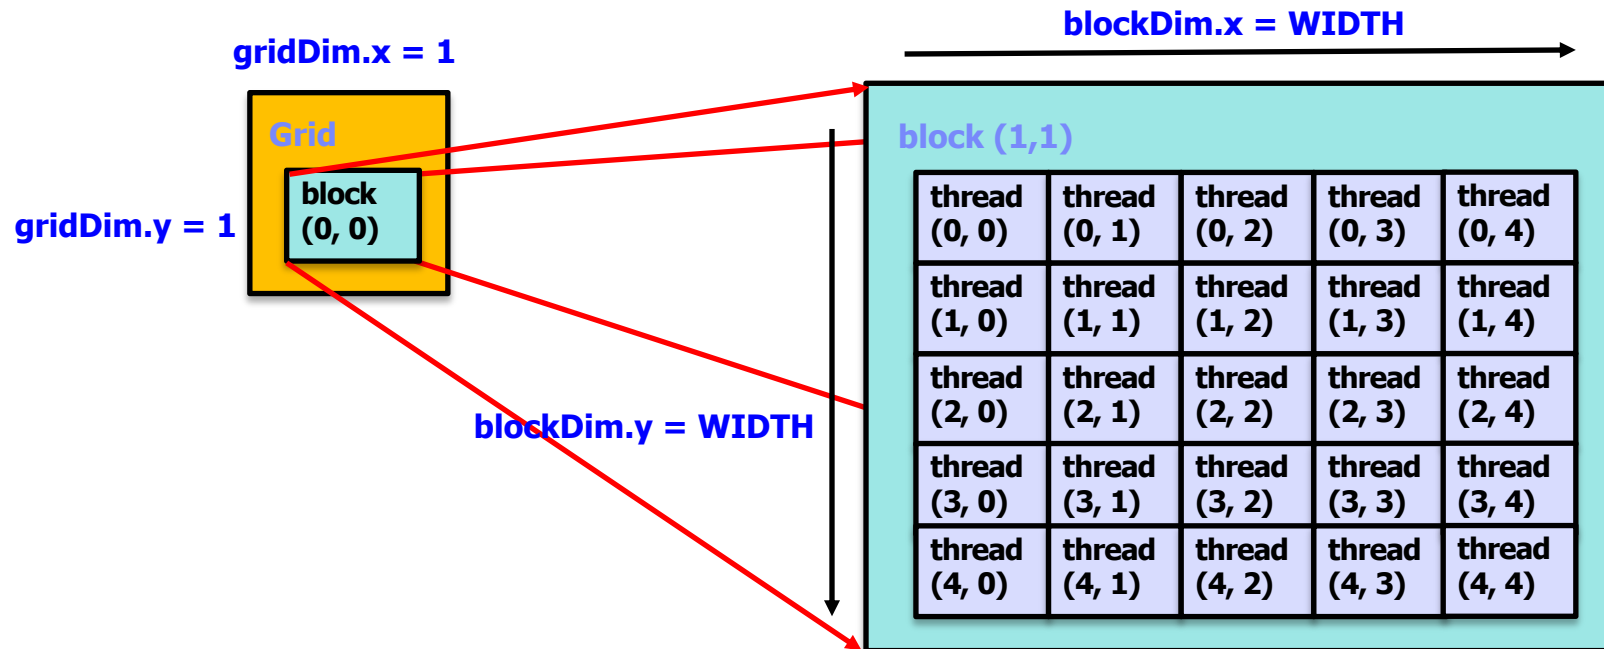


Image from <http://developer.amd.com/zones/OpenCLZone/courses/pages/Introductory-OpenCL-SAAHPC10.aspx>

# Any Problem?

---

- We used only one thread block...
- Each thread block can execute at most 1024 threads
  - some old architecture can execute only 256 or 512 threads
- So, maximum matrix size is...
  - $32 \times 32 = 1024$
  - with a single thread block...
- Many global memory accesses
  - Global memory is slow
  - We will discuss it later!
- solution?
  - use multiple thread blocks !

# Thread Organization (layout)

- matrix  $\rightarrow$  2D layout
- tiled approach : use multiple blocks

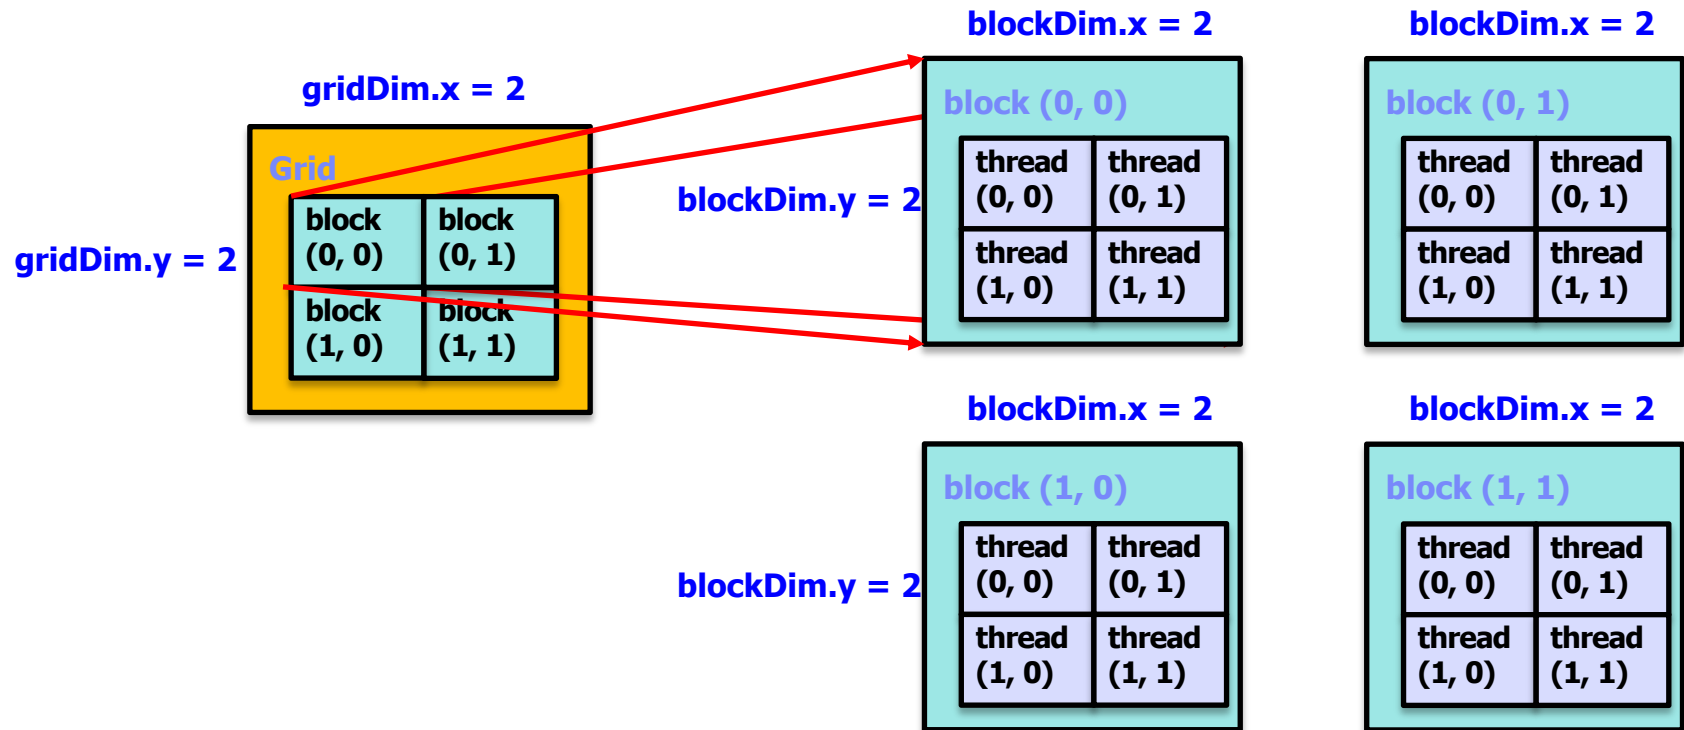


Image from <http://developer.amd.com/zones/OpenCLZone/courses/pages/Introductory-OpenCL-SAAHPC10.aspx>

# Simplified Thread Organization

- assumption: square matrices
- global size :  $WIDTH \times WIDTH$
- tile  $\rightarrow$  a block :  $TILE\_WIDTH \times TILE\_WIDTH$
- grid :  $\text{ceil}(WIDTH / TILE\_WIDTH) \times \text{ceil}(WIDTH / TILE\_WIDTH)$

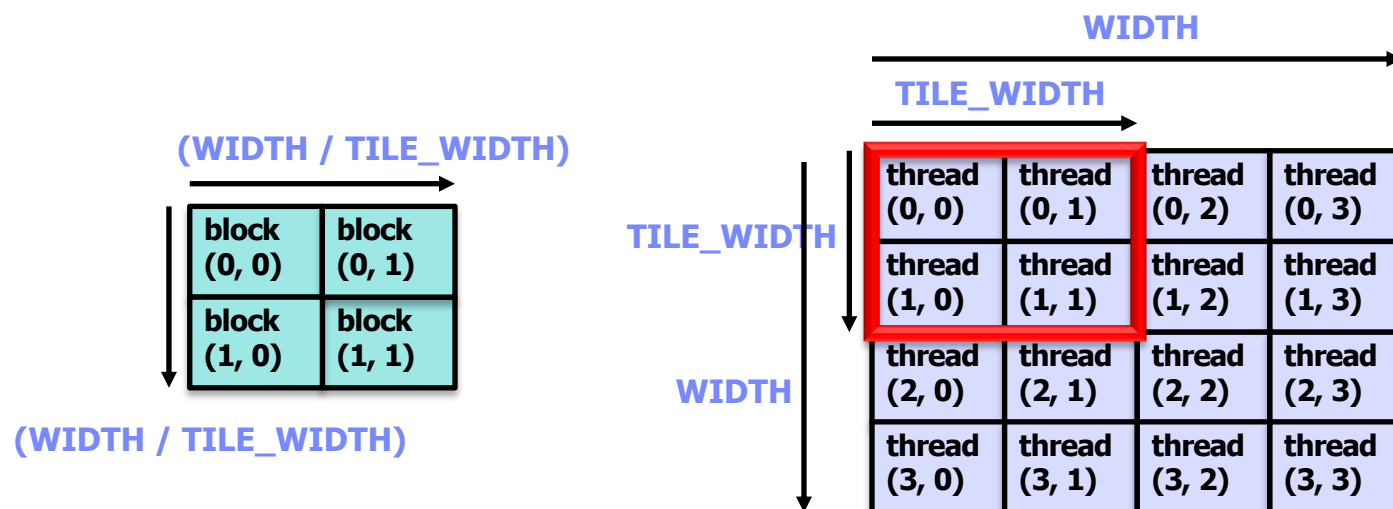


Image from <http://developer.amd.com/zones/OpenCLZone/courses/pages/Introductory-OpenCL-SAAHPC10.aspx>

# An Example of Thread Organization

- WIDTH = 8
- TILE\_WIDTH = 2
- grid dimension = 4 x 4

block (0,0)	block (0,1)	block (0,2)	block (0,3)
block (1,0)	block (1,1)	block (1,2)	block (1,3)
block (2,0)	block (2,1)	block (2,2)	block (2,3)
block (3,0)	block (3,1)	block (3,2)	block (3,3)

P <sub>0,0</sub>	P <sub>0,1</sub>	P <sub>0,2</sub>	P <sub>0,3</sub>	P <sub>0,4</sub>	P <sub>0,5</sub>	P <sub>0,6</sub>	P <sub>0,7</sub>
P <sub>1,0</sub>	P <sub>1,1</sub>	P <sub>1,2</sub>	P <sub>1,3</sub>	P <sub>1,4</sub>	P <sub>1,5</sub>	P <sub>1,6</sub>	P <sub>1,7</sub>
P <sub>2,0</sub>	P <sub>2,1</sub>	P <sub>2,2</sub>	P <sub>2,3</sub>	P <sub>2,4</sub>	P <sub>2,5</sub>	P <sub>2,6</sub>	P <sub>2,7</sub>
P <sub>3,0</sub>	P <sub>3,1</sub>	P <sub>3,2</sub>	P <sub>3,3</sub>	P <sub>3,4</sub>	P <sub>3,5</sub>	P <sub>3,6</sub>	P <sub>3,7</sub>
P <sub>4,0</sub>	P <sub>4,1</sub>	P <sub>4,2</sub>	P <sub>4,3</sub>	P <sub>4,4</sub>	P <sub>4,5</sub>	P <sub>4,6</sub>	P <sub>4,7</sub>
P <sub>5,0</sub>	P <sub>5,1</sub>	P <sub>5,2</sub>	P <sub>5,3</sub>	P <sub>5,4</sub>	P <sub>5,5</sub>	P <sub>5,6</sub>	P <sub>5,7</sub>
P <sub>6,0</sub>	P <sub>6,1</sub>	P <sub>6,2</sub>	P <sub>6,3</sub>	P <sub>6,4</sub>	P <sub>6,5</sub>	P <sub>6,6</sub>	P <sub>6,7</sub>
P <sub>7,0</sub>	P <sub>7,1</sub>	P <sub>7,2</sub>	P <sub>7,3</sub>	P <sub>7,4</sub>	P <sub>7,5</sub>	P <sub>7,6</sub>	P <sub>7,7</sub>

# An Example of Thread Organization

- WIDTH = 8
- TILE\_WIDTH = 4
- grid dimension = 2 x 2

block (0,0)	block (0,1)
block (1,0)	block (1,1)

P <sub>0,0</sub>	P <sub>0,1</sub>	P <sub>0,2</sub>	P <sub>0,3</sub>	P <sub>0,4</sub>	P <sub>0,5</sub>	P <sub>0,6</sub>	P <sub>0,7</sub>
P <sub>1,0</sub>	P <sub>1,1</sub>	P <sub>1,2</sub>	P <sub>1,3</sub>	P <sub>1,4</sub>	P <sub>1,5</sub>	P <sub>1,6</sub>	P <sub>1,7</sub>
P <sub>2,0</sub>	P <sub>2,1</sub>	P <sub>2,2</sub>	P <sub>2,3</sub>	P <sub>2,4</sub>	P <sub>2,5</sub>	P <sub>2,6</sub>	P <sub>2,7</sub>
P <sub>3,0</sub>	P <sub>3,1</sub>	P <sub>3,2</sub>	P <sub>3,3</sub>	P <sub>3,4</sub>	P <sub>3,5</sub>	P <sub>3,6</sub>	P <sub>3,7</sub>
P <sub>4,0</sub>	P <sub>4,1</sub>	P <sub>4,2</sub>	P <sub>4,3</sub>	P <sub>4,4</sub>	P <sub>4,5</sub>	P <sub>4,6</sub>	P <sub>4,7</sub>
P <sub>5,0</sub>	P <sub>5,1</sub>	P <sub>5,2</sub>	P <sub>5,3</sub>	P <sub>5,4</sub>	P <sub>5,5</sub>	P <sub>5,6</sub>	P <sub>5,7</sub>
P <sub>6,0</sub>	P <sub>6,1</sub>	P <sub>6,2</sub>	P <sub>6,3</sub>	P <sub>6,4</sub>	P <sub>6,5</sub>	P <sub>6,6</sub>	P <sub>6,7</sub>
P <sub>7,0</sub>	P <sub>7,1</sub>	P <sub>7,2</sub>	P <sub>7,3</sub>	P <sub>7,4</sub>	P <sub>7,5</sub>	P <sub>7,6</sub>	P <sub>7,7</sub>

# Kernel Launch

---

```
const int  WIDTH = 8;
```

```
const int  TILE_WIDTH = 4;
```

```
// Setup the execution configuration
```

```
dim3 dimGrid( ceil(WIDTH / TILE_WIDTH), ceil(WIDTH / TILE_WIDTH), 1 );
```

```
dim3 dimBlock( TILE_WIDTH, TILE_WIDTH, 1 );
```

```
// Launch the device computation threads!
```

```
MatrixMulKernel <<<dimGrid, dimBlock>>> (dev_c, dev_a, dev_b, WIDTH);
```



# Local Index vs Global Index

- For thread (1,2) of Block (0, 1)

- `blockIdx.y = 0`
- `blockIdx.x = 1`

- local index

- `threadIdx.y = 1`
- `threadIdx.x = 0`

- global index

- $y = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$

$$0 * 2 + 1 \rightarrow 1$$

- $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$

$$1 * 2 + 0 \rightarrow 2$$

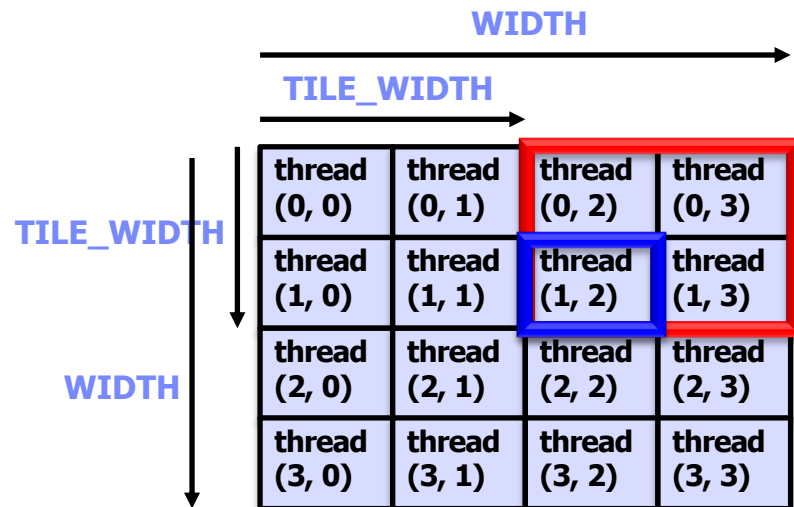
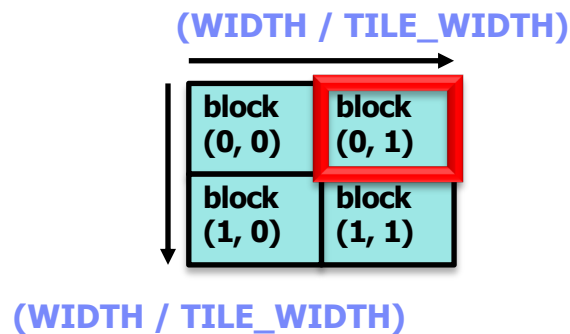
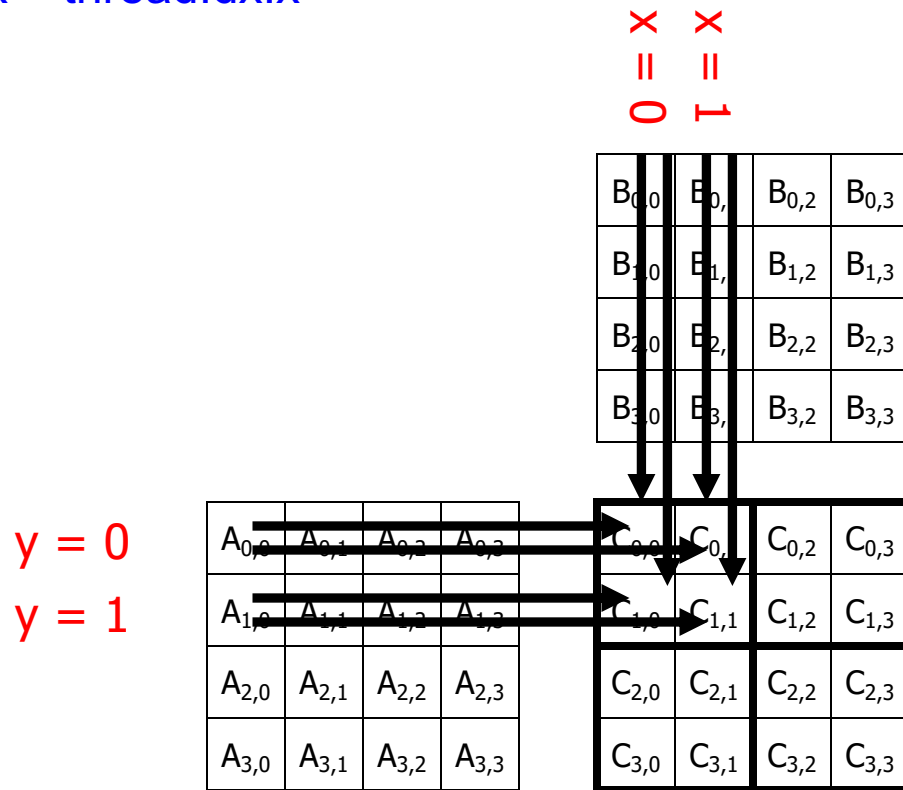


Image from <http://developer.amd.com/zones/OpenCLZone/courses/pages/Introductory-OpenCL-SAAHPC10.aspx>

# Indices for Block (0,0)

- general calculation:
  - $y = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
  - $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
- in the block (0,0)
  - $y = 0 * 2 + \text{threadIdx.y}$
  - $x = 0 * 2 + \text{threadIdx.x}$

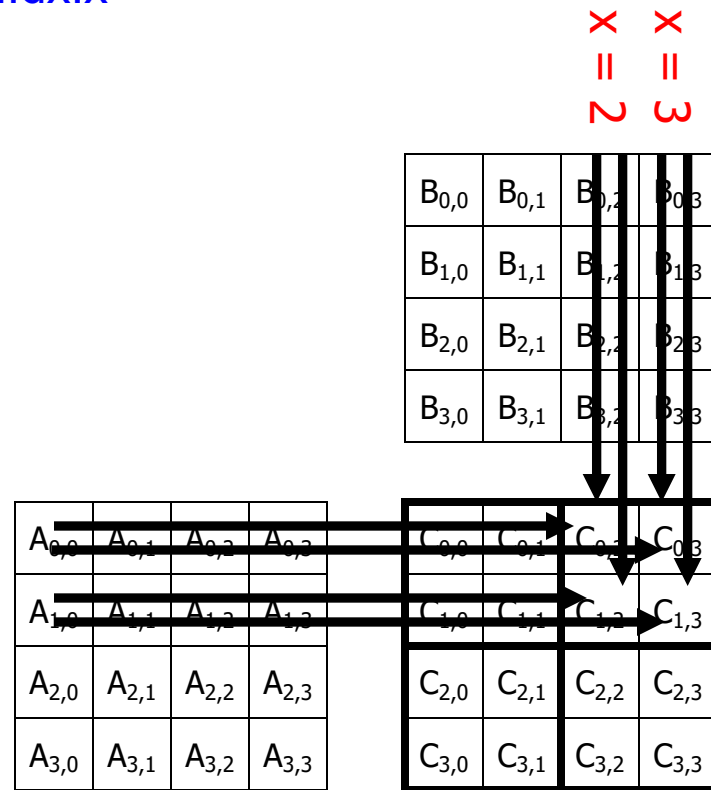


# Indices for Block (0,1)

- general calculation:
  - $y = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
  - $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
- in the block (0,1)
  - $y = 0 * 2 + \text{threadIdx.y}$
  - $x = 1 * 2 + \text{threadIdx.x}$

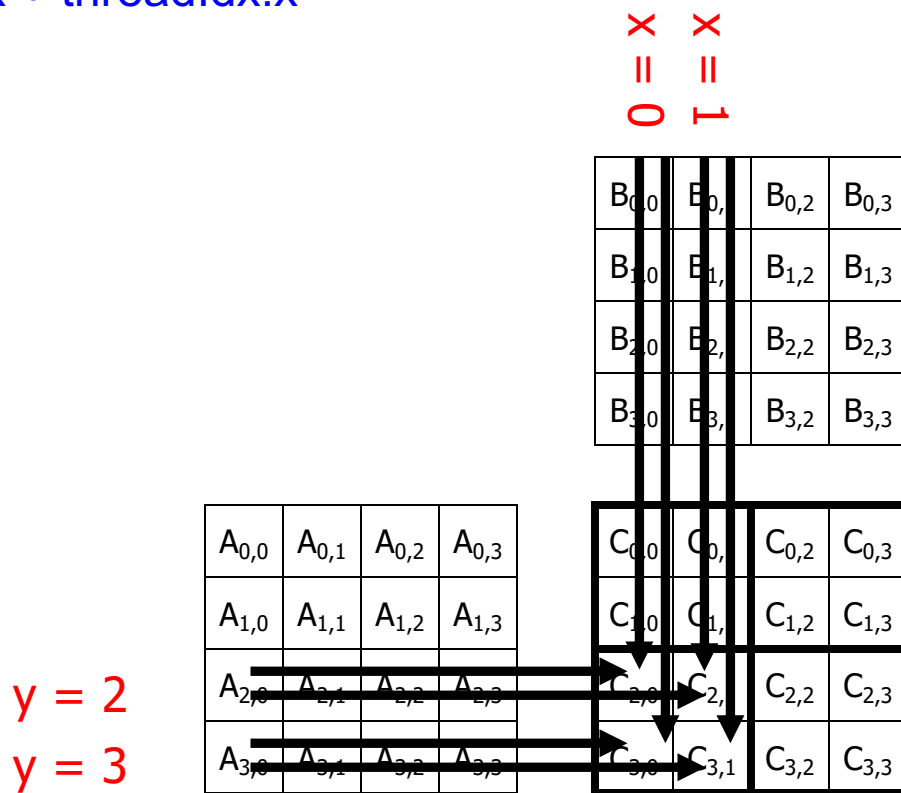
$y = 0$

$y = 1$



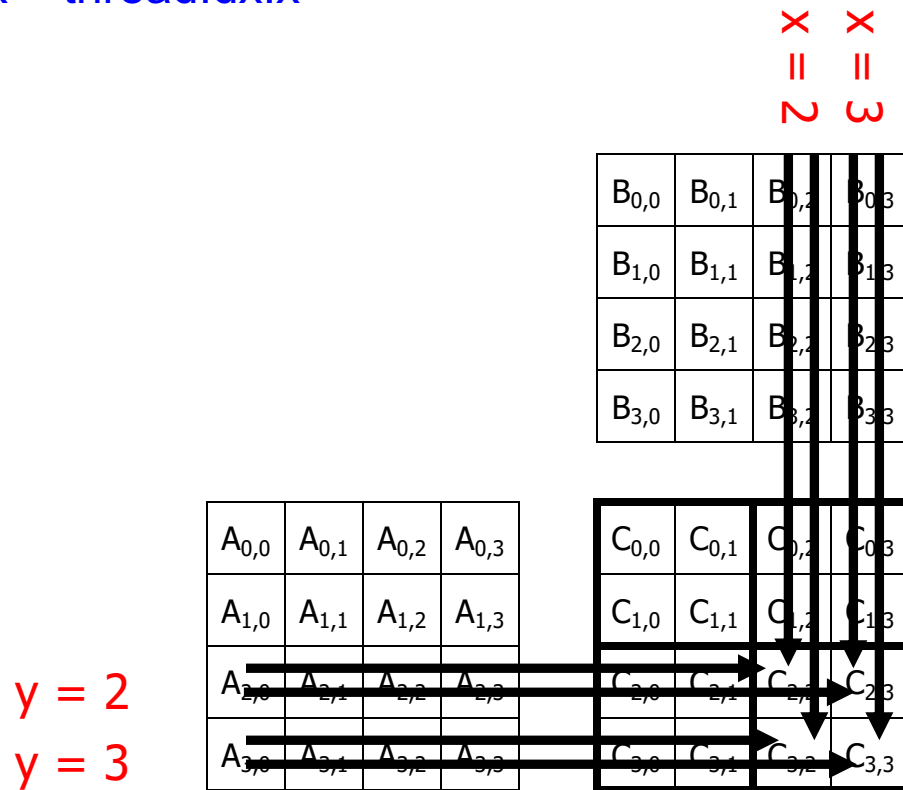
# Indices for Block (1,0)

- general calculation:
  - $y = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
  - $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
- in the block (1,0)
  - $y = 1 * 2 + \text{threadIdx.y}$
  - $x = 0 * 2 + \text{threadIdx.x}$



# Indices for Block (1,1)

- general calculation:
  - $y = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y}$
  - $x = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$
- in the block (1,1)
  - $y = 1 * 2 + \text{threadIdx.y}$
  - $x = 1 * 2 + \text{threadIdx.x}$



# Matrix Multiplication Kernel

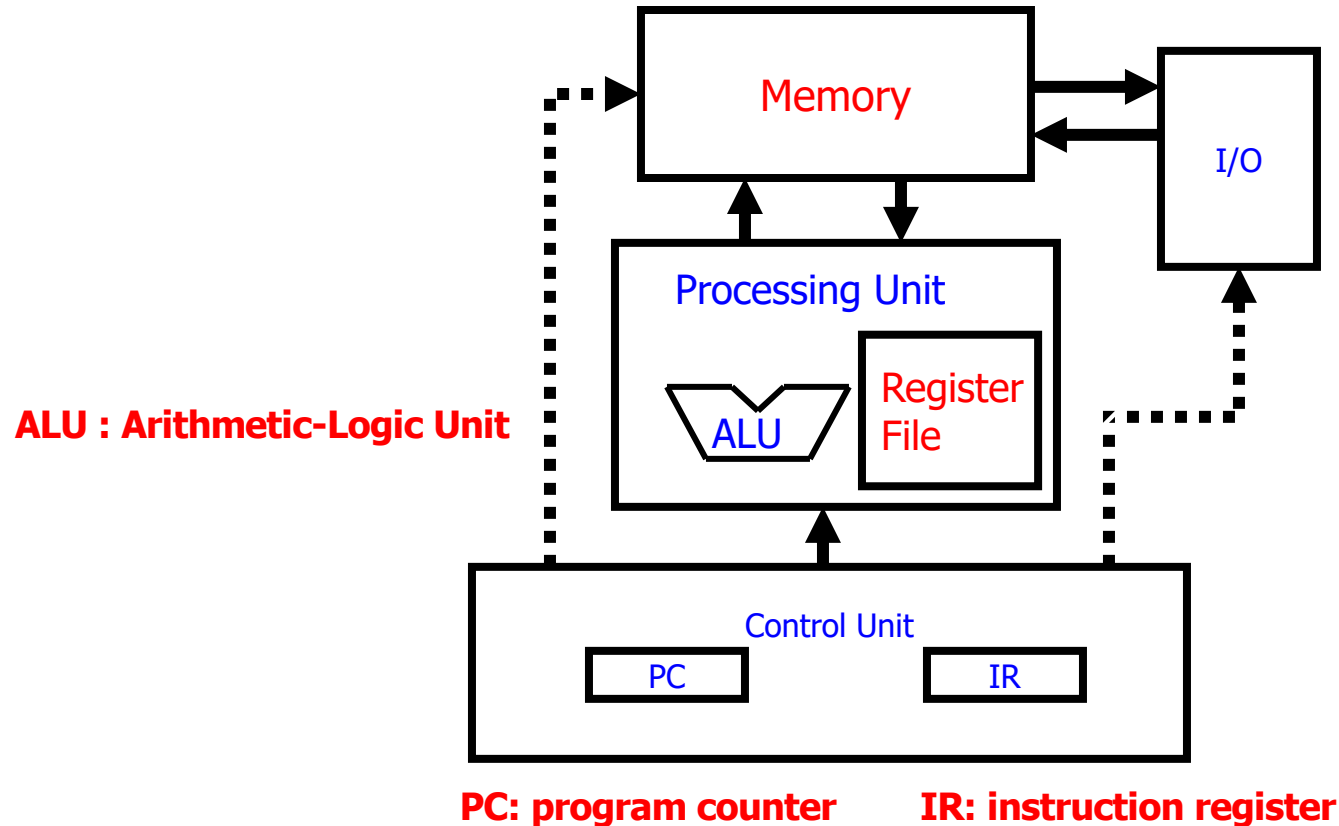
---

```
__global__ void matmul(float* c, const float* a, const float* b, const int width) {  
    int y = blockIdx.y * blockDim.y + threadIdx.y;  
    int x = blockIdx.x * blockDim.x + threadIdx.x;  
    float sum = 0.0F;  
    for (register int k = 0; k < width; ++k) {  
        float lhs = a[y * width + k];  
        float rhs = b[k * width + x];  
        sum += lhs * rhs;  
    }  
    c[y * width + x] = sum;  
}
```

# Memory Hierarchy

# Why Memory?: Von-Neumann Model

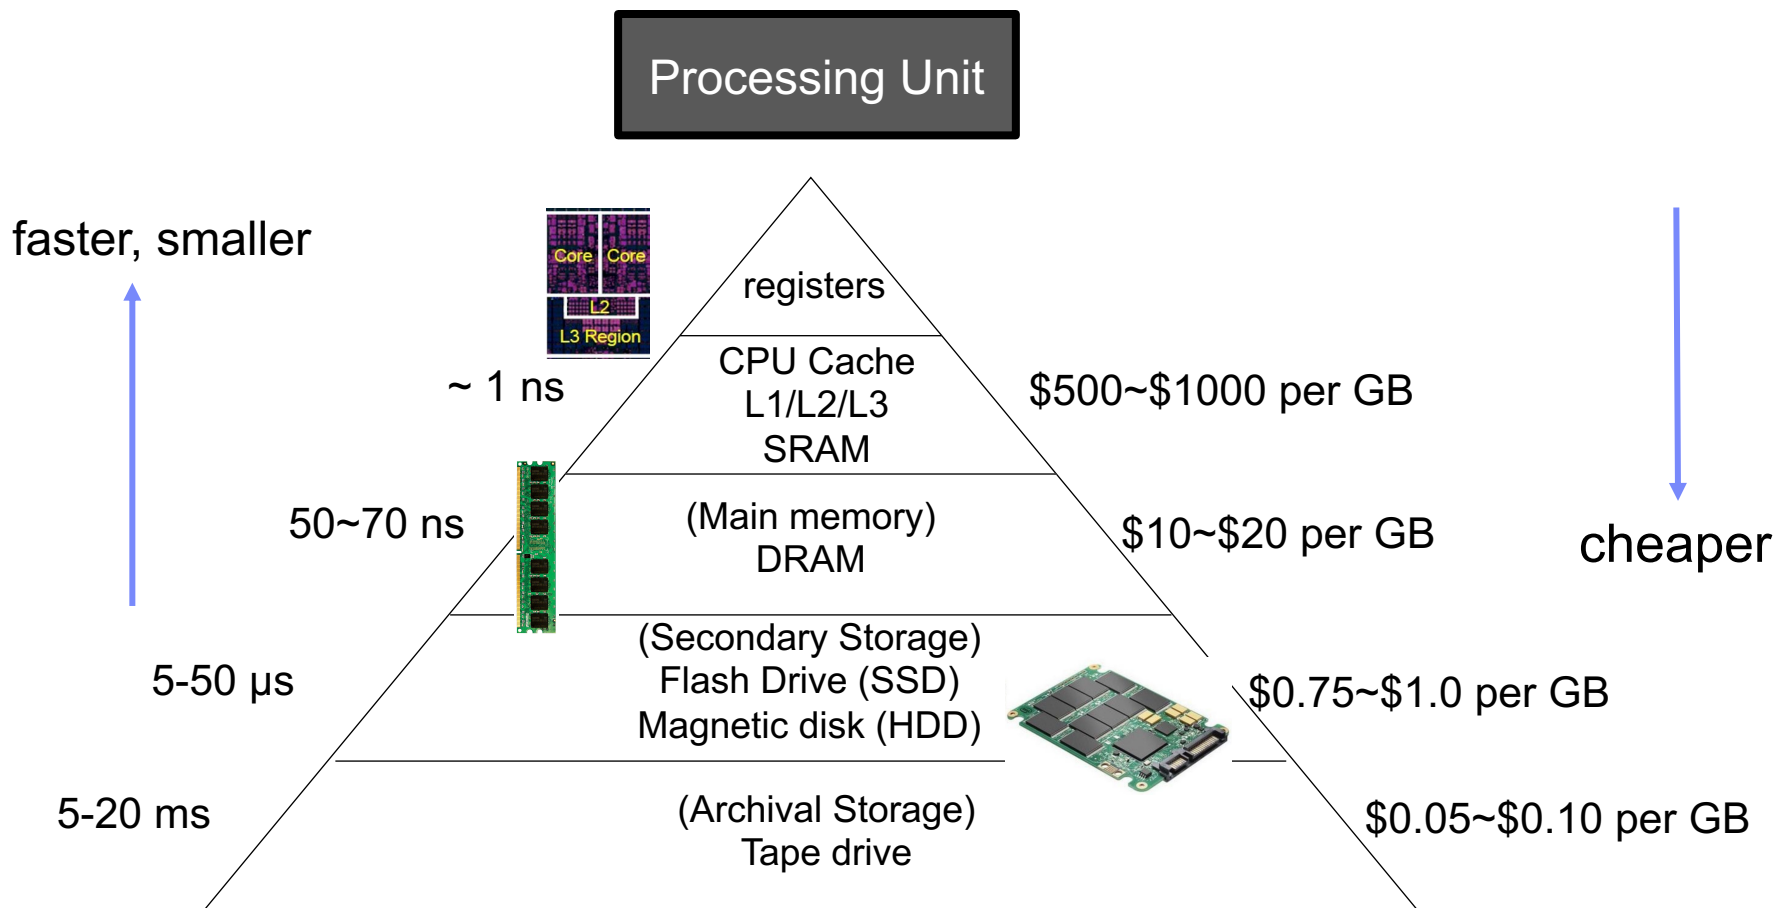
---



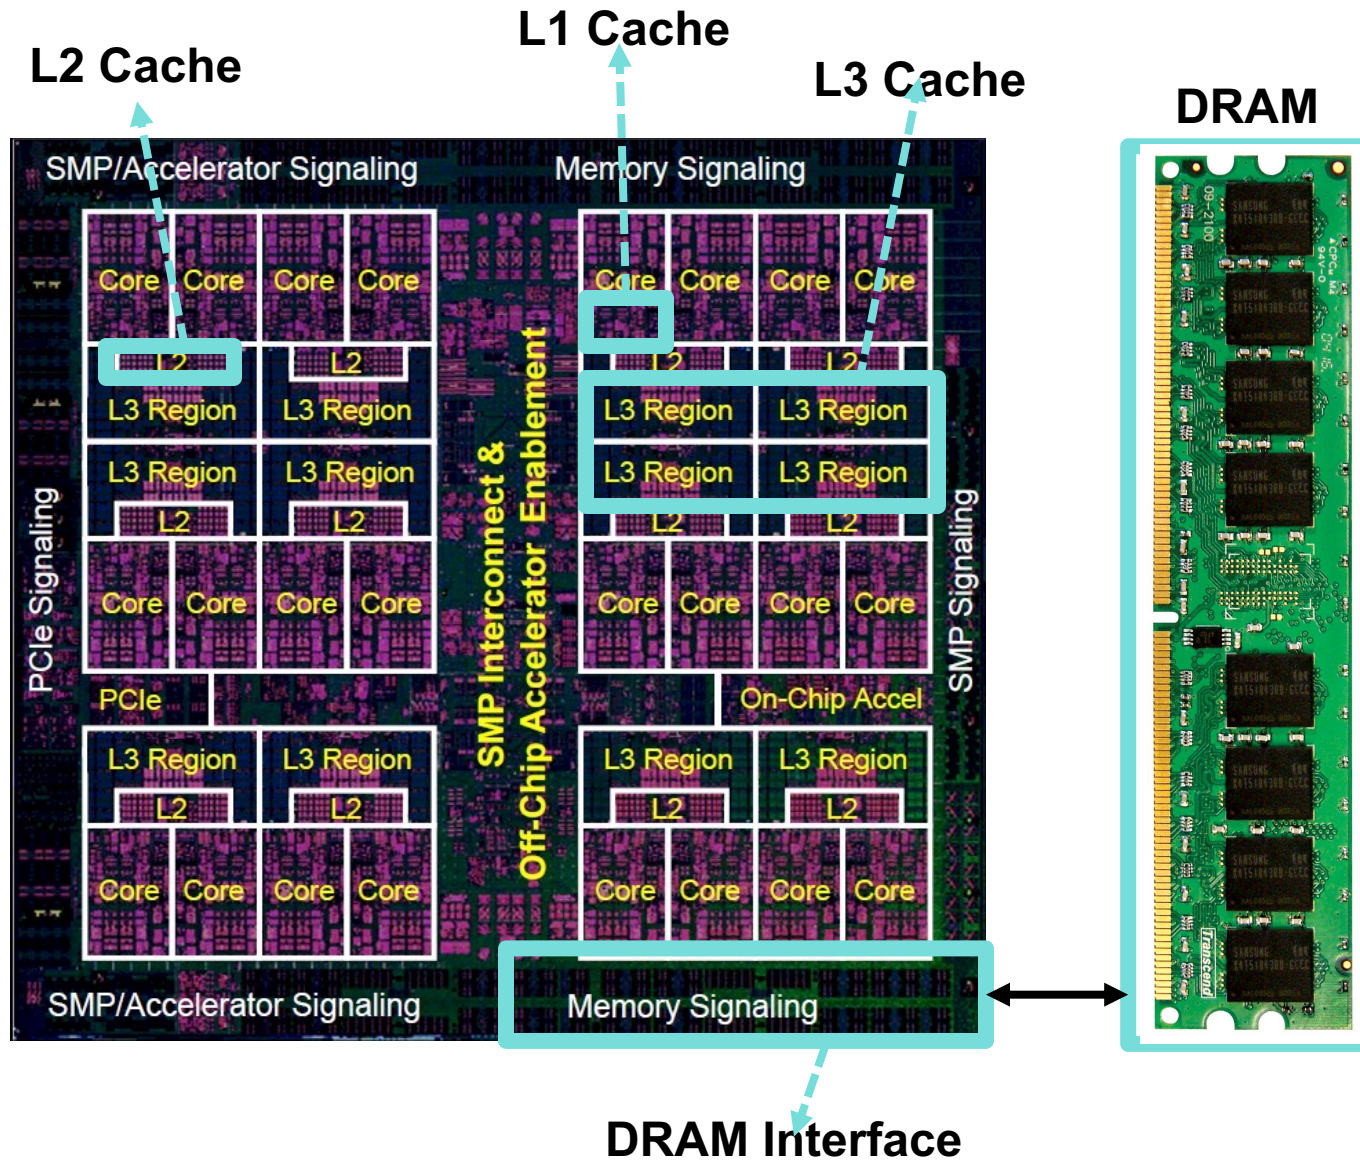


# Memory Hierarchy

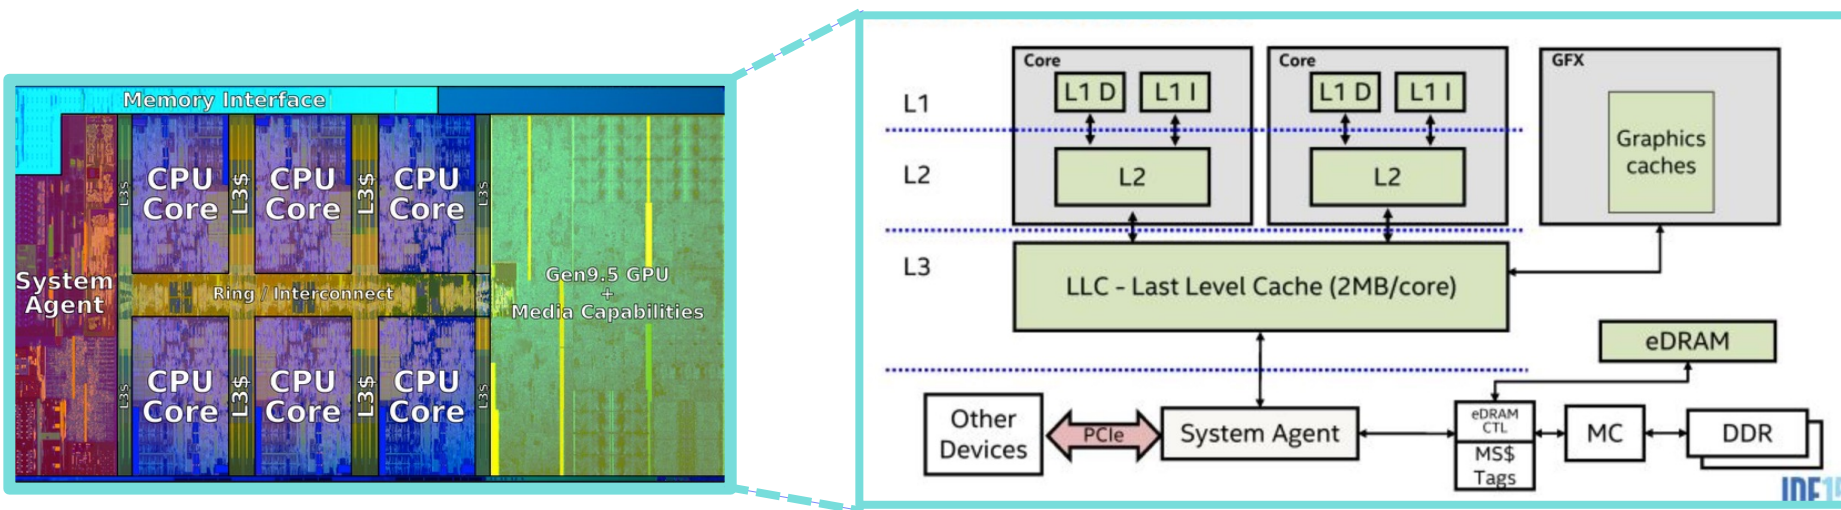
- How to create illusion of the ideal memory (fast, big, cheap)?
  - Have **multiple levels** of memory **with different** speeds and sizes and ensure most of the **data is kept in the fast level**



# CPU Memory Hierarchy



# CPU Memory Hierarchy



- **L1I Cache:** 32~64 KB/core
- **L1D Cache:** 32~64KB/core
- **L2 Cache:** 256KB~1MB
- **L3 Cache:** 1~2MB/core
- Some CPUs use L3 Cache

# Why does the Memory Hierarchy work well?

---

- Answer: **Locality**

- Programs access a relatively small portion of their address space at any instant of time

```
for (i = 0; i < 10; ++i) {  
    j = 17 * i;  
}
```

- Two different types of locality

- **Temporal locality**

- Data accessed recently are likely to be accessed again soon
- E.g., instructions in a loop, induction variables

- **Spatial locality**

- Data near those accessed recently are likely to be accessed soon
- E.g., sequential instruction access, array data

# **Importance of Memory Access Efficiency**

# Poor Memory Access → Poor Performance

---

- **The poor performance is attributable to**

- The long access latencies (hundreds of clock cycles)
- Finite access bandwidth (Giga bytes per second) of global memory

- **Compute-to-global-memory-access-ratio**

- The number of calculations performed per each access to the global memory
- used to calculate the expected performance level of a kernel code
- **High compute-to-global-memory-access-ratio**
  - **High performance**

# Expected Performance of the Matrix Multiplication Kernel

---

## ■ Assumption

- Matrix A and B are in the global memory (slow but big)
- Bandwidth of Global memory : 1000GB/s = 250 giga single-precision numbers per second

## ■ Compute-to-global-memory-access-ratio of for-loop

- # of Computation: 1 Addition + 1 Multiplication = 2
- # of Global memory accesses : 1 for A and 1 for B = 2
- **Compute-to-global-memory-access-ratio = 1**

```
.....  
    for (register int k = 0; k < width; ++k)  
    {  
        float lhs = a[y * width + k];  
        float rhs = b[k * width + x];  
        sum += lhs * rhs;  
    }  
.....
```

## ■ Expected Performance of for-loop = 250 GFLOP

- Because compute-to-global-memory-access-ratio is 1 and 250 giga single-precision numbers can be brought from global memory

## ■ If peak performance of GPU is 12TFLOPS,

- **compute-to-global-memory-access-ratio should be 48 or higher**

# Agenda

---

- **Matrix Multiplication**
  - Simple Version
  - Tiled Version
- Review: Memory Hierarchy
- Importance of Memory Access Efficiency
- GPU Memory Hierarchy
- Improving Tiled Matrix Multiplication
- Impact of Memory on Parallelism