

CUDA Memory Model 2

Prof. Seokin Hong

Agenda

- **Matrix Multiplication**
 - Basic Version
 - Tiled Version
- Review: Memory Hierarchy
- Importance of Memory Access Efficiency
- **GPU Memory Hierarchy**
- Improving Tiled Matrix Multiplication
- Impact of Memory on Parallelism

GPU Memory Hierarhcy

Review: Thread Organization

- Hierarchical Structure
- Grid, block, thread

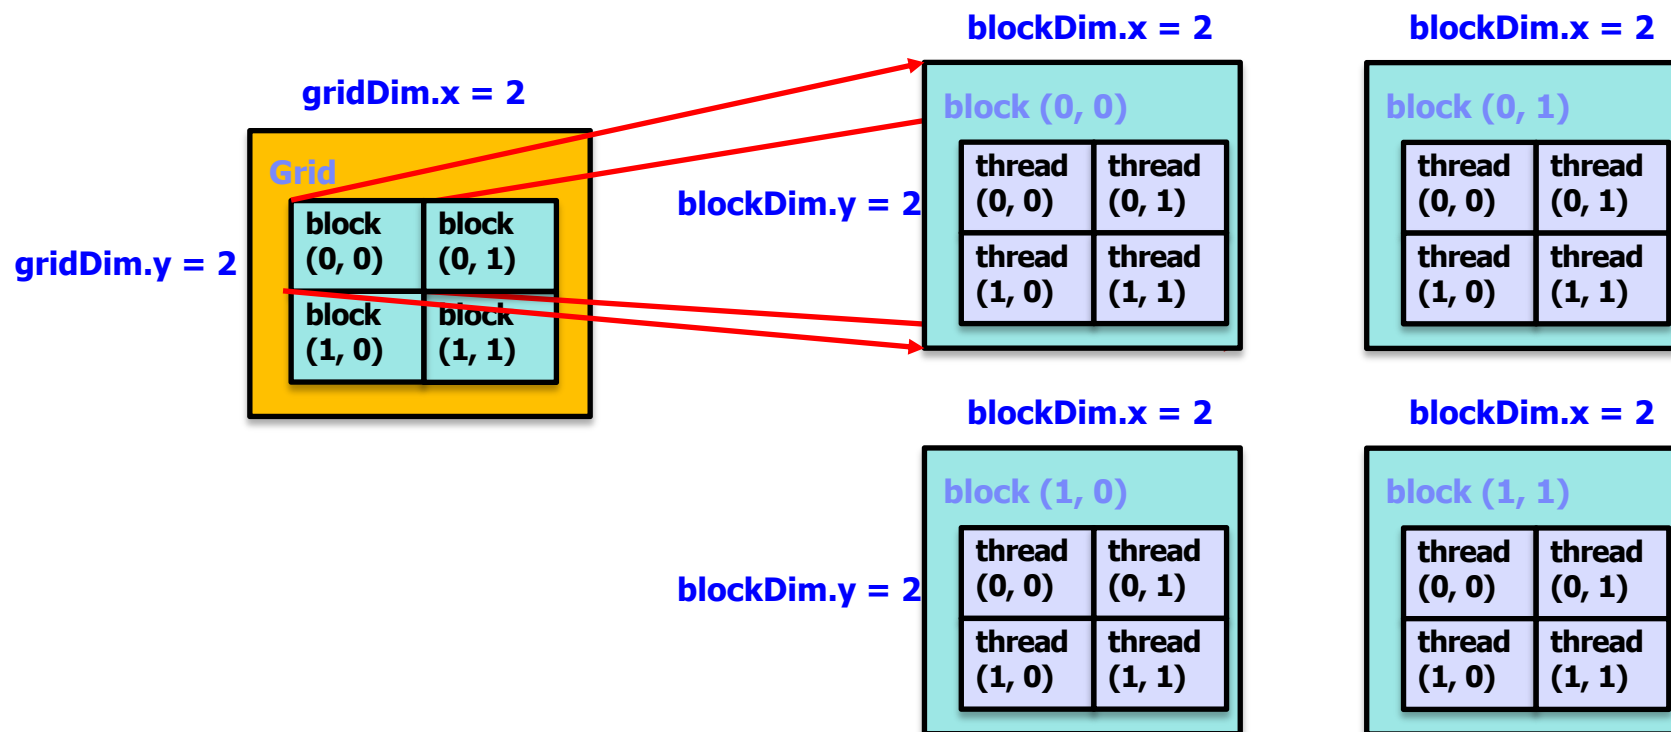


Image from <http://developer.amd.com/zones/OpenCLZone/courses/pages/Introductory-OpenCL-SAAHPC10.aspx>

Programmer's View of CUDA Memories

- Each thread can:

- per-thread **registers**

- (~1 cycle)

- per-block **shared memory**

- (~5 cycles)

- per-grid **global memory**

- (~500 cycles)

- per-thread **local memory**

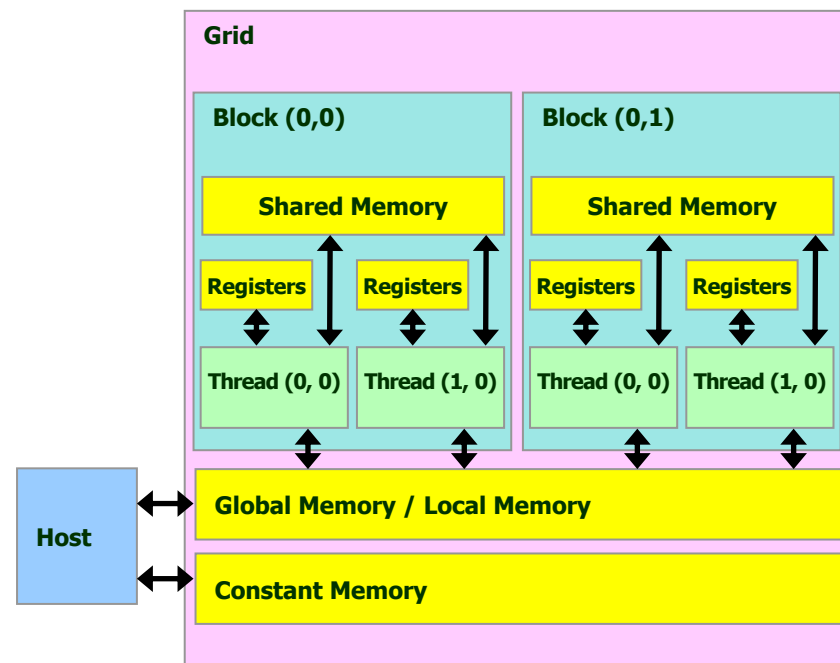
- (~500 cycles)

- actually, located on the global memory

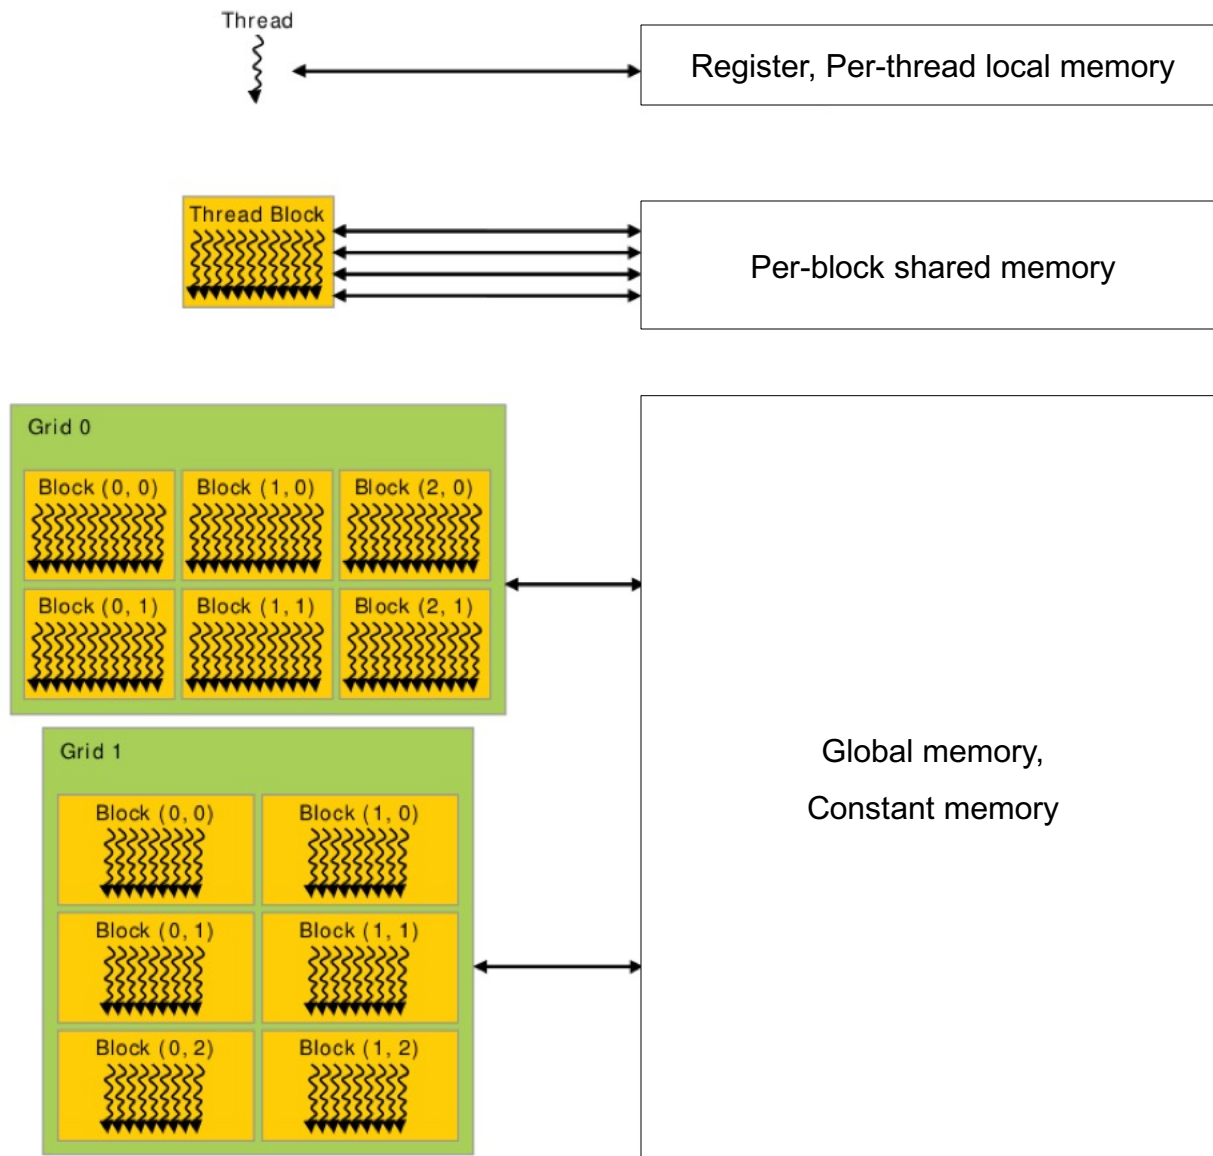
- per-grid **constant memory**

- (~5 cycles with caching)

- Read-only, allocated by host on device, cached on on-chip memory (fast)

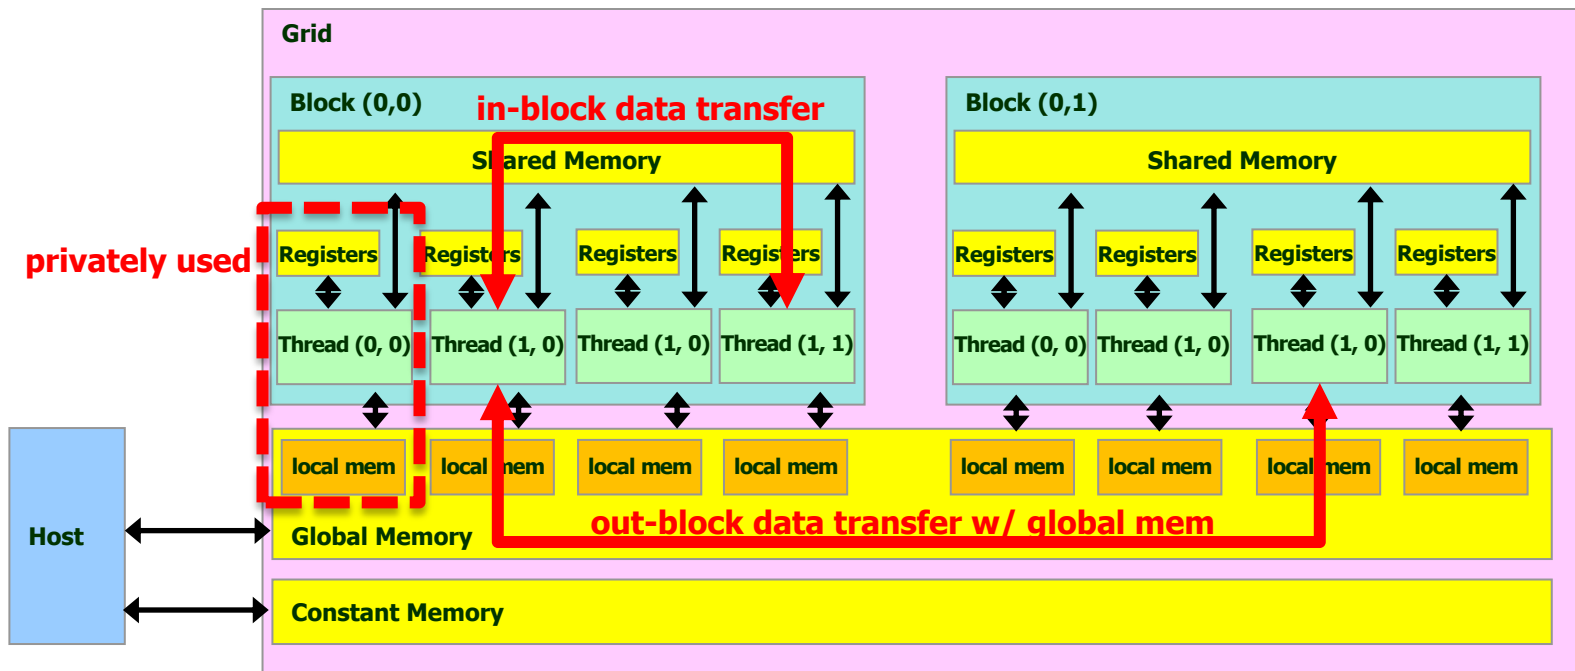


Programmer's View of CUDA Memories




Programmer's View of CUDA Memories

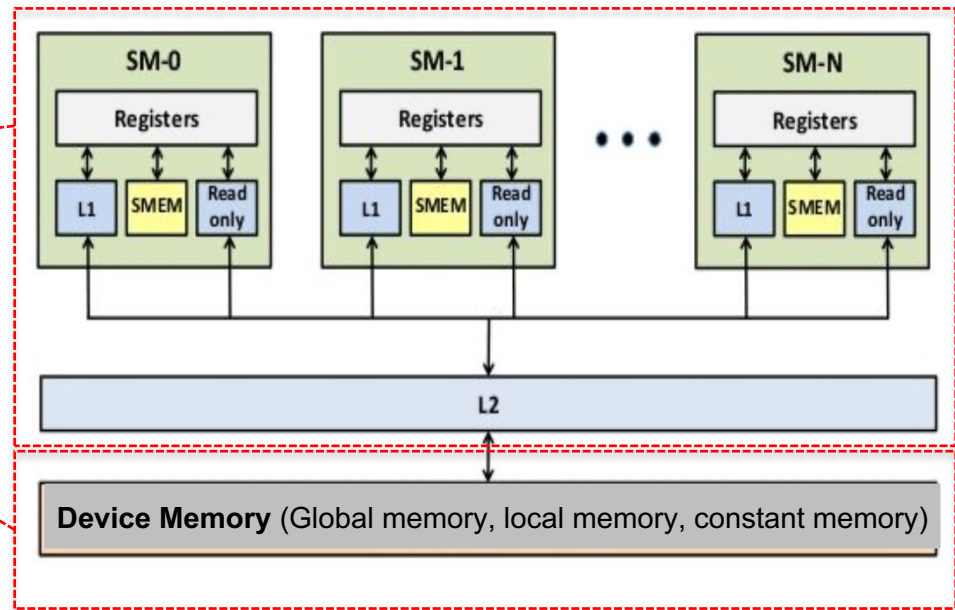
- Each thread has its own private registers and local memory
- **Threads within a block share the shared memory**
 - Shared memory is used for Inter-thread communication within a block
- **Thread blocks share the global memory**
 - Global memory is used for Inter-block (or inter-grid) communication



Architectural View of CUDA Memories



Green box is GK110, red lines are global memory

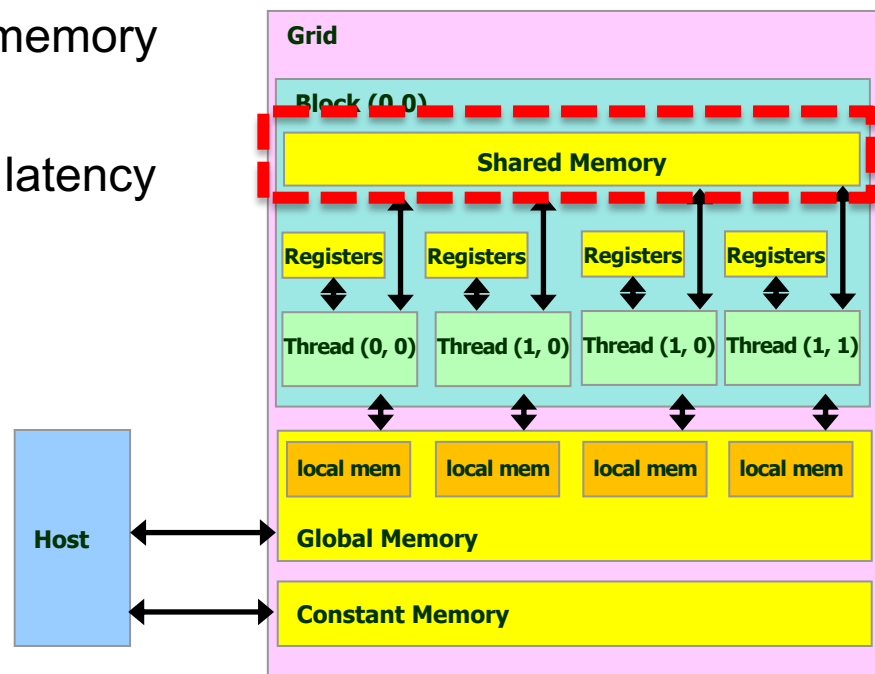


	Global Memory	Shared Memory	Constant Memory	Register
Capacity	16GB (Device memory)	64KB/SM	64KB/SM (Device memory)	256KB / SM
Cache	L1, L2	None	Special cache	None
Access	GPU-wide	SM-wide	GPU-wide (cached on each SM)	Private to each thread
Latency	200~400 cycles	1~5 cycle	1 cycle (hit), 200~400 cycle (miss)	1 cycle

* Numbers can be different depending on GPU architecture

Shared Memory

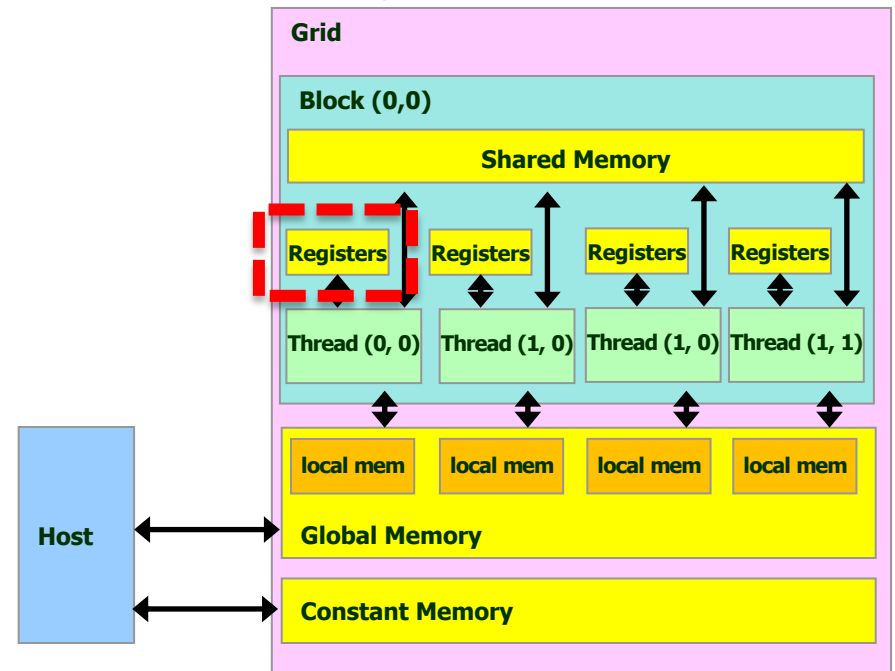
- A special type of memory whose contents are **explicitly declared**
 - **Shared by threads of the same block**
 - Located in the processor (SM)
 - Accessed **at much higher speed** (in both latency and throughput)
 - **Shared memory is partitioned among the blocks**
 - Commonly referred to as scratchpad memory in computer architecture
 - Same hardware as L1 cache, ~5ns of latency



Registers

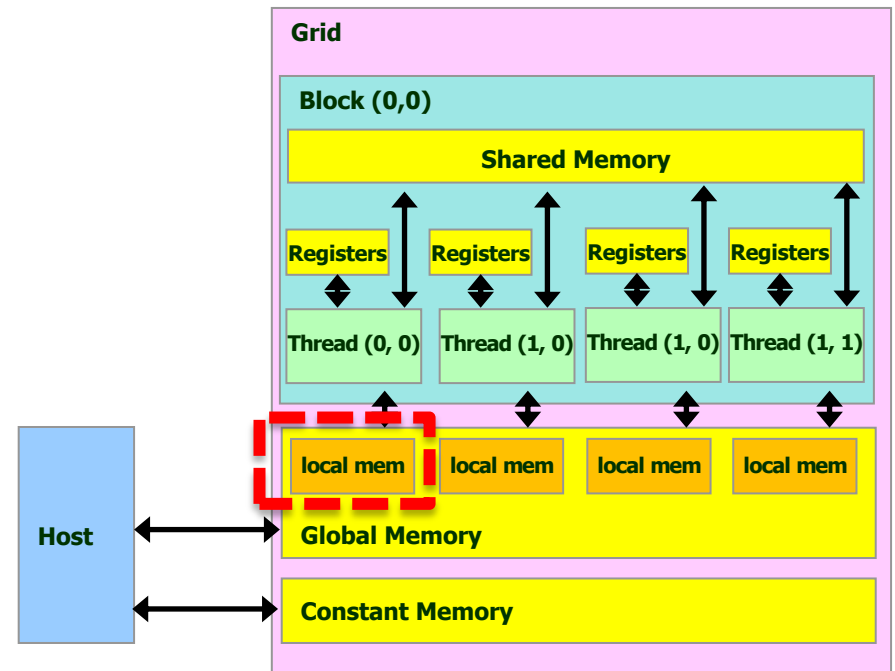
■ Fastest memory

- **Private per thread**
- Located in the processor (SM)
- About 10x faster than shared memory
- **Registers are partitioned among the threads**
- Most local variables declared in kernels are stored in registers (e.g., float x)



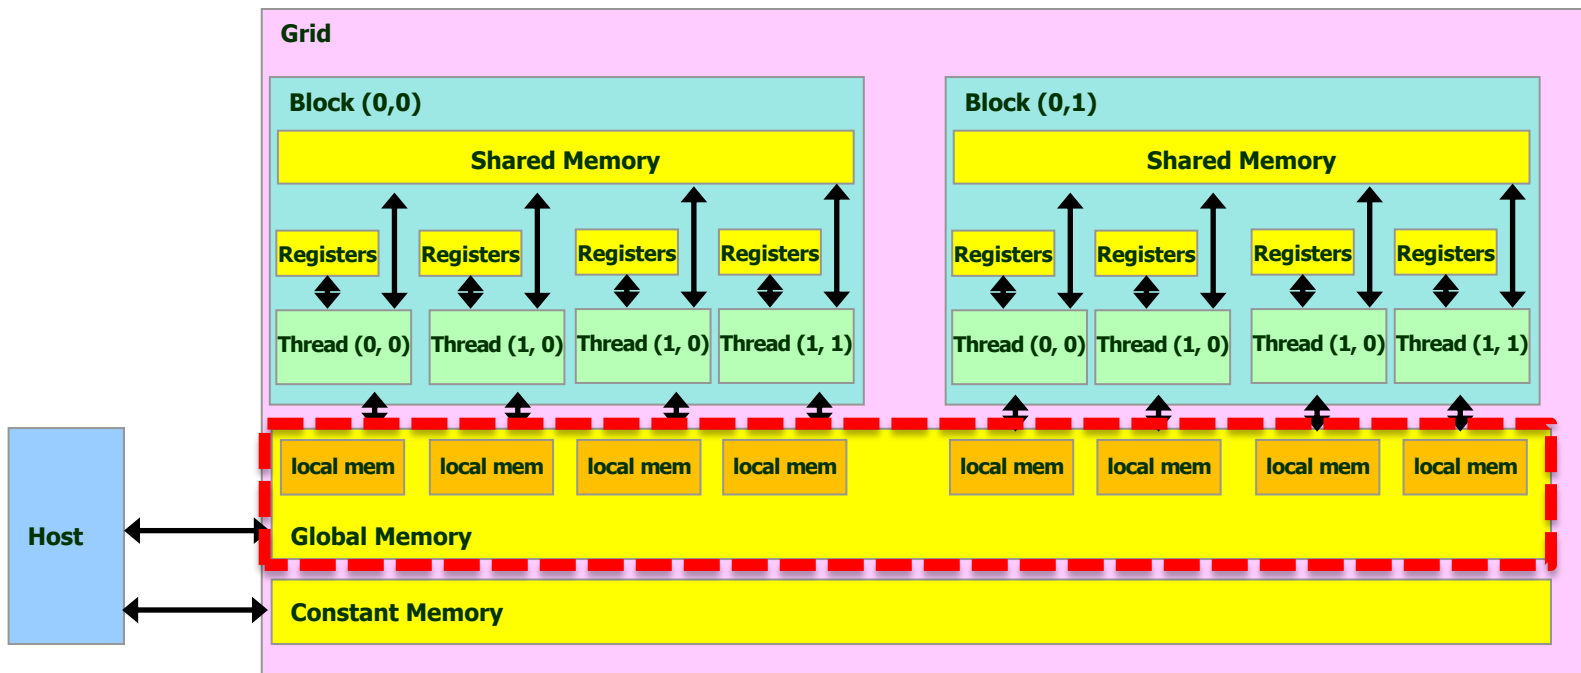
Local Memory

- Local memory stores data that can't fit in registers
 - Private per thread**
 - Stored in the global memory → much slower than registers
 - Used for local variables, register spilling



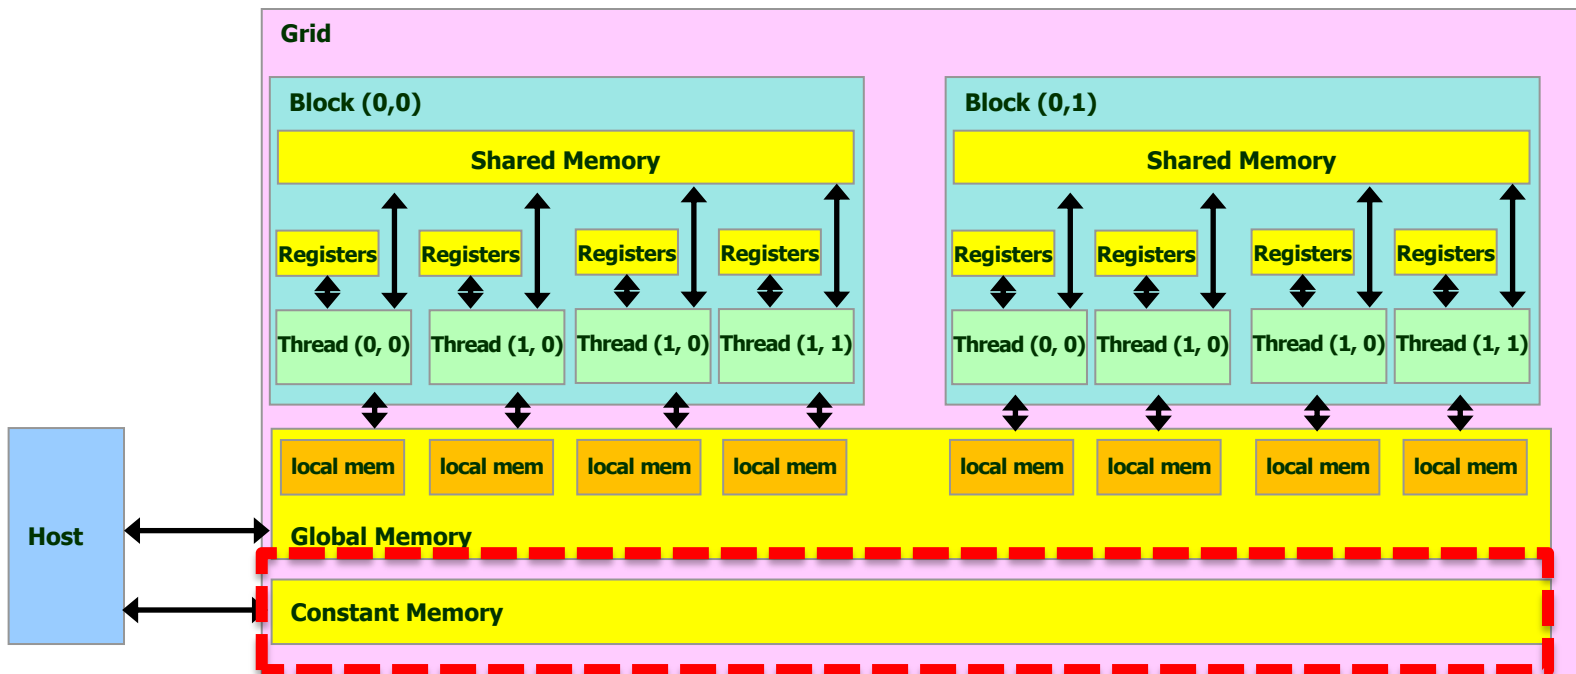
Global Memory

- Global memory is separate hardware from the GPU core
 - **Shared by all threads**
 - Located in off-chip device memory
 - → much slower than registers
 - Majority of data is in global memory
 - Both Host and GPU can access



Constant Memory

- **Used for constant values (read-only data)**
 - **Shared by all threads**
 - Located in off-chip device memory
 - → much slower than registers
 - Both Host and GPU can access
 - Constants must be set from host before running kernel



L1/L2 Cache (Hardware-managed)

■ On-chip Cache memory

- Store recently accessed data in the global memory

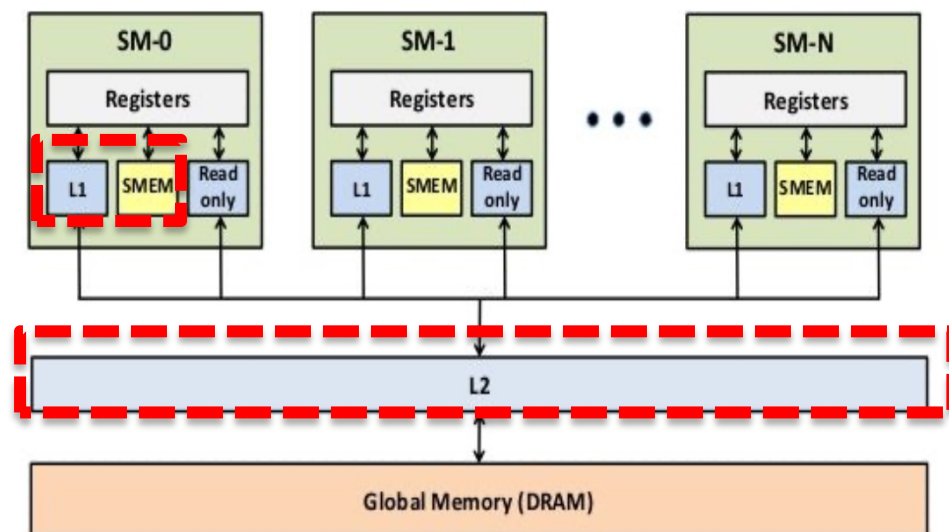
- Store local & global memory data

- **L1 Cache**

- Same hardware as shared memory
 - Configurable size (16, 32, 48KB)
 - Each SM has its own L1 cache

- **L2 Cache**

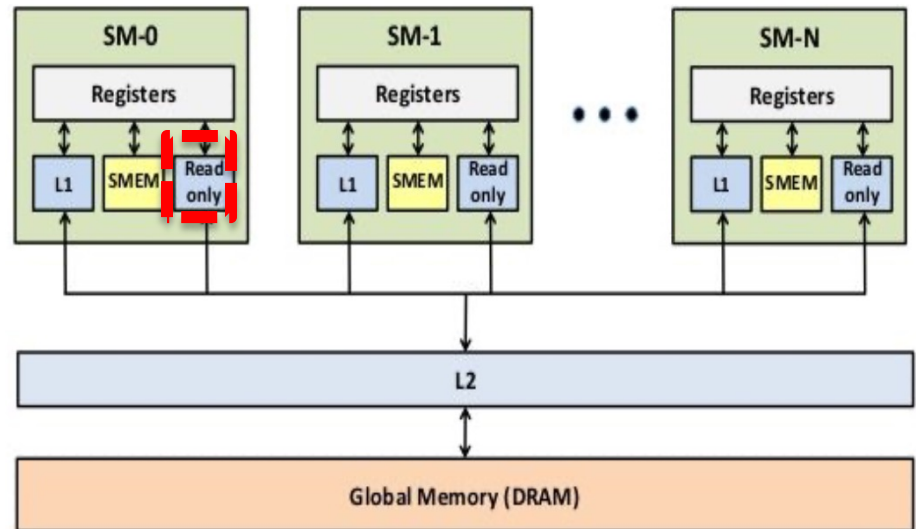
- 1MB~ in size
 - Shared by all SM's



Constant Cache (Hardware-managed)

■ On-chip Cache memory

- Store recently accessed constant in the constant memory
- Hardware-managed



GPU Memory in AWS server

```
seokin@ip-172-31-41-4: /usr/local/cuda/samples/1_Uutilities/deviceQuery$ ./deviceQuery
./deviceQuery Starting...
```

```
CUDA Device Query (Runtime API) version (CUDART static linking)
```

```
Detected 1 CUDA Capable device(s)
```

```
Device 0: "Tesla T4"
```

```
CUDA Driver Version / Runtime Version      11.0 / 11.0
```

```
CUDA Capability Major/Minor version number: 7.5
```

```
Total amount of global memory:             15110 MBytes (15843721216 bytes)
```

```
(40) Multiprocessors, (64) CUDA Cores/MP:   2560 CUDA Cores
```

```
GPU Max Clock rate:                        1590 MHz (1.59 GHz)
```

```
Memory Clock rate:                         5001 Mhz
```

```
Memory Bus Width:                           256-bit
```

```
L2 Cache Size:                             4194304 bytes
```

```
Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
```

```
Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
```

```
Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
```

```
Total amount of constant memory:           65536 bytes
```

```
Total amount of shared memory per block:   49152 bytes
```

```
Total number of registers available per block: 65536
```

```
Warp size:                                  32
```

```
Maximum number of threads per multiprocessor: 1024
```

```
Maximum number of threads per block:        1024
```

```
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
```

```
Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
```

```
Maximum memory pitch:                       2147483647 bytes
```

```
Texture alignment:                           512 bytes
```

```
Concurrent copy and kernel execution:       Yes with 3 copy engine(s)
```

```
Run time limit on kernels:                   No
```

```
Integrated GPU sharing Host Memory:         No
```

```
Support host page-locked memory mapping:    Yes
```

```
Alignment requirement for Surfaces:         Yes
```

```
Device has ECC support:                     Enabled
```

```
Device supports Unified Addressing (UVA):   Yes
```

```
Device supports Managed Memory:             Yes
```

```
Device supports Compute Preemption:         Yes
```

```
Supports Cooperative Kernel Launch:         Yes
```

```
Supports MultiDevice Co-op Kernel Launch:   Yes
```

```
Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 30
```

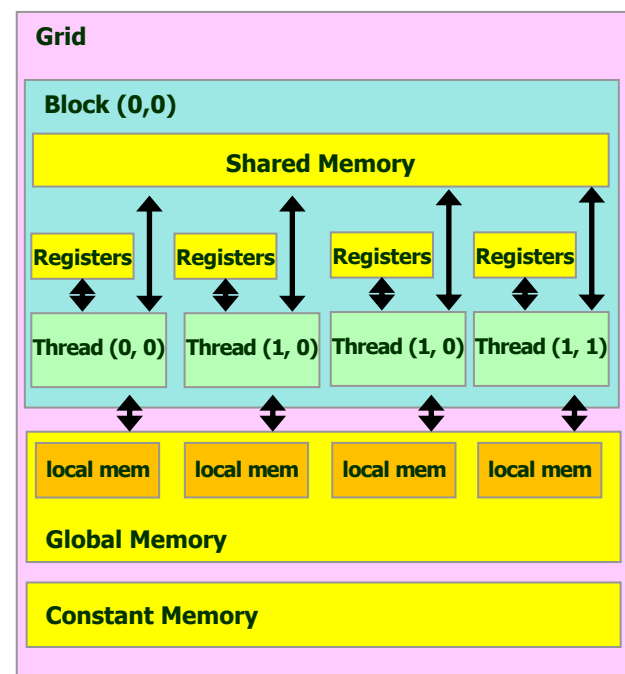
```
Compute Mode:
```

```
< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
```


CUDA Variable Type Qualifiers

Variable declaration	Memory	Scope	Lifetime
<code>int var;</code>	register	thread	thread
<code>int array_var[10];</code>	local	thread	thread
<code>__shared__ int shared_var;</code>	shared	block	block
<code>__device__ int global_var;</code>	global	grid	application
<code>__constant__ int constant_var;</code>	constant	grid	application

- **scalar variables** without qualifier reside in a register
 - compiler will spill to thread-local memory
- **array variables** without qualifier reside in local memory

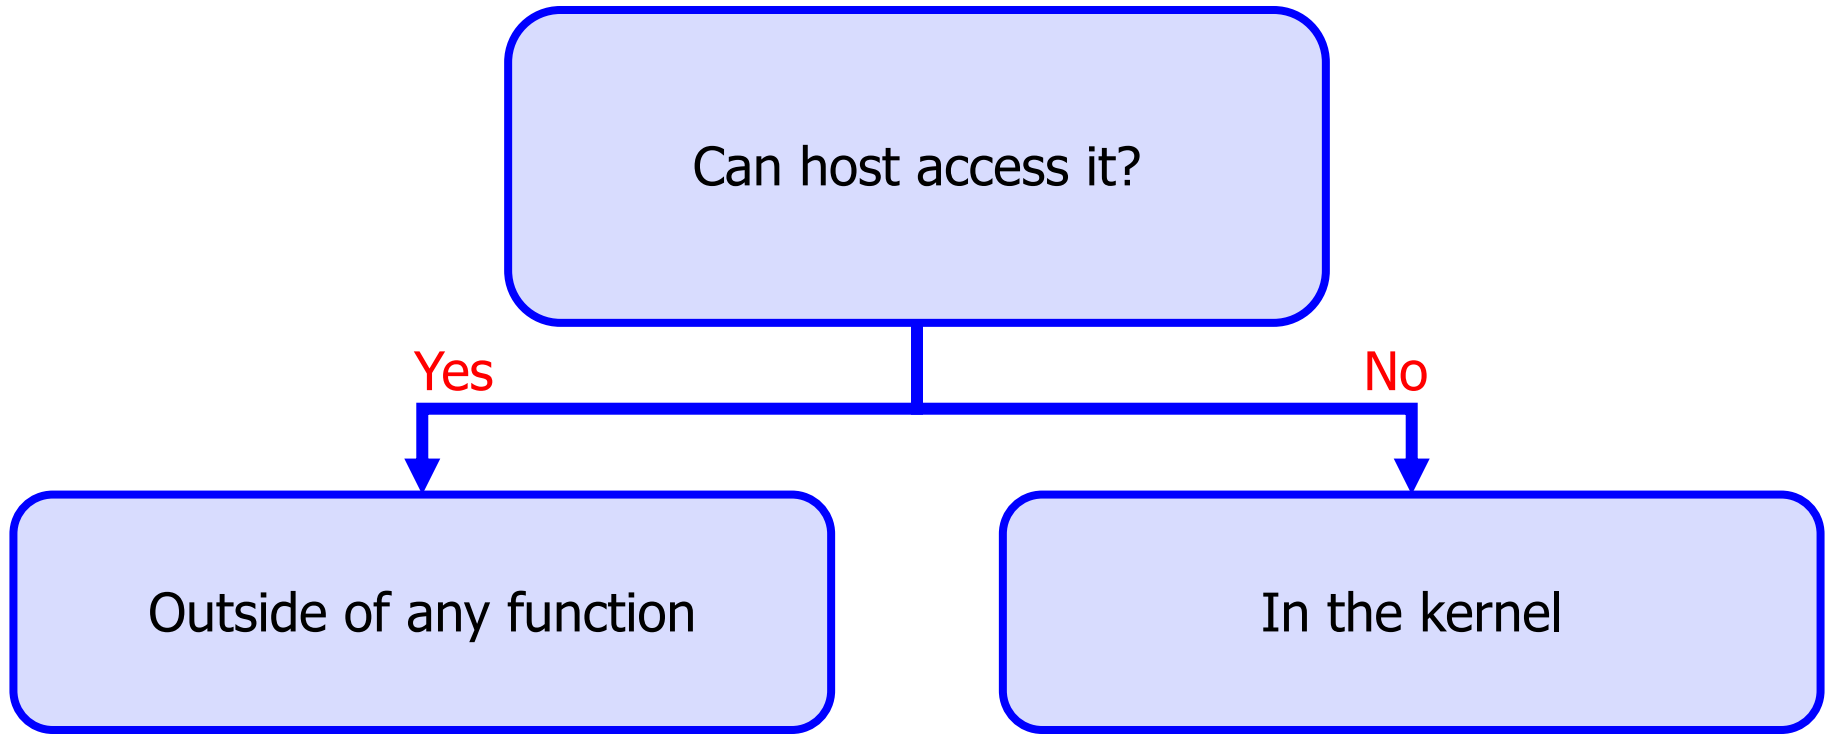


CUDA Variable Type Performance

Variable declaration	Memory	Penalty
<code>int var;</code>	register	1x
<code>int array_var[10];</code>	local	100x
<code>__shared__ int shared_var;</code>	shared	1x
<code>__device__ int global_var;</code>	global	100x
<code>__constant__ int constant_var;</code>	constant	1x

- scalar variables reside in on-chip registers → fast
- shared variables reside in on-chip memories → fast
- Thread-local & global variables reside in uncached off-chip memory → slow
- Thread-local & global variables reside in on-chip cache memory (L1, L2) → fast
- constant variables reside in on-chip cache memory (constant cache) → fast

Where to declare variables?

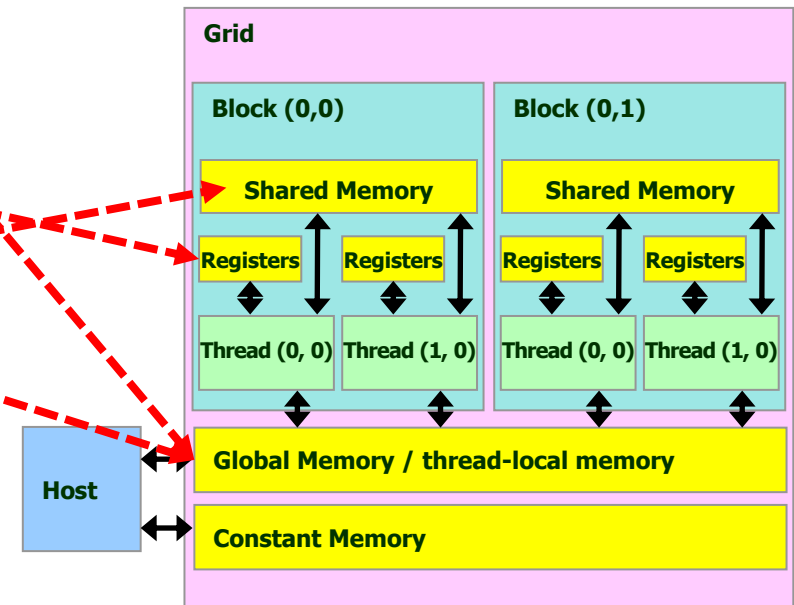


```
__constant__ int constant_var;  
__device__  int global_var;
```

```
int var;  
int array_var[10];  
__shared__ int shared_var;
```

Example: thread-local variables

```
__global__ void kernelFunc(float* dst, const float* src) {  
    // p goes in a register  
    float p = src[threadIdx.x];  
    // per-thread heap goes in off-chip memory  
    float heap[10];  
    // shared variables  
    __shared__ float partial_sum[1024];  
    // now actions  
    ...  
}
```

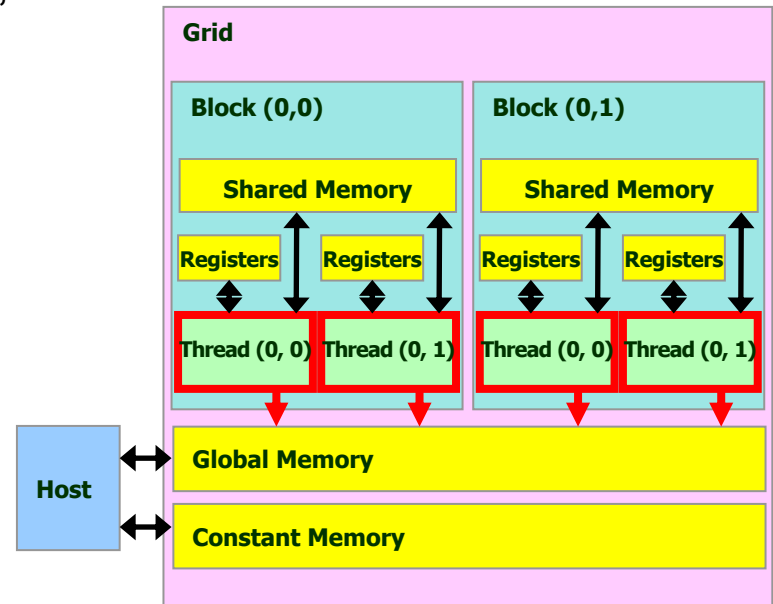


Race Condition: Global Memory Case

- Question:

```
__global__ void raceGlobal( int* dst ) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    dst[0] = idx;  
    // what is the value of dst[0] ?  
}
```

- undecidable !

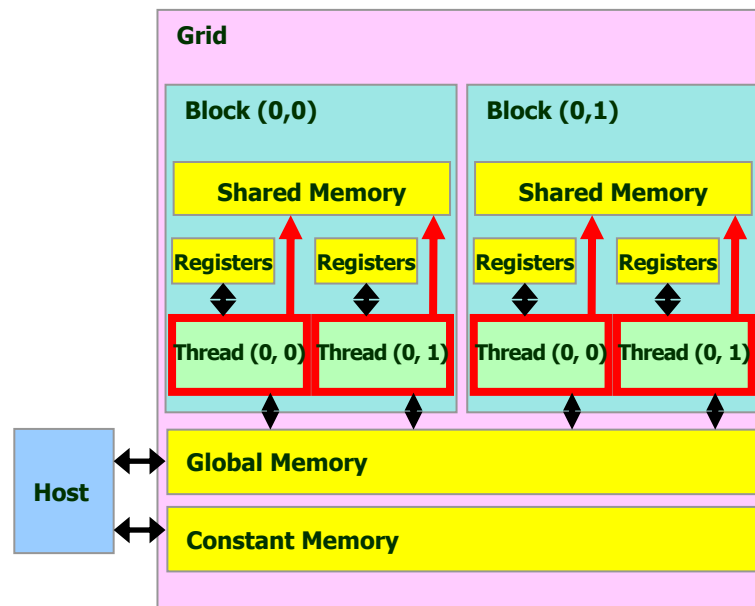


Race Condition: Global Memory Case

■ Question:

```
__global__ void raceShared(void) {  
    __shared__ int shared_dst;  
    shared_dst = threadIdx.x;  
    // what is the value of shared_dst ?  
}
```


■ undecidable !



Communication Through Memory

- The order in which threads access the variable is **undefined** without explicit coordination
- Use **barriers**
 - for shared variables
 - **__syncthreads()** function in CUDA
 - All threads in the same block must reach the **__syncthreads()** before any can move on
- or **atomic operations**
 - for shared variables and global variables → explained later

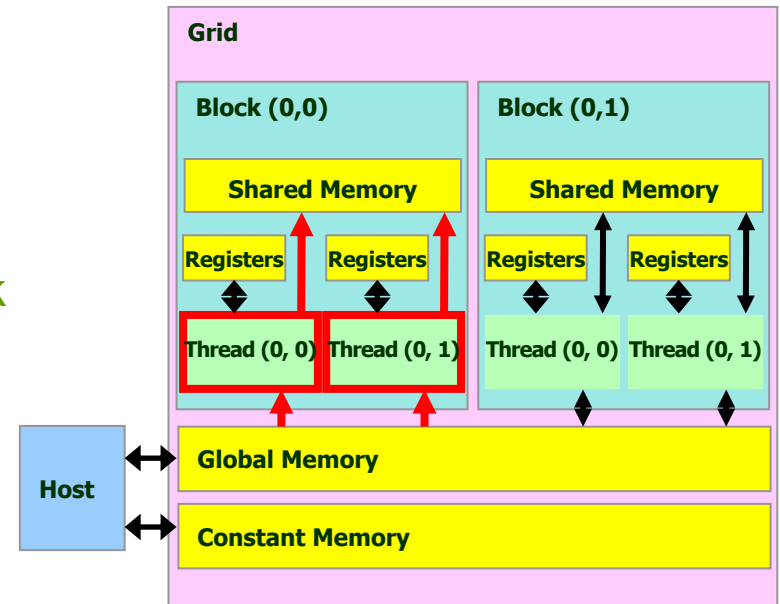
Barrier Synchronization



Barrier Synchronization

- Use `__syncthreads` to ensure data is ready for access


```
__global__ void kernelFunc( int* g_input ) {  
    __shared__ int s_data[BLOCK_SIZE];  
    s_data[threadIdx.x] = g_input[threadIdx.x];  
    __syncthreads();  
    // all data available for all threads in the block  
}
```




Barrier Synchronization

- Use `__syncthreads` to ensure data is ready for access

```
__global__ void kernelFunc( int* g_input ) {  
    __shared__ int s_data[BLOCK_SIZE];  
    s_data[threadIdx.x] = g_input[threadIdx.x];  
    __syncthreads();  
    // all data available for all threads in the block  
  
    ... actions ...  
    // every thread can use the shared data  
}
```



`__syncthreads()`

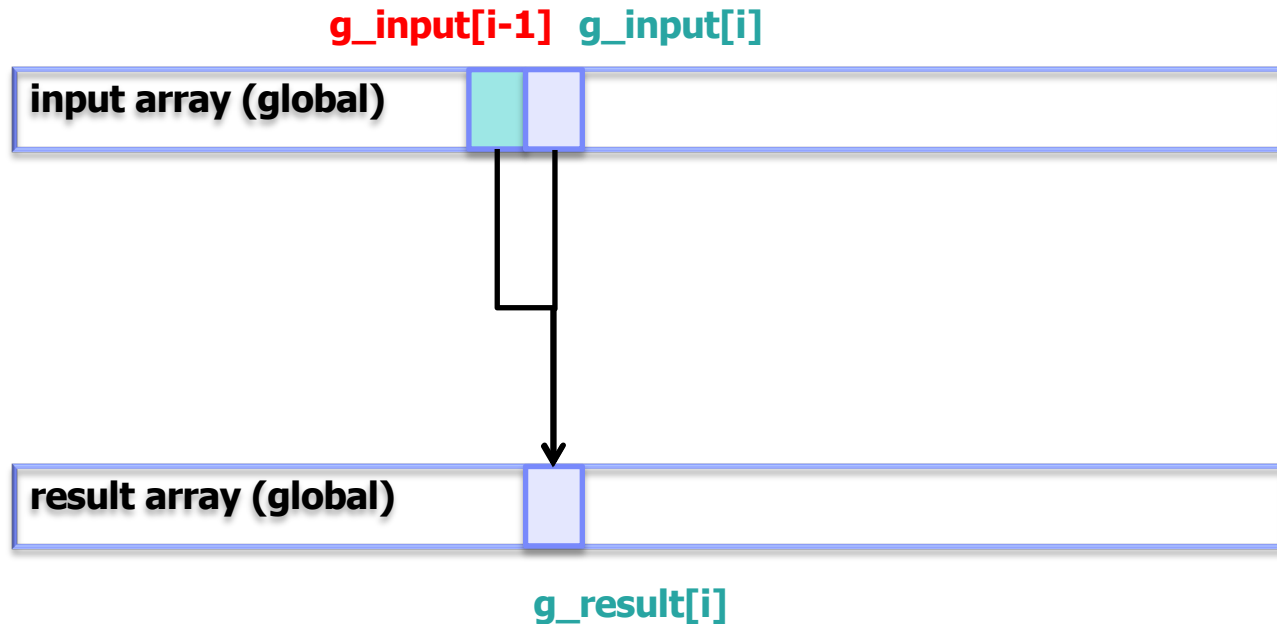


Barrier Synchronization

- Use barriers such as `__syncthreads` to wait until `__shared__` data is ready
- **Don't synchronize or serialize unnecessarily**
 - degrade the program speed ...
 - heavy operation !

Adjacent Difference Example

- $\text{g_result}[i] = \text{g_input}[i] - \text{g_input}[i - 1];$



```
void getDiff(float* dst, const float* src, unsigned int size) {  
    for (int i = 1; i < size; ++i) {  
        dst[i] = src[i] - src[i-1];  
    }  
}
```

Example: CPU version

■ Complete Version (adj_diff_cpu.cu)

```
#include <stdio.h>
#include <stdlib.h> // for rand(), malloc(), free()
#include <fcntl.h>   // for open(), write()
#include <sys/stat.h>
#include "common.h"
#include <sys/time.h>

#define GRIDSIZE          (64 * 1024)
#define BLOCKSIZE 1024
#define TOTALSIZE         (GRIDSIZE * BLOCKSIZE) // 32M byte needed!

void genData(float* ptr, unsigned int size) {
    while (size-- > 0) {
        *ptr++ = (float)(rand() % 1000) / 1000.0F;
    }
}

// compute result[i] = input[i] - input[i-1]
void getDiff(float* dst, const float* src, unsigned int size) {
    for (int i = 1; i < size; ++i) {
        dst[i] = src[i] - src[i-1];
    }
}

....
```

Example: CPU version

■ Complete Version (adj_diff_cpu.cu)

```
.....
int main(void) {
    float* pSource = NULL;
    float* pResult = NULL;
    int i;
    struct timeval start_time, end_time;

    // malloc memories on the host-side
    pSource = (float*)malloc(TOTALSIZE * sizeof(float));
    pResult = (float*)malloc(TOTALSIZE * sizeof(float));
    // generate source data
    genData(pSource, TOTALSIZE);
    // get current time
    gettimeofday(&start_time, NULL);
    getDiff(pResult, pSource, TOTALSIZE);
    // get end time
    gettimeofday(&end_time, NULL);
    double operating_time = (double)(end_time.tv_sec)+(double)(end_time.tv_usec)/1000000.0 -
    ((double)(start_time.tv_sec)+(double)(start_time.tv_usec)/1000000.0);
    printf("Elapsed: %f seconds\n", (double)operating_time);

    // print sample cases
    i = 1;
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);
    i = TOTALSIZE - 1;
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);
    i = TOTALSIZE / 2;
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);
    // free the memory
    free(pSource);
    free(pResult);
}
```

Example: GPU version using global memory

```
// compute result[i] = input[i] - input[i-1]
__global__ void adj_diff_naive( float* g_result, float* g_input ) {
    // compute this thread's global index
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;

    if (i > 0) {
        // each thread loads two elements from global memory
        int x_i    = g_input[i];
        int x_i_minus_1 = g_input[i-1];

        g_result[i] = x_i - x_i_minus_1;
    }
}
```

Example: GPU version using global memory

```
// compute result[i] = input[i] – input[i-1]
```

```
__global__ void adj_diff_naive( float* g_result, float* g_input ) {
```

```
    // compute this thread's global index
```

```
    unsigned int i = blockDim.x * blockIdx.x + threadIdx.x;
```

```
    if (i > 0) {
```

```
        // each thread loads two elements from global memory
```

```
        int x_i = g_input[i];
```

```
        int x_i_minus_1 = g_input[i-1];
```

Two Global Memory Reads

```
        g_result[i] = x_i – x_i_minus_1;
```

One Global Memory Write

Global memory is slow!
g_input[i] is used twice !

Example: GPU version using global memory

■ Complete Version (adj_diff_naive.cu)

```
#include <stdio.h>
#include <stdlib.h> // for rand(), malloc(), free()
#include <fcntl.h>   // for open(), write()
#include <sys/stat.h>
#include "common.h"
#include <sys/time.h>

#define GRIDSIZE          (8 * 1024)
#define BLOCKSIZE 1024
#define TOTALSIZE        (GRIDSIZE * BLOCKSIZE) // 32M byte needed!

void genData(float* ptr, unsigned int size) {
    while (size-- > 0) {
        *ptr++ = (float)(rand() % 1000) / 1000.0F;
    }
}

// compute result[i] = input[i] - input[i-1]
__global__ void adj_diff_naive( float* g_result, float* g_input ) {
    .....
}
.....
```

Example: GPU version using global memory

■ Complete Version: Host code

```
int main(void) {
    float* pSource = NULL;
    float* pResult = NULL;
    int i;
    struct timeval start_time, end_time;

    // malloc memories on the host-side
    pSource = (float*)malloc(TOTALSIZE * sizeof(float));
    pResult = (float*)malloc(TOTALSIZE * sizeof(float));
    // generate source data
    genData(pSource, TOTALSIZE);
    // CUDA: allocate device memory
    float* pSourceDev = NULL;
    float* pResultDev = NULL;
    CUDA_CHECK( cudaMalloc((void**)&pSourceDev, TOTALSIZE * sizeof(float)) );
    CUDA_CHECK( cudaMalloc((void**)&pResultDev, TOTALSIZE * sizeof(float)) );
    // CUDA: copy from host to device
    CUDA_CHECK( cudaMemcpy(pSourceDev, pSource, TOTALSIZE * sizeof(float), cudaMemcpyHostToDevice) );
    // get current time
    cudaThreadSynchronize();
    gettimeofday(&start_time, NULL);
    // CUDA: launch the kernel: result[i] = input[i] - input[i-1]
    dim3 dimGrid(GRIDSIZE, 1, 1);
    dim3 dimBlock(BLOCKSIZE, 1, 1);
    adj_diff_naive<<<dimGrid, dimBlock>>>(pResultDev, pSourceDev);
    // get end time
    cudaThreadSynchronize();
    gettimeofday(&end_time, NULL);
    double operating_time = (double)(end_time.tv_sec)+(double)(end_time.tv_usec)/1000000.0 -
    ((double)(start_time.tv_sec)+(double)(start_time.tv_usec)/1000000.0);
    printf("Elapsed: %f seconds\n", (double)operating_time);
}
```

Example: GPU version using global memory

■ Complete Version: Host code

```
.....  
    // CUDA: copy from device to host  
    CUDA_CHECK( cudaMemcpy(pResult, pResultDev, TOTALSIZE * sizeof(float), cudaMemcpyDeviceToHost) );  
  
    // print sample cases  
    i = 1;  
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);  
    i = TOTALSIZE - 1;  
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);  
    i = TOTALSIZE / 2;  
    printf("i=%2d: %f = %f - %f\n", i, pResult[i], pSource[i], pSource[i - 1]);  
    // CUDA: free the memory  
    CUDA_CHECK( cudaFree(pSourceDev) );  
    CUDA_CHECK( cudaFree(pResultDev) );  
    // free the memory  
    free(pSource);  
    free(pResult);  
}
```

Shared Memory Use Strategy


- Global memory resides in device memory (DRAM)
 - Much slower access than shared memory
- Tile data to take advantage of **fast shared memory**:

Shared Memory Use Strategy




- Partition data into subsets that fit into shared memory

Shared Memory Use Strategy




- Handle each data subset with **one thread block**

Shared Memory Use Strategy




- Load the subset from global memory to shared memory, using multiple threads to exploit memory-level parallelism

Shared Memory Use Strategy



- Perform the computation on the subset from **shared memory**

Shared Memory Use Strategy



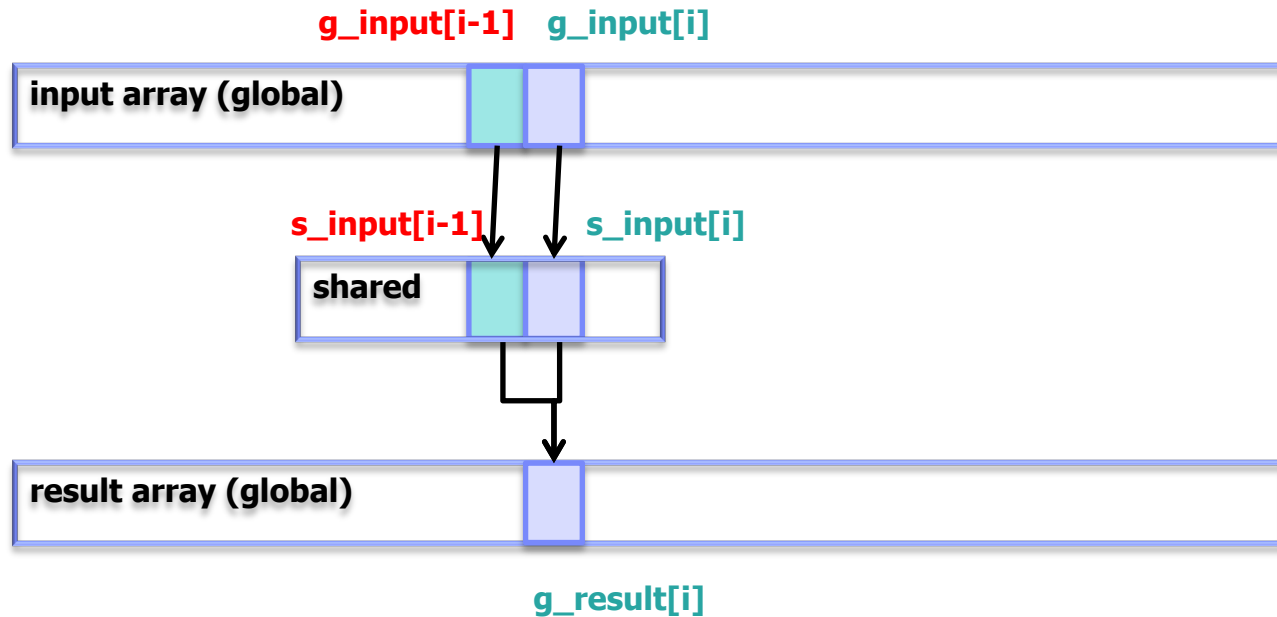
- Copy the result from **shared memory** back to global memory

Shared Memory Use Strategy

- Carefully partition data according to access patterns
- Read-only → `__constant__` memory (fast)
- R/W & shared within block → `__shared__` memory (fast)
- R/W within each thread → `registers` (fast)
- R/W inputs/results → `__global__` memory (slow)


Example: GPU version using shared memory

- $g_result[i] = s_input[i] - s_input[i - 1];$



Example: GPU version using shared memory

```
__global__ void adj_diff(float* g_result, float* g_input) {  
    int tx = threadIdx.x;  
  
    // allocate a __shared__ array, one element per thread  
    __shared__ float s_data[BLOCKSIZE];  
  
    // each thread reads one element to s_data  
    unsigned int i = blockDim.x * blockIdx.x + tx;  
    s_data[tx] = g_input[i];  
  
    // avoid race condition: ensure all loads complete before continuing  
    __syncthreads();  
  
    // now action  
    if (tx > 0) {  
        g_result[i] = s_data[tx] - s_data[tx-1];  
    } else if (i > 0) {  
        // handle thread block boundary (for tx == 0 case)  
        g_result[i] = s_data[tx] - g_input[i-1];  
    }  
}
```



Optimization Analysis

Implementation	Original	Improved
Global Loads	$2N$	$N + N/\text{BLOCK_SIZE}$
Global Writes	N	N
Throughput	36.8 GB/s	57.5 GB/s
source line of codes (SLOC)	18	35
Relative Improvement	1x	1.57x
Improvement/SLOC	1x	0.81x

- Optimizations tend to come with a development cost

Compare the Execution Time

- Host version: 0.260803 seconds
- CUDA global memory: 0.002964 seconds
- CUDA shared memory: 0.000005 seconds
- Use shared memory if possible!!!

Dynamically allocated shared memory

- when the size of the array **isn't** known at compile time...

```
__global__ void adj_diff(int *result, int *input) {  
    // use extern to indicate a __shared__ array will be  
    // allocated dynamically at kernel launch time  
    extern __shared__ int s_data[];  
  
    ...  
}  
  
// pass the size of the per-block array, in bytes, as the third  
// argument to the triple chevrons  
adj_diff<<<num_blocks, block_size, block_size * sizeof(int)>>>(r,i);
```

Next?

- **Matrix Multiplication**
 - Basic Version
 - Tiled Version
- Review: Memory Hierarchy
- Importance of Memory Access Efficiency
- GPU Memory Hierarchy
- **Improving Tiled Matrix Multiplication using Shared Memory**
- **Impact of Memory on Parallelism**