

Parallel Patterns: Parallel Histogram Computation with Atomic Operations

Prof. Seokin Hong

Objective

- To understand atomic operations
 - Read-modify-write in parallel computation
 - Use of atomic operations in CUDA

- Parallel Histogram Computation as an example application of atomic operations

Atomic Operations

Race Condition

Example: Mem[x] += 1

thread1: Old \leftarrow Mem[x]
 New \leftarrow Old + 1
 Mem[x] \leftarrow New

thread2: Old \leftarrow Mem[x]
 New \leftarrow Old + 1
 Mem[x] \leftarrow New

- If Mem[x] was initially 0,
 - What would the value of Mem[x] be after threads 1 and 2 have completed?
 - What does each thread get in their Old variable?

- The answer may vary due to data races.

Race Condition(cont'd)

Timing Scenario #1

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3	(1) Mem[x] \leftarrow New	
4		(1) Old \leftarrow Mem[x]
5		(2) New \leftarrow Old + 1
6		(2) Mem[x] \leftarrow New

- Thread 1 Old = 0
- Thread 2 Old = 1
- Mem[x] = 2 after the sequence

Race Condition(cont'd)

Timing Scenario #2

Time	Thread 1	Thread 2
1		(0) Old \leftarrow Mem[x]
2		(1) New \leftarrow Old + 1
3		(1) Mem[x] \leftarrow New
4	(1) Old \leftarrow Mem[x]	
5	(2) New \leftarrow Old + 1	
6	(2) Mem[x] \leftarrow New	

- Thread 1 Old = 1
- Thread 2 Old = 0
- Mem[x] = 2 after the sequence

Race Condition(cont'd)

Timing Scenario #3

Time	Thread 1	Thread 2
1	(0) Old \leftarrow Mem[x]	
2	(1) New \leftarrow Old + 1	
3		(0) Old \leftarrow Mem[x]
4	(1) Mem[x] \leftarrow New	
5		(1) New \leftarrow Old + 1
6		(1) Mem[x] \leftarrow New

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Race Condition(cont'd)

Timing Scenario #4

Time	Thread 1	Thread 2
1		(0) Old \leftarrow Mem[x]
2		(1) New \leftarrow Old + 1
3	(0) Old \leftarrow Mem[x]	
4		(1) Mem[x] \leftarrow New
5	(1) New \leftarrow Old + 1	
6	(1) Mem[x] \leftarrow New	

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Race Condition(cont'd)

- In the parallel execution,
 - Execution Sequence is not defined by the programming model,
 - It can be arbitrary
- Then, what is the solution?
- **Atomic operations !**
 - CUDA provides **atomic operations** to deal with this problem

Atomic Operations

- To Ensure Good Outcomes

thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New


Or

thread1: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New


thread2: Old \leftarrow Mem[x]
New \leftarrow Old + 1
Mem[x] \leftarrow New

Atomic Operations (cont'd)

- memory locking




- atomic operation




Atomic Operations

■ atomic operation #1



■ atomic operation #2



Atomic Operations

- **The operation is atomic.**

- It can not be divided into sub-steps. It is **a single step.**

- The name **atomic** comes from the fact that **it is uninterruptable**

- An atomic operation guarantees that **only a single thread has access to a piece of memory until an operation completes.**

- Different types of atomic instructions

- **Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor,...**

CUDA atomic CAS operation

- CAS : **compare and swap**
 - the most fundamental operation among all atomic operations
- **int atomicCAS(int* address, int expected, int newVal);**
 - step 1. oldVal = read (*address)
 - step 2. (*address) = (oldVal == expected) ? newVal : oldVal
 - step 3. return oldVal
 - **do it in an atomic manner !**
- How to use it?
 - oldVal = *address;
 - readback = atomicCAS(address, oldVal, newVal);
 - if (oldVal == readback) success!
 - else fail... **(another thread did another CAS)**

atomicCAS example

```
__device__ inline void MyAtomicAdd(float *address, float value) {  
    int oldval, newval, readback;  
    oldval = __float_as_int(*address);  
    newval = __float_as_int(__int_as_float(oldval) + value);  
    while ((readback=atomicCAS((int *)address, oldval, newval)) != oldval) {  
        oldval = readback;  
        newval = __float_as_int(__int_as_float(oldval) + value);  
    }  
}
```

CUDA Atomic Functions

- `int atomicAdd(int* address, int val);`
- `int atomicSub(int* address, int val);`
- `int atomicExch(int* address, int val);`
- `int atomicMin(int* address, int val);`
- `int atomicMax(int* address, int val);`
- `unsigned int atomicInc(unsigned int* address, unsigned int val);`
- `unsigned int atomicDec(unsigned int* address, unsigned int val);`
- `int atomicCAS(int* address, int compare, int val);`
- `int atomicAnd(int* address, int val);`
- `int atomicOr(int* address, int val);`
- `int atomicXor(int* address, int val);`

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

Example: Count

- make many threads
- set count = 0
- for each thread, $\text{count} = \text{count} + 1$
- at the end, show the count value

Count : non-atomic version

```
#include <stdio.h>
#include <stdlib.h>

#define GRIDSIZE    (32 * 1024)
#define BLOCKSIZE   1024
#define TOTALSIZE   (GRIDSIZE * BLOCKSIZE)

__global__ void kernel(unsigned long long int* pCount) {
    (*pCount) = (*pCount) + 1; // we met race condition !
}

int main(void) {
    unsigned long long int aCount[1];
    // prepare timer
    cudaEvent_t start;
    cudaEvent_t stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
```

Count : non-atomic version

```
// CUDA: allocate device memory
```

```
unsigned long long int* pCountDev = NULL;
```

```
cudaMalloc((void**)&pCountDev, sizeof(unsigned long long int));
```

```
cudaMemset(pCountDev, 0, sizeof(unsigned long long int));
```

```
// start timer
```

```
cudaEventRecord(start, 0);
```

```
// CUDA: launch the kernel
```

```
dim3 dimGrid(GRIDSIZE, 1, 1);
```

```
dim3 dimBlock(BLOCKSIZE, 1, 1);
```

```
kernel<<<dimGrid, dimBlock>>>(pCountDev);
```

```
// CUDA: copy from device to host
```

```
cudaMemcpy(aCount, pCountDev, sizeof(unsigned long long int),  
          cudaMemcpyDeviceToHost);
```

```
printf("total number of threads = %llu\n", TOTALSIZE);
```

```
printf("count = %llu\n", aCount[0]);
```

Count : non-atomic version

```
// end timer
```

```
float time;
```

```
cudaEventRecord(stop, 0);
```

```
cudaEventSynchronize(stop);
```

```
cudaEventElapsedTime(&time, start, stop);
```

```
printf("elapsed time = %f msec\n", time);
```

```
cudaEventDestroy(start);
```

```
cudaEventDestroy(stop);
```

```
// CUDA: free the memory
```

```
cudaFree(pCountDev);
```

```
}
```

Count : non-atomic version

- execution result:

total number of threads = 33554432 ← 32M threads

count = 6323 ← failure !

elapsed time = 4.196864 msec

Count : Atomic Version

- with a new atomic kernel:

```
__global__ void kernel(unsigned long long int* pCount) {  
    atomicAdd(pCount, 1ULL);  
}
```

- no change in the main().

Compilation and Execution Result

- atomic operations are only supported new CUDA architectures...

- So, NVCC with “**-arch sm_61**” option:

```
nvcc -w -arch sm_61 -o count-atomic count-atomic.cu
```

- execution result

total number of threads = 33554432

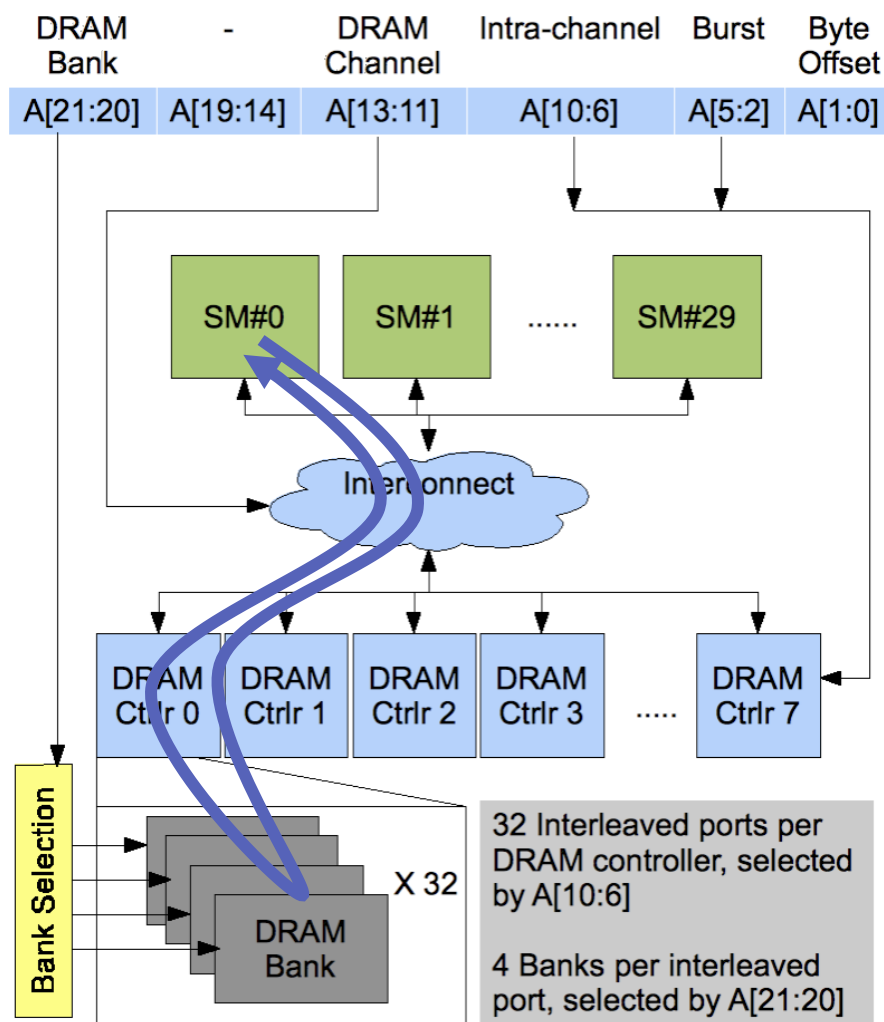
count = 33554432

elapsed time = **452.034241 msec**

- success !

- but, **too slow**... → Why?

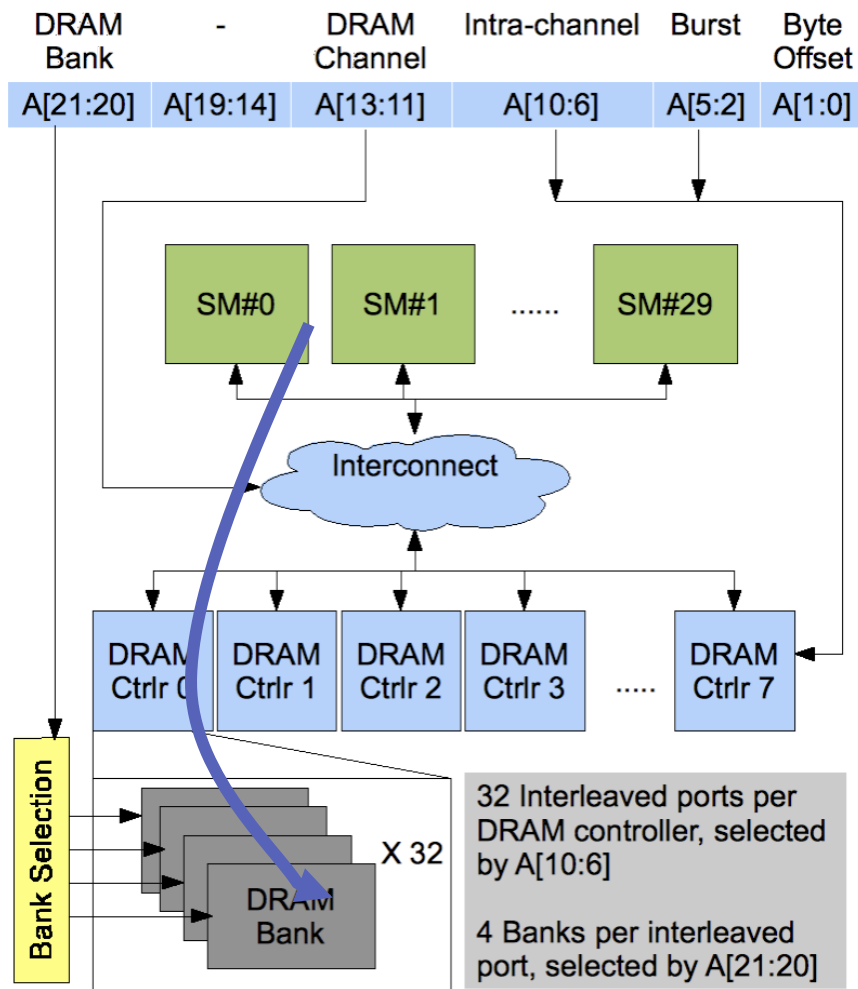
Atomic Operations on DRAM (Global memory)



- An atomic operation starts with a read, with a latency of a few hundred cycles



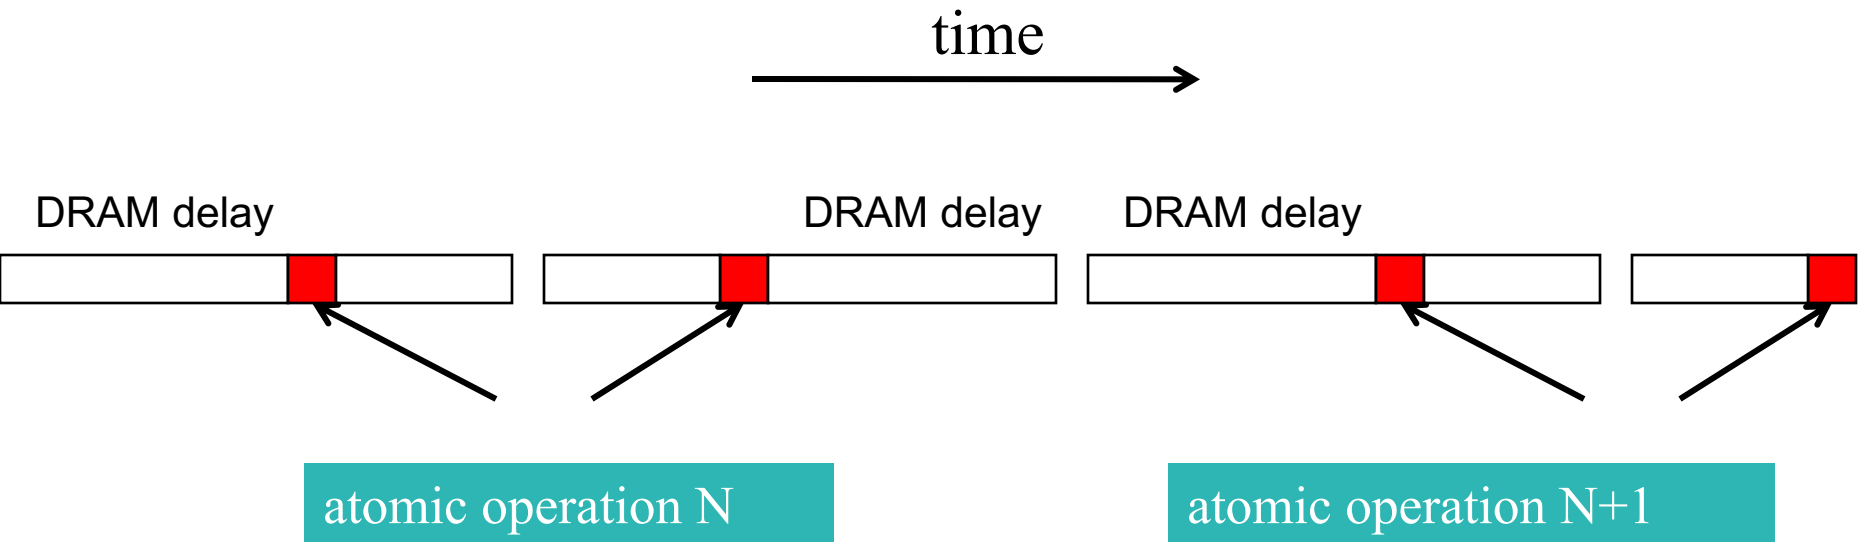
Atomic Operations on DRAM (Global memory)



- An atomic operation starts with a read, with a latency of a few hundred cycles
- The atomic operation ends with a write, with a latency of a few hundred cycles
- **During this whole time, no one else can access the location**

Atomic Operations on DRAM (Global memory)

- Each Load-Modify-Store has two full memory access delays
 - All atomic operations on the same variable (RAM location) are **serialized**



Latency determines throughput of atomic operations

- **Throughput** of an atomic operation is the **rate** at which the application can execute an atomic operation on a particular location.
- **The rate is limited by the total latency** of the read-modify-write sequence, **typically more than 1000 cycles for global memory (DRAM) locations.**
- This means that if many threads attempt to do atomic operation on the same location (contention), **the memory bandwidth is reduced to < 1/1000!**

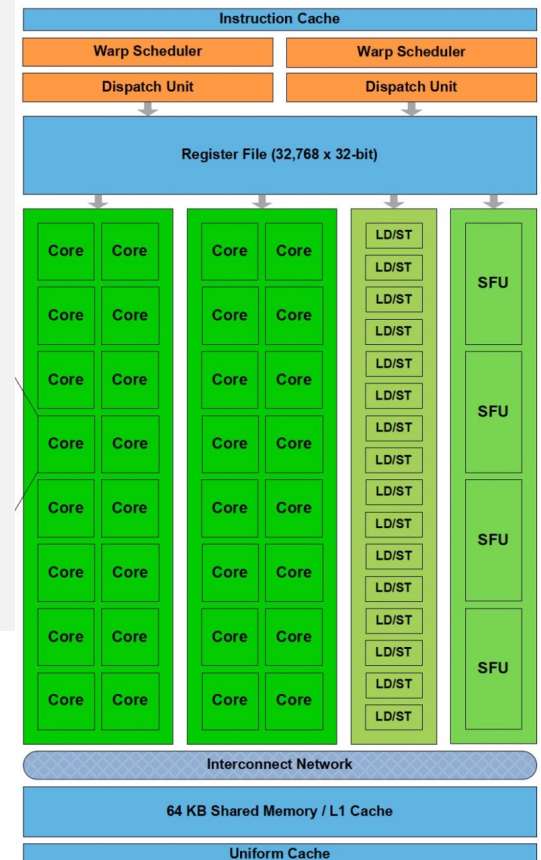
Hardware Improvements

- Atomic operations on **L2 cache**
 - medium latency, but still serialized
- Atomic operations on **Shared Memory**
 - Very short latency, but still serialized
 - Private to each thread block
 - Need algorithm work by programmers

Count : shared memory version

- new kernel with shared memory

```
__global__ void kernel(int* pCount) {  
    __shared__ int nCountShared;  
  
    if (threadIdx.x == 0) {  
        nCountShared = 0;  
    }  
    __syncthreads();  
  
    atomicAdd(&nCountShared, 1);  
  
    __syncthreads();  
  
    if (threadIdx.x == 0) {  
        atomicAdd(pCount, nCountShared);  
    }  
}
```



Execution Result

- execution result

total number of threads = 33554432

count = 33554432


elapsed time = 319.288818 msec

- We achieved some speed ups!

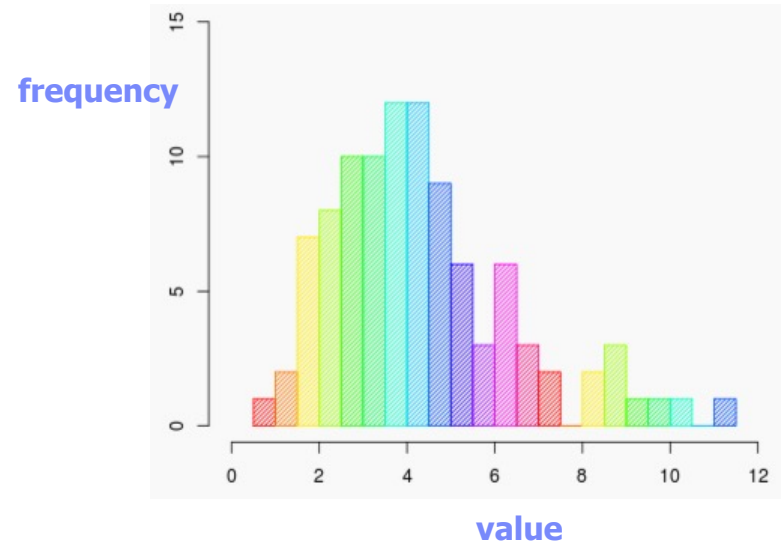

Parallel Histogram Computation

Another Example: Histogram

- histogram analysis
 - a basic tool for image processing



target image



Another Example: Histogram

- in this example,
 - make 1,024 x 1,024 image (integer values)
 - each pixel contains: 0 ~ 14 values randomly
 - make its histogram
 - histogram[0] = number of pixels with value 0
 - histogram[1] = number of pixels with value 1
 - histogram[2] = number of pixels with value 2
 - ...

Histogram : host version

```
#include <stdio.h>
#include <stdlib.h>

#define GRIDSIZE      1024
#define BLOCKSIZE     1024
#define TOTALSIZE     (GRIDSIZE * BLOCKSIZE)
#define      NUMHIST      16

void genData(unsigned int* ptr, unsigned int size) {
    while (size--) {
        *ptr++ = (unsigned int)(rand() % (NUMHIST - 1)); // 0 ~ 14
    }
}

void kernel(unsigned int* hist, unsigned int* img, unsigned int size) {
    while (size--) {
        unsigned int pixelVal = *img++;
        hist[pixelVal] = hist[pixelVal] + 1;
    }
}
```

Histogram : host version

```
int main(void) {  
    unsigned int* pImage = NULL;  
    unsigned int* pHistogram = NULL;  
    int i;  
  
    // malloc memories on the host-side  
    pImage = (unsigned int*)malloc(TOTALSIZE * sizeof(unsigned int));  
    pHistogram = (unsigned int*)malloc(NUMHIST * sizeof(unsigned int));  
    for (i = 0; i < NUMHIST; ++i) {  
        pHistogram[i] = 0;  
    }  
    // generate source data  
    genData(pImage, TOTALSIZE);  
}
```

Histogram : host version

```
// perform the action
kernel(pHistogram, pImage, TOTALSIZE);

...

// print the histogram
long total = 0L;
for (i = 0; i < NUMHIST; ++i) {
    printf("%2d: %10d\n", i, pHistogram[i]);
    total += pHistogram[i];
}
printf("total: %10ld (should be %ld)\n", total, TOTALSIZE);

...
}
```

Execution Result:

elapsed time = 4339.358362 usec

0: 69692

1: 69634

2: 70121

3: 70277

4: 70215

5: 69479

6: 69988

7: 70344

8: 69984

9: 69976

10: 70099

11: 69686

12: 69810

13: 69909

14: 69362

15: 0

total: 1048576 (should be 1048576)

Histogram : GPU version

```
#include <stdio.h>

#include <stdlib.h>


#define GRIDSIZE    1024
#define BLOCKSIZE   1024
#define TOTALSIZE   (GRIDSIZE * BLOCKSIZE)
#define    NUMHIST           16


void genData(unsigned int* ptr, unsigned int size) {
    while (size--) {
        *ptr++ = (unsigned int)(rand() % (NUMHIST - 1));
    }
}

__global__ void kernel(unsigned int* hist, unsigned int* img, unsigned int size) {
    unsigned int i = blockIdx.x * blockDim.x + threadIdx.x;

    unsigned int pixelVal = img[i];
    atomicAdd(&(hist[pixelVal]), 1);
}
```

Histogram : GPU version

```
int main(void) {
    unsigned int* pImage = NULL;
    unsigned int* pHistogram = NULL;
    int i;

    // prepare timer
    cudaEvent_t start;
    cudaEvent_t stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);

    // malloc memories on the host-side
    pImage = (unsigned int*)malloc(TOTALSIZE * sizeof(unsigned int));
    pHistogram = (unsigned int*)malloc(NUMHIST * sizeof(unsigned int));

    // generate source data
    genData(pImage, TOTALSIZE);
}
```

Histogram : GPU version

// CUDA: allocate device memory

```
unsigned int* pImageDev;
```

```
unsigned int* pHistogramDev;
```

```
cudaMalloc((void**)&pImageDev, TOTALSIZE * sizeof(unsigned int));
```

```
cudaMalloc((void**)&pHistogramDev, NUMHIST * sizeof(unsigned int));
```

```
cudaMemset(pHistogramDev, 0, NUMHIST * sizeof(unsigned int));
```

// CUDA: copy from host to device

```
cudaMemcpy(pImageDev, pImage, TOTALSIZE * sizeof(unsigned int), cudaMemcpyHostToDevice);
```

// start the timer

```
cudaEventRecord(start, 0);
```

// perform the action

```
dim3 dimGrid(GRIDSIZE, 1, 1);
```

```
dim3 dimBlock(BLOCKSIZE, 1, 1);
```

```
kernel<<<dimGrid, dimBlock>>>(pHistogramDev, pImageDev, TOTALSIZE);
```


Histogram : GPU version

```
// end the timer
```

```
float time;
```

```
cudaEventRecord(stop, 0);
```

```
cudaEventSynchronize(stop);
```

```
cudaEventElapsedTime(&time, start, stop);
```

```
printf("elapsed time = %f msec\n", time);
```

```
cudaEventDestroy(start);
```

```
cudaEventDestroy(stop);
```

```
// CUDA: copy from device to host
```

```
cudaMemcpy(pHistogram, pHistogramDev, NUMHIST * sizeof(unsigned int),  
cudaMemcpyDeviceToHost);
```

Histogram : GPU version

```
// print the histogram
```

```
long total = 0L;
```

```
for (i = 0; i < NUMHIST; ++i) {
```

```
    printf("%2d: %10d\n", i, pHistogram[i]);
```

```
    total += pHistogram[i];
```

```
}
```

```
printf("total: %10ld (should be %ld)\n", total, TOTALSIZE);
```

```
// CUDA: free the memory
```

```
cudaFree(pImageDev);
```

```
cudaFree(pHistogramDev);
```

```
// free the memory
```

```
free(pImage);
```

```
free(pHistogram);
```

```
}
```

Execution Result

elapsed time = 1.612192 msec

0: 69692

1: 69634

2: 70121

3: 70277

4: 70215

5: 69479

6: 69988

7: 70344

8: 69984

9: 69976

10: 70099

11: 69686

12: 69810

13: 69909

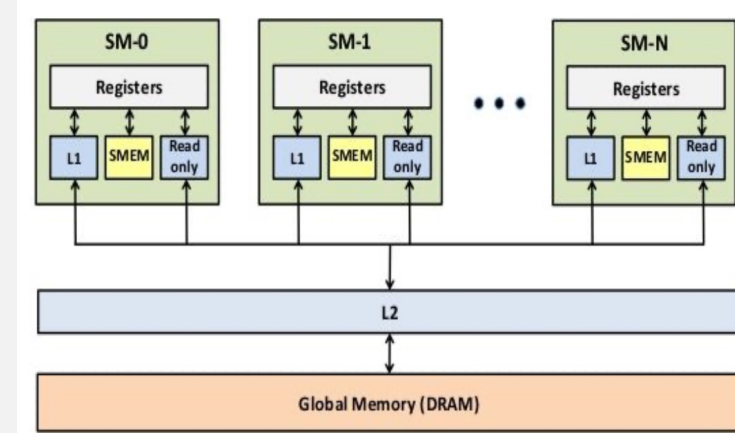
14: 69362

15: 0

total: 1048576 (should be 1048576)

Histogram : shared memory

```
__global__ void kernel(unsigned int* hist, unsigned int* img, unsigned int size) {  
    __shared__ int histShared[NUMHIST];  
    if (threadIdx.x < NUMHIST) {  
        histShared[threadIdx.x] = 0;  
    }  
    __syncthreads();  
    unsigned int i = blockIdx.x * blockDim.x + threadIdx.x;  
    unsigned int pixelVal = img[i];  
    atomicAdd(&(histShared[pixelVal]), 1);  
    __syncthreads();  
    if (threadIdx.x < NUMHIST) {  
        atomicAdd(&(hist[threadIdx.x]), histShared[threadIdx.x]);  
    }  
}
```



Next?

- **Parallel Patterns: Convolution**