



# **Parallel Patterns: Convolution**

Prof. Seokin Hong

# Convolution?



- A mathematical operation which describes a rule of how to **combine two functions or pieces of information** to **form a third function**.
- The convolution of  $f$  and  $g$  is written  $f * g$

$$(f * g)(t) \triangleq \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$




# Convolution Applications

- Convolution is a popular array operation that is used in various forms in signal processing, digital recording, image processing, video processing, and computer vision.
- Convolution is often performed as a **filter** that **transforms signals and pixels** into more **desirable values**.
  - Some filters smooth out the signal values so that one can see the trend
  - Gaussian filters can be used to sharpen boundaries and edges of objects in images.



Input image




Convolution Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



Feature map



Before




After



# Convolution Computation


- An **array operation** where each **output** data element is a **weighted sum** of a collection of **neighboring** input elements
- The **weights** used in the weighted sum calculation are defined by an **input mask array**, commonly referred to as the **convolution kernel** (*aka filter or mask*)

$$(f * g)(t) \triangleq \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau.$$



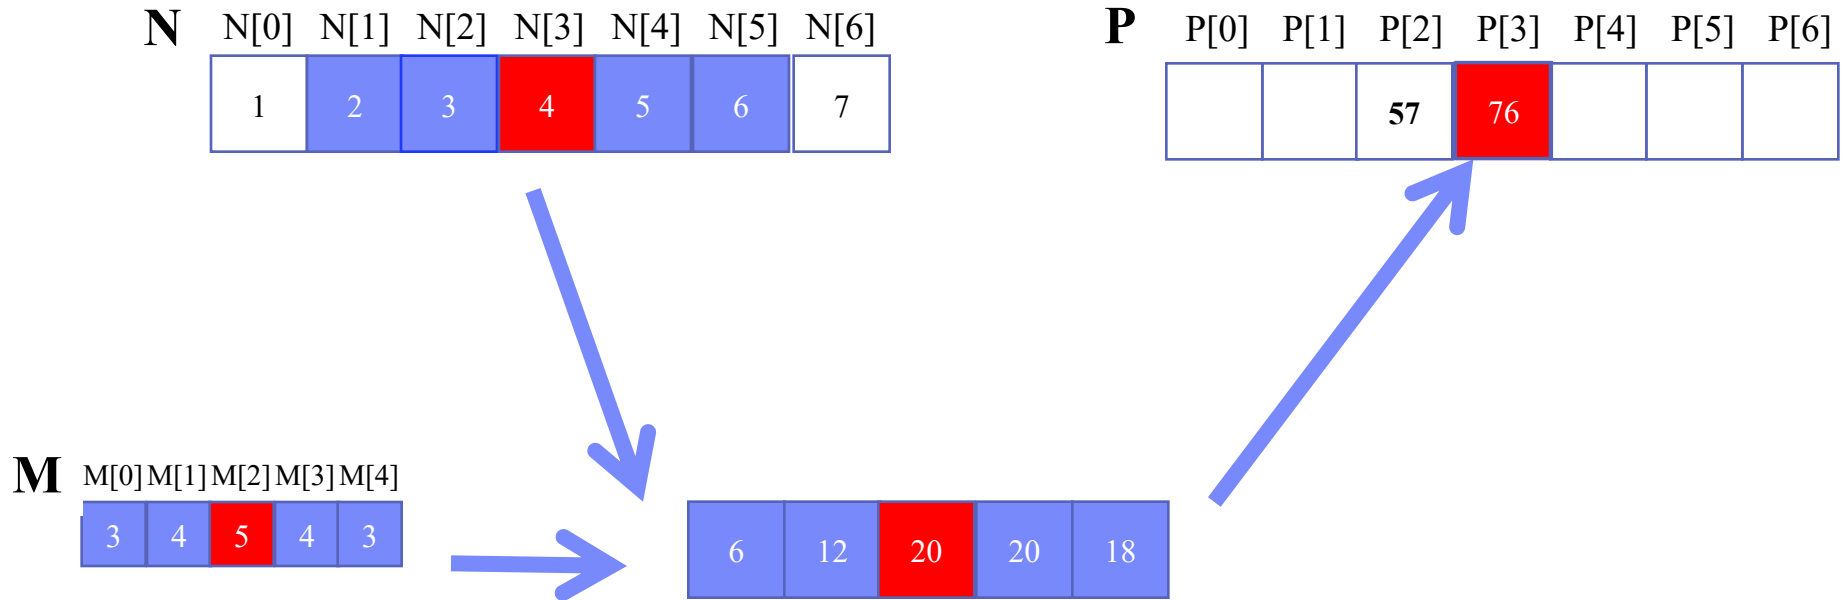
# 1D Convolution Example

- Commonly used for audio processing
  - Mask size is usually an **odd** number of elements for symmetry (5 in this example)
- Calculation of  $P[2]$



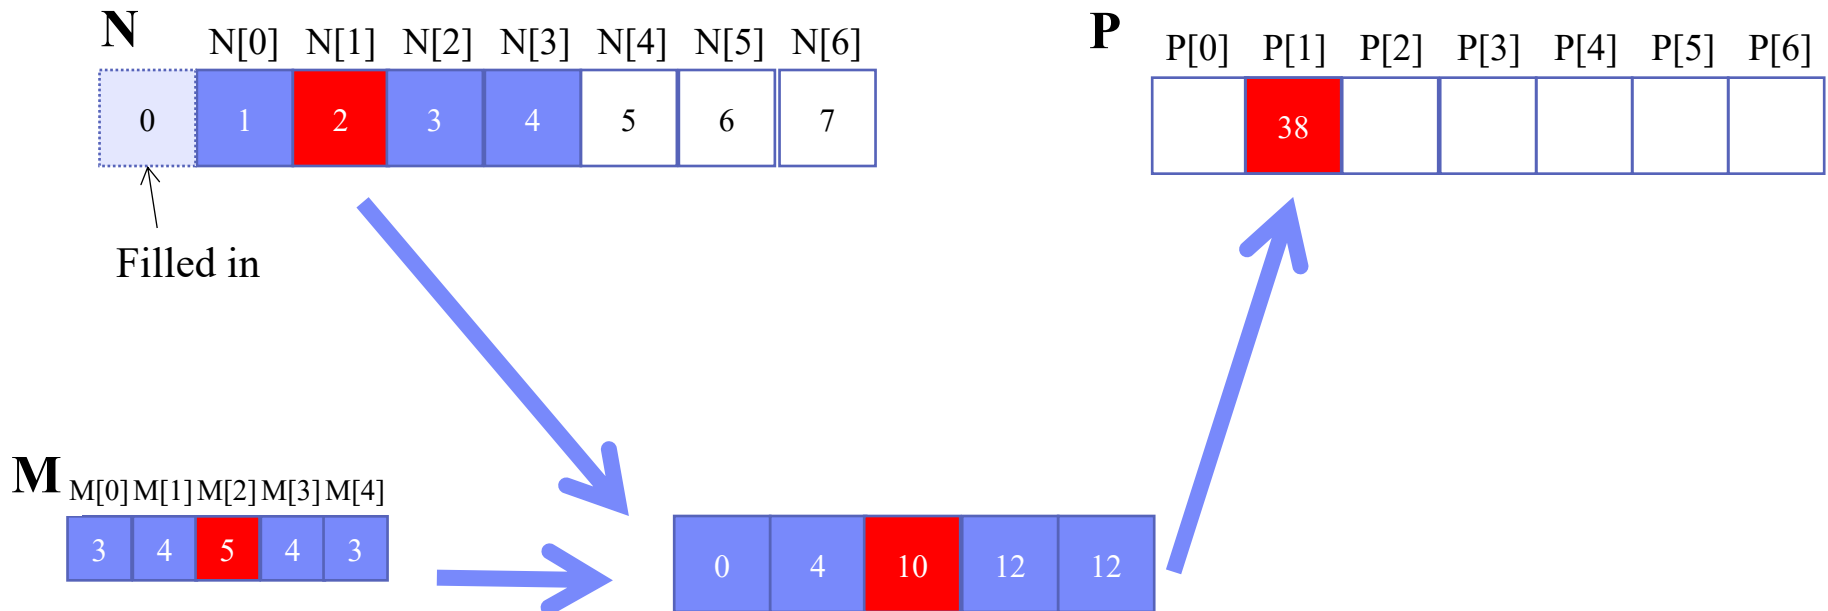
# 1D Convolution Example - more on inside elements

- Calculation of  $P[3]$



# 1D Convolution with Boundary Condition Handling

- Calculation of **output elements near the boundaries** (beginning and end) of the input array need to **deal with “ghost” elements**
  - Different policies (0, replicates of boundary values, etc.)



# 1D Convolution with Boundary Condition Handling

## ■ 1D Convolution CUDA Kernel

- This kernel forces all elements outside the valid data index range to 0


```
__global__ void convolution_1D_basic_kernel(float *N, float *M, float *P, int Mask_Width, int Width) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;

    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);

    for (int j = 0; j < Mask_Width; j++) {
        if (N_start_point + j >= 0 && N_start_point + j < Width) {
            Pvalue += N[N_start_point + j]*M[j];
        }
    }
    P[i] = Pvalue;

}
```





# 2D Convolution

**N**

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

**P**

		321				


**M**

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1




1	4	9	8	5
4	9	16	15	12
9	16	25	24	21
8	15	24	21	16
5	12	21	16	5

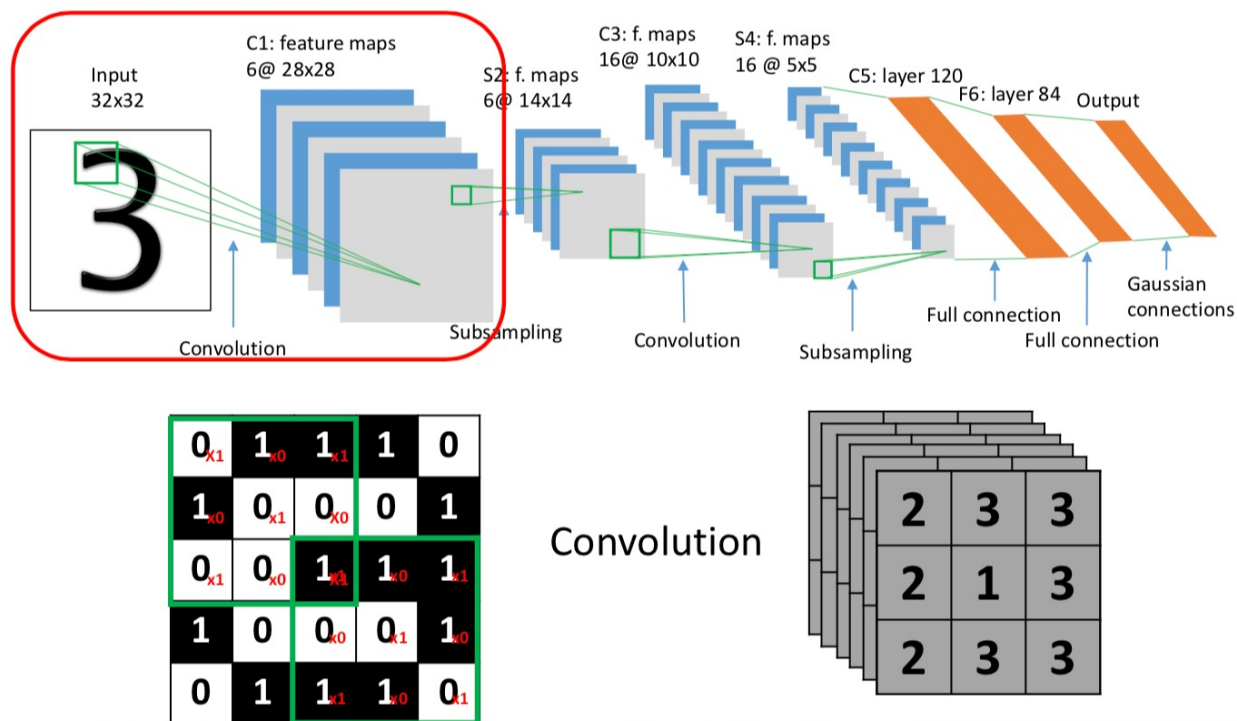
# 2D Convolution with Boundary Condition Handling



# 2D Convolution – Ghost Cells

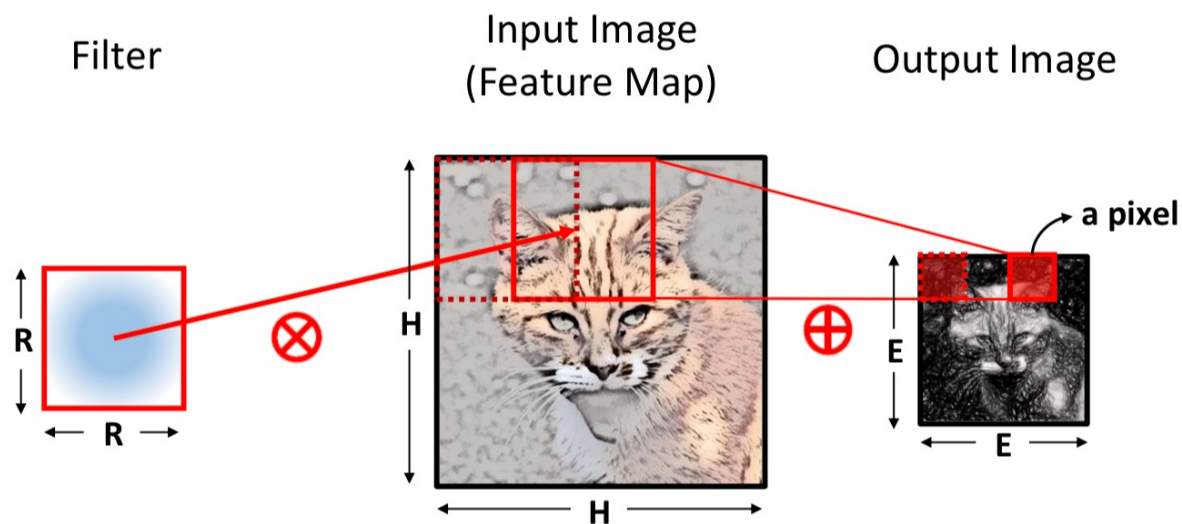


# Convolution Neural Network (CNN)



# Convolution Neural Network (CNN)

- Convolution



**Sliding Window Processing**

Eyeriss: A Spatial Architecture for Energy-Efficient Dataflow for Convolutional Neural Networks (ISCA 2016)

# Convolution Neural Network (CNN)

## ■ Convolution Computation

```
for(r=0; r<R; r++)    //output feature map
  for(q=0; q<Q; q++)    //input feature map
    for(m=0; m<M; m++)    //row in feature map
      for(n=0; n<N; n++)    //column in feature map
        for(k=0; k<K; k++)    //row in convolution kernel
          for(l=0; l<L; l++)    //column in convolution kernel
            Y[r][m][n] += W[r][q][k][l] * X[q][m+k][n+l];
```


$$Y = Y + W * X$$

**“Convolution” account for over 90% of the CNN operations and dominates runtime**

# Access Pattern for M

- **M** is referred to as mask (a.k.a. kernel, filter, etc.)
  - Elements of M are called mask (kernel, filter) coefficients
- Calculation of all output P elements needs **M**
- **M** is not changed during a execution of CUDA kernel that performs the convolution
- In addition, **M** elements are accessed in the same order when calculating all P elements
- M is a good candidate for **Constant Memory**

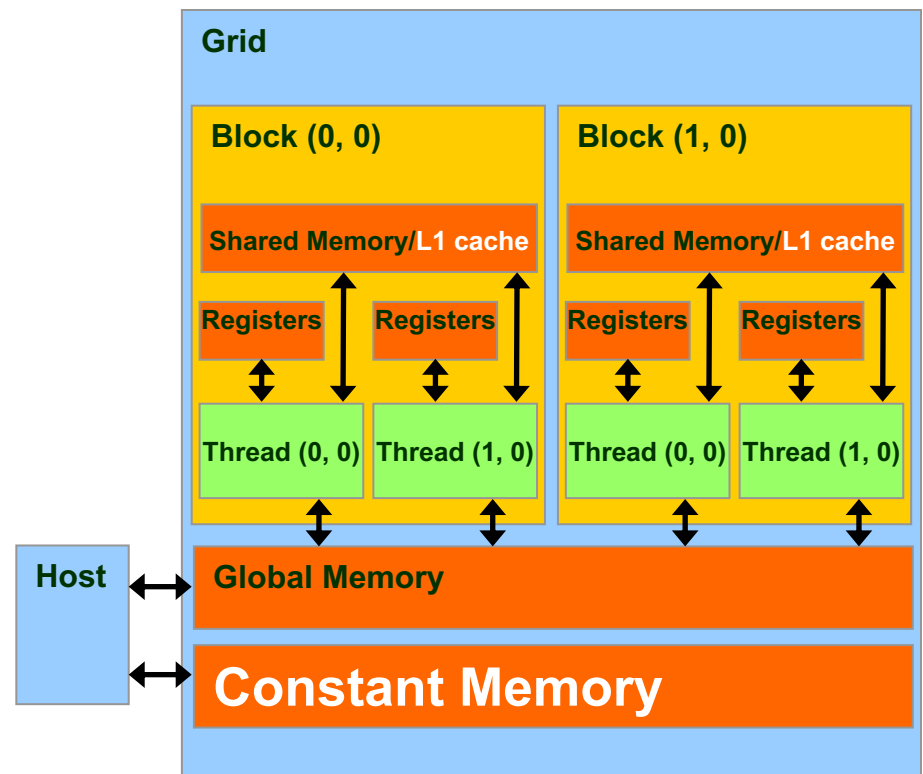
1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1



Sliding Window Processing

# Programmer View of CUDA Memories (Review)


- Each thread can:
  - Read/write per-thread **registers** (~1 cycle)
  - Read/write per-block **shared memory** (~5 cycles)
  - Read/write per-grid **global memory** (~500 cycles)
  - Read/only per-grid **constant memory** (~5 cycles with caching)



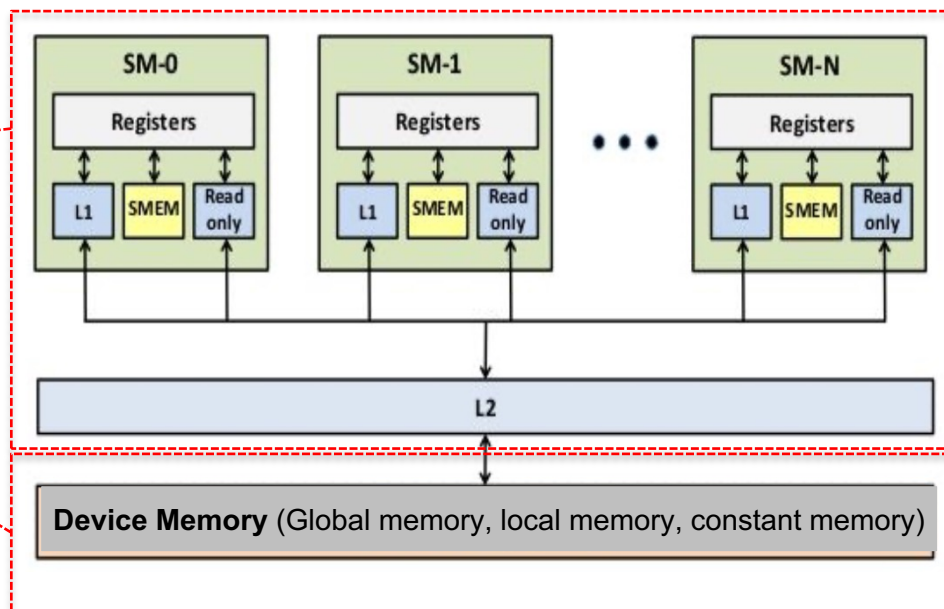


# Memory Hierarchies

- If we had to go to global memory to access data all the time, the execution speed of GPUs would be limited by the global memory bandwidth
- One solution: Caches



Green box is GK110, red lines are global memory



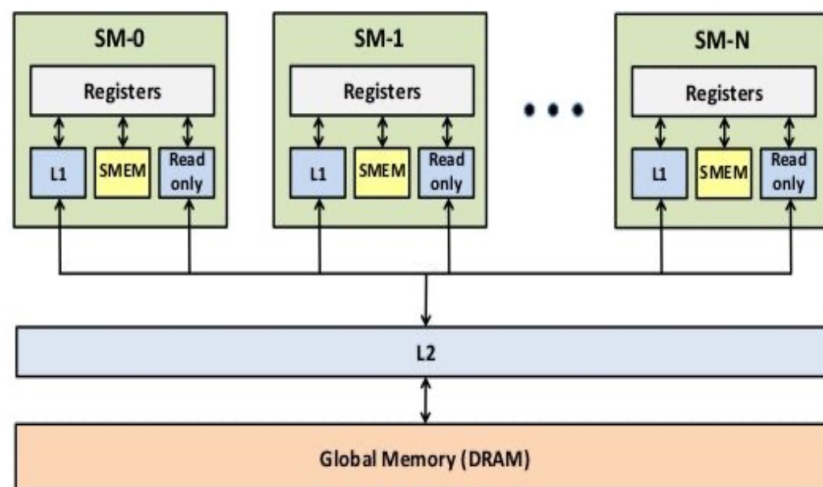
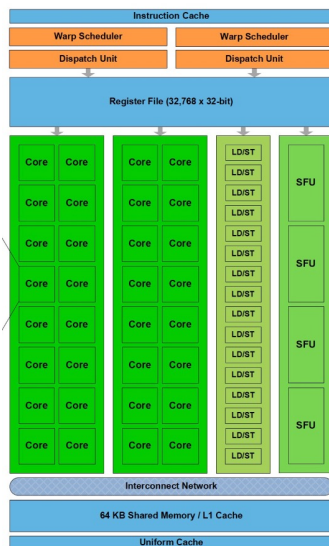
# Cache (review)

---

- A cache is an “array” of cache lines
  - A cache line can usually hold data from several consecutive memory addresses
- When data is requested from the global memory, **an entire cache line that includes the data being accessed is loaded into the cache**, in an attempt to reduce global memory requests
  - The data in the cache is a “copy” of the original data in global memory
- **Locality**
  - **Spatial locality**: when the data elements stored in consecutive memory locations are access consecutively
  - **Temporal locality**: when the same data element is access multiple times in short period of time
- Both spatial locality and temporal locality improve the performance of caches

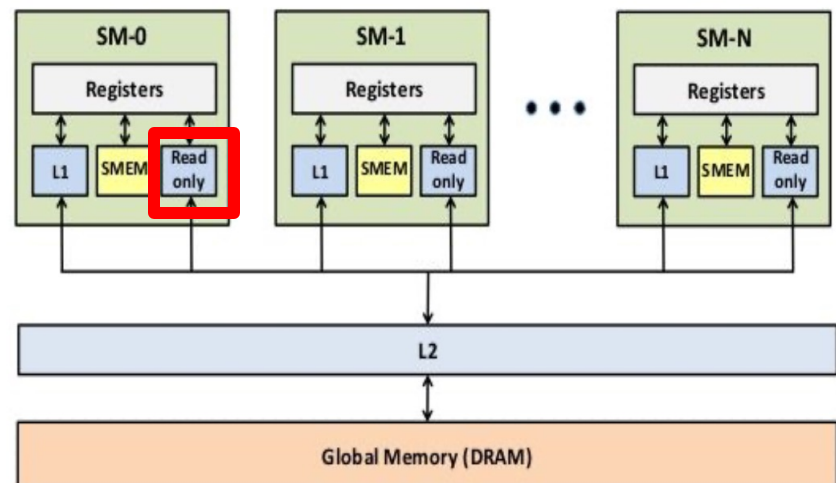
# Scratchpad vs. Cache

- **Scratchpad** (**shared memory** in CUDA) is another type of temporary storage used to relieve main memory contention.
  - In terms of distance from the processor, scratchpad is similar to L1 cache.
- **Unlike cache**, **scratchpad does not necessarily hold a copy of data** that is also in main memory
- **Scratchpad requires explicit data transfer instructions** into locations in the scratchpad, whereas cache doesn't



# Constant Cache in GPUs

- Modification to cached data needs to be (eventually) reflected back to the original data in global memory
  - Requires logic to track the modified status, etc.
- **Constant cache** is a **special cache for constant data** that will not be modified during kernel execution
  - Data declared in the constant memory will not be modified during kernel execution.



# How to Use Constant Memory

---

- Host code allocates, initializes variables the same way as any other variables that need to be copied to the device
- Use `cudaMemcpyToSymbol(dest, src, size)` to copy the variable into the device memory
- This copy function tells the device that `the variable will not be modified by the kernel and can be safely cached.`

# Some Header File Stuff for M

---

```
#define MASK_WIDTH 5
```

```
// Matrix Structure declaration
```

```
typedef struct {
```

```
    unsigned int width;
```

```
    unsigned int height;
```

```
    unsigned int pitch;
```

```
    float* elements;
```

```
} Matrix;
```

# AllocateMatrix()

---

```
// Allocate a device matrix of dimensions height*width
//           If init == 0, initialize to all zeroes.
//           If init == 1, perform random initialization.
// If init == 2, initialize matrix parameters, but do not allocate memory
```

**Matrix AllocateMatrix(int height, int width, int init)**

```
{
    Matrix M;
    M.width = M.pitch = width;
    M.height = height;
    int size = M.width * M.height;
    M.elements = NULL;
    // don't allocate memory on option 2
    if(init == 2) return M;
    int size = height * width;
    M.elements = (float*) malloc(size*sizeof(float));
    for(unsigned int i = 0; i < M.height * M.width; i++)
    {
        M.elements[i] = (init == 0) ? (0.0f) :
                                (rand() / (float)RAND_MAX);
        if(rand() % 2) M.elements[i] = - M.elements[i]
    }
    return M;
}
```

# Host Code

---

```
// global variable, outside any kernel/function
```

```
__constant__ float Mc[MASK_WIDTH][MASK_WIDTH];
```

```
...
```

```
// allocate N, P, initialize N elements, copy N to Nd
```

```
Matrix M;
```

```
M = AllocateMatrix(MASK_WIDTH, MASK_WIDTH, 1);
```

```
// initialize M elements
```

```
....
```

```
cudaMemcpyToSymbol(Mc, M.elements,
```

```
    MASK_WIDTH*MASK_WIDTH*sizeof(float));
```

```
ConvolutionKernel<<<dimGrid, dimBlock>>>(Nd, Pd);
```



# Next?

---

- **Tiled Convolution**