
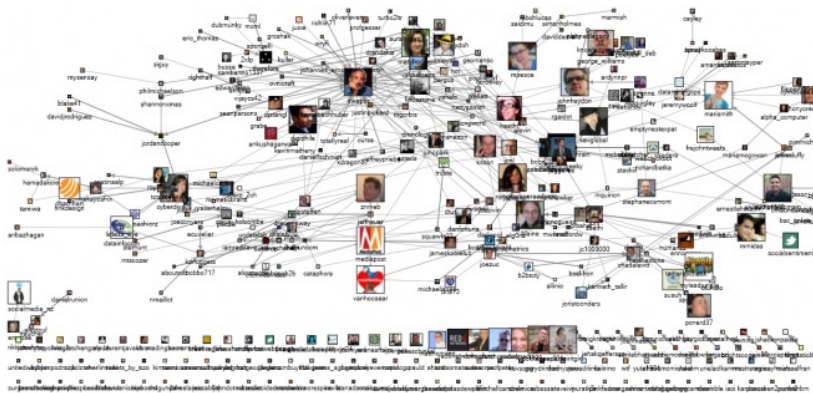


Parallel Patterns: Graph Search

Prof. Seokin Hong

Graph

- A **graph data structure** represents the **relations** between **entries**
 - Example
 - **Social network:** **entities** are **users** and **relations** are **connections** between users
 - **Driving direction app:** **entities** are **locations** and the **relations** are the **roadways** between them
- Relations are bi-directional or directional




Graph with directional edges

Representation of a Graph

- **Adjacency matrix** can be used to represent a graph
 - Assign an unique number to each vertex
 - When there is an edge going from vertex i to vertex j , the value of element $A[i][j]$ is 1. Otherwise, it is 0.

	0	1	2	3	4	5	6	7	8
0		1	1						
1				1	1				
2						1	1	1	
3					1				1
4						1			1
5							1		
6									1
7	1						1		
8									

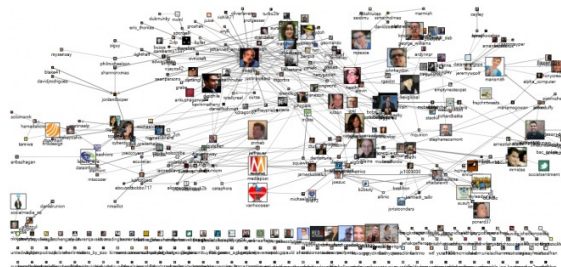


Graph with directional edges

Sparsely-connected Graph


- Usually, graph is **sparsely** connected

- If the number of vertex is N , the average number of outgoing edges from each vertex is **much smaller** than $N-1$
- Many real-world graphs are sparsely connected
 - Ex) a social network such as Facebook



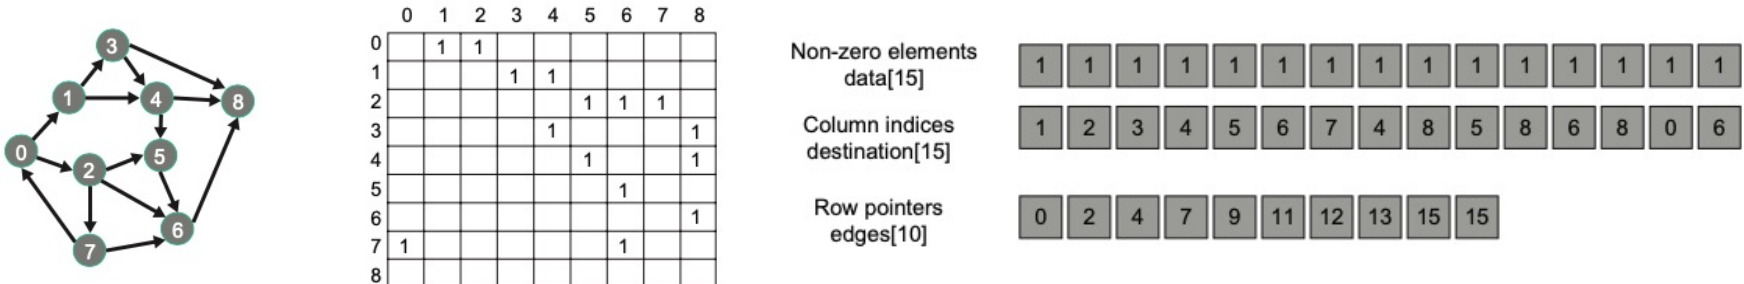
- So, the number of **non-zero elements in the adjacency matrix** is **much smaller** than the total number of elements
- **Using sparse matrix representation** (e.g., CSR) to represent a sparsely-connected graph drastically **reduce the amount of storage** and **the number of wasted operations**

	0	1	2	3	4	5	6	7	8
0		1	1						
1				1	1				
2						1	1	1	
3					1				1
4						1			1
5							1		
6									1
7	1							1	
8									




Sparse matrix representation of sparsely-connected Graph

- **Row pointer** array gives starting location of non-zero elements (**edges**)
 - Example
 - `edges[3]` gives the starting location (=7) of non-zero elements in row 3 (vertex 3)
 - `edges[4]` gives the starting location (=9) of non-zero elements in row 4 (vertex 4)
 - So, we expect to find non-zero elements for row 3 in `data[7]` and `data[8]` and the column indices for these elements in `destination[7]` and `destination[8]`
- **Column index** array gives **destination** of each edge
 - Ex) the destination of the two edges for source vertex 3 are `destination[7]=4` and `destination[8]=8`
- **Data array** can be used to store additional information about the relationship
 - Ex) distance between two locations or date when two social network users became connected




Breadth-First Search (BFS)

- **BFS** is a graph traversing algorithm that **traverses a graph layerwise** thus exploring the neighbor nodes (nodes which are directly connected to source node).
- BFS is often **used to discover the shortest path** between two vertices
 - **Each vertex is labeled** with the **smallest number of edges that one needs to traverse** in order to go **from the source** to the vertex



Desired BFS result with vertex 0 as the source




Desired BFS result with vertex 0 as the source


Breadth-first Search (BFS)

- Example of labeling vertices with BFS : source is vertex 0


Step0




Step1



Step2




Step3



Breadth-first Search (BFS)

■ Finding shortest path

- Start from destination vertex and trace back to the source
- Select a vertex that has smallest number of edges



Sequential BFS function

- The graph is **represented in the CSR format**
- **The function receives**
 - the index of the **source vertex**
 - **edges array** (row pointer array of CSR)
 - **destination array** (column index array of CSR)
 - **label array** whose elements are used to **store the visit status information** for the vertices
- **Before search,**
 - **Source vertex is labeled to 0**, indicating it is a level 0 vertex
 - All other label elements are initialized to -1, indicating that they are not visited
- **At the end of the search,**
 - all vertices **reachable** from the source are labeled with a **positive number**
 - **If the label of a vertex is -1**, it means that the vertex is **unreachable** from the source

Implementation of Sequential BFS function

```
void BFS_sequential(int source, int *edges, int *dest, int *label)
{
    int frontier[2][MAX_FRONTIER_SIZE]; //two frontier arrays
    int *c_frontier = &frontier[0];    /*c_frontier array stores the frontier vertices that are
                                         being discovered in the current iteration */
    int c_frontier_tail = 0;             /*stores the index of position at which a newly discovered
                                         frontier vertex can be added in c_frontier array*/
    int * p_frontier = &frontier[1]; /*p_frontier array stores the frontier vertices discovered in the
                                         previous iteration*/
    int p_frontier_tail = 0;             //number of elements that have been inserted into the p_frontier

    insert_frontier(source, p_frontier, &p_frontier_tail); //place source vertex into p_frontier and
                                                         increment p_frontier_tail to 1

    label[source] = 0; //label of a source vertex is initialized to 0

    .....
```

A frontier : vertices that are not yet visited but are adjacent to a visited vertex and visit a vertex only when it is in the frontier

Implementation of Sequential BFS function

.....


```

while (p_frontier_tail > 0) {
    for(int f=0; f<p_frontier_tail; f++){
        c_vertex= p_frontier[f];
        for (int i= edges[c_vertex]; i<edges[c_vertex+1];i++) { //for all its edges
            if(label[dest[i]] == -1){ //the dest vertex has not been visited
                insert_frontier(dest[i], c_frontier, &c_frontier_tail); //insert discovered vertex into c_frontier
                label[dest[i]] = label[c_vertex]+1;
            }
        }
    }

    int temp= c_frontier; c_frontier = p_frontier; p_frontier = temp; //swap previous and current frontier
    p_frontier_tail = c_frontier_tail;
    c_frontier_tail = 0;
}
    
```

Non-zero elements data[15]	1	1	1	1	1	1	1	1	1	1	1	1	1	1	
Column indices destination[15]	1	2	3	4	5	6	7	4	8	5	8	6	8	0	6
Row pointers edges[10]	0	2	4	7	9	11	12	13	15	15					

Step0



label



c_frontier



p_frontier



Parallel BFS Function

- Parallelize each iteration of the while-loop
 - Multiple threads collaboratively process the previous frontier array and assemble the current frontier array
 - How?
 - Assign a section of the previous frontier array to each thread block

```
while (p_frontier_tail > 0) {  
    for(int f=0; f<p_frontier_tail; f++){  
        c_vertex= p_frontier[f];  
        for (int i= edges[c_vertex]; i<edges[c_vertex+1];i++) {  
            if(label[dest[i]] == -1){  
                insert_frontier(dest[i], c_frontier & c_frontier_tail);  
            }  
        }  
    }  
}
```

//visit all previous frontier vertices

//pick up one of the previous frontier vertex

//for all its edges

//the dest vertex has not been visited

//insert discovered vertex into c_frontier

Parallelize

Implementation of Parallel BFS Function

■ Host code

```
void BFS_host(unsigned int source, unsigned int*edge, unsigned int *dest, unsigned int *label)
{
    //allocate edges_d, dest_d, label_d, and visited_d in device global memory
    //copy edges, dest, and label to device global memory
    // allocate frontier_d, c_frontier_tail_d, p_frontier_tail_d in device global memory
    .....
    unsigned int *c_frontier_d = &frontier_d[0];           //c_frontier_d point the first half
    unsigned int *p_frontier_d = &frontier_d[MAX_FRONTIER_SIZE]; //c_frontier_d point the second half

    //launch a simple kernel to initialize the following in the device global memory
    //initialize all visited_d elements to 0 except source to 1
    /*c_frontier_tail_d = 0;
    //p_frontier_d[0] = source;
    /*p_frontier_tail_d = 1;
    //label_d[source] =0;
    p_frontier_tail = 1;
    ....
```

Implementation of Parallel BFS Function

■ Host code (cont'd)

....

```
while(p_frontier_tail>0) {
```

```
    int num_blocks = ceil(p_frontier_tail/float(BLOCK_SIZE));
```

```
    BFS_Bqueue_kernel<<num_blocks, BLOCK_SIZE>>>(p_frontier_d, p_frontier_tail_d,  
c_frontier_d, c_frontier_tail_d, edges_d, dest_d, label_d, visited_d);
```

```
//use cudaMemcpy to read the *c_frontier_tail value back to host
```

```
// and assign it to p_frontier_tail for the while-loop condition test
```

```
    int* temp = c_frontier_d; c_frontier_d=p_frontier_d; p_frontier_d=temp; //swap the roles
```

```
// launch a simple kernel to set *p_frontier_tail_d = *c_frontier_tail_d; *c_frontier_tail_d=0;
```

```
}
```

```
}
```

Implementation of Parallel BFS function

■ Kernel

```
__global__ void BFS_Bqueue_kernel(unsigned int *p_frontier, unsigned int *p_frontier_tail,
unsigned int *c_frontier, unsigned int *c_frontier_tail, unsigned int *edges, unsigned int *dest,
unsigned int *label, unsigned int* visited)
{
    __shared__ unsigned int c_frontier_s[BLOCK_QUEUE_SIZE];
    __shared__ unsigned int c_frontier_tail_s, our_c_frontier_tail;



    if(threadIdx.x == 0) c_frontier_tail_s = 0;
    __syncthreads();

    const unsigned int tid = blockIdx.x*blockDim.x +threadIdx.x;
    .....
```

Implementation of Parallel BFS function

Kernel (cont'd)

```
....
if (tid < *p_frontier_tail){
    const unsigned int my_vertex = p_frontier[tid];
    for(unsigned int i = edges[my_vertex]; i < edges[my_vertex+1] ; i++)
    {
        const unsigned int was_visited = atomicExch(&(visited[dest[i]],1);
        if(!was_visited){
            label[dest[i]] = label[my_vertex]+1;
            const unsigned int my_tail = atomicAdd(&c_frontier_tail_s,1);
            if(my_tail < BLOCK_QUEUE_SIZE) { // add discovered vertex to the local queue (shared memory)
                c_frontier_s[my_tail]= dest[i];
            } else { //if full, add it to the global queue directly (global memory)
                c_frontier_tail_s =BLOCK_QUEUE_SIZE;
                const unsigned int my_global_tail = atomicAdd(c_frontier_tail,1);
                c_frontier[my_global_tail] = dest[i];
            }
        }
    }
}
....
```



Implementation of Parallel BFS function

■ Kernel (cont'd)

```
....  
__syncthreads();  
    if(threadIdx.x ==0){  
        our_c_frontier_tail = atomicAdd(c_frontier_tail, c_frontier_tail_s);  
    }  
__syncthreads();  
for(unsigned int i= threadIdx.x; i<c_frontier_tail_s; i+=blockDim.x){  
    c_frontier[our_c_frontier_tail +i]=c_frontier_s[i];  
}  
}
```

