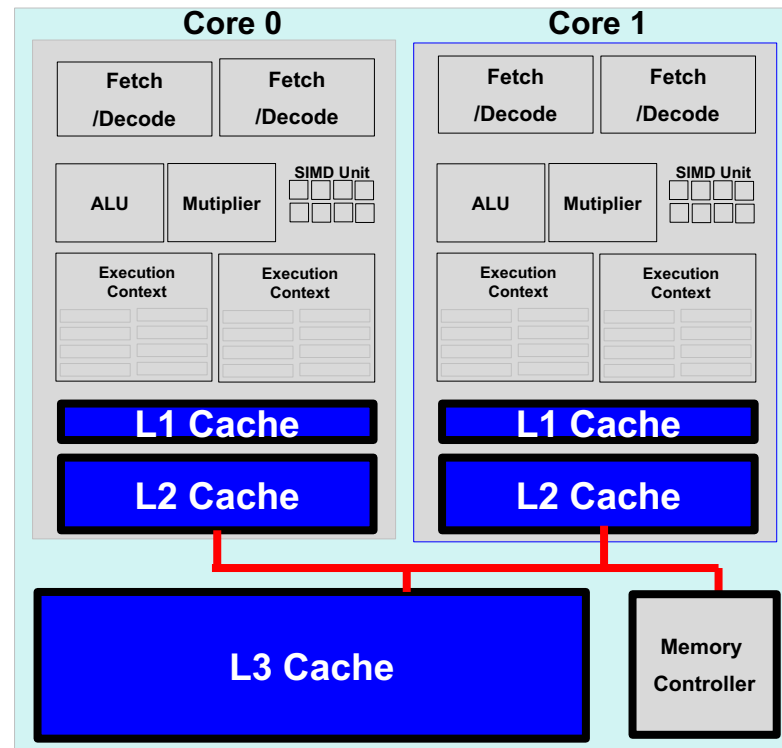


Cache Coherence

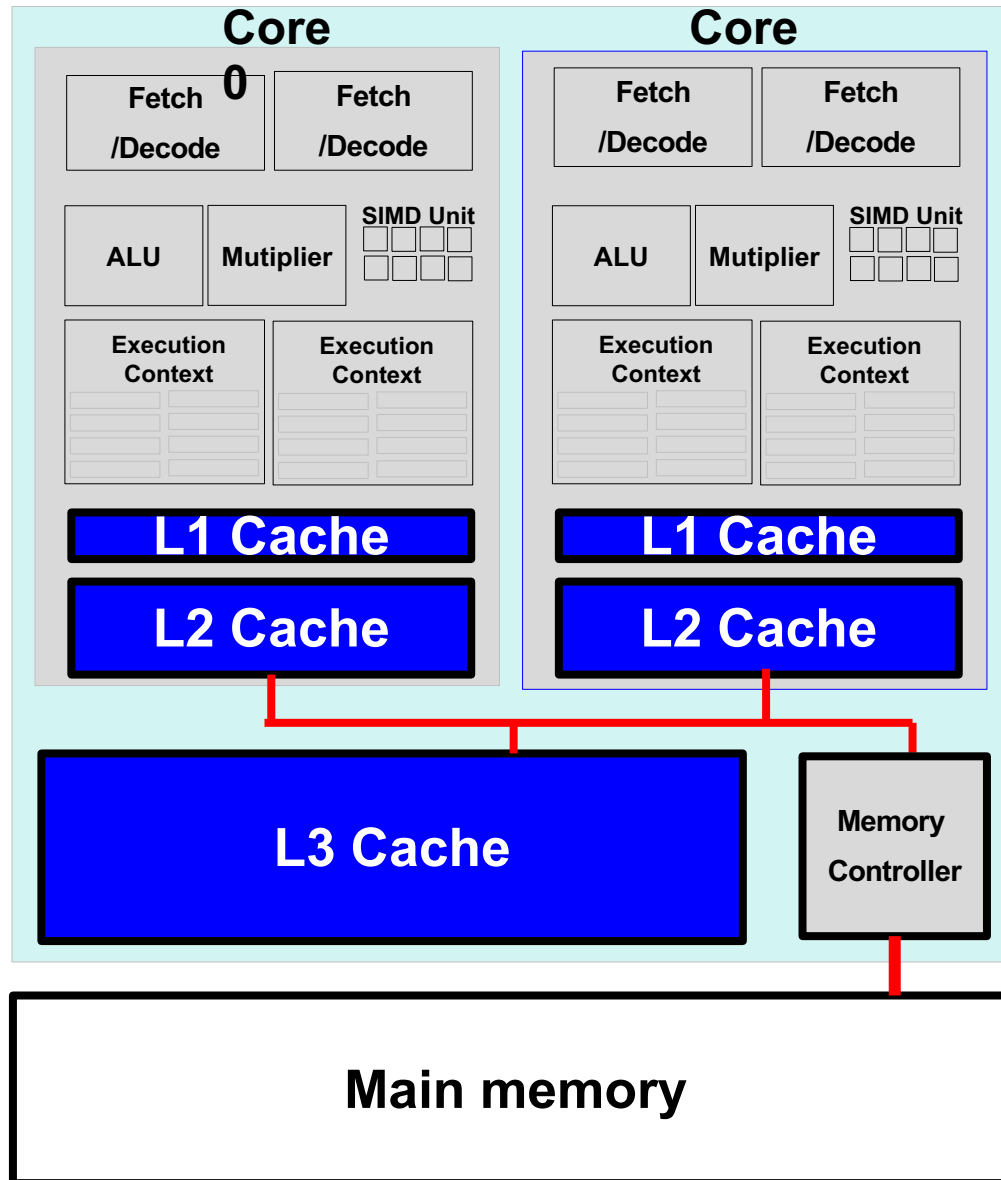
Prof. Seokin Hong

Next ?

- **Cores share memory (shared-memory multiprocessor model)**
 - Cache Coherence and Memory Consistency
- **Cores are connected each other through intra-chip network**
 - Interconnection Network




Shared and Private Memories



Cache Coherence

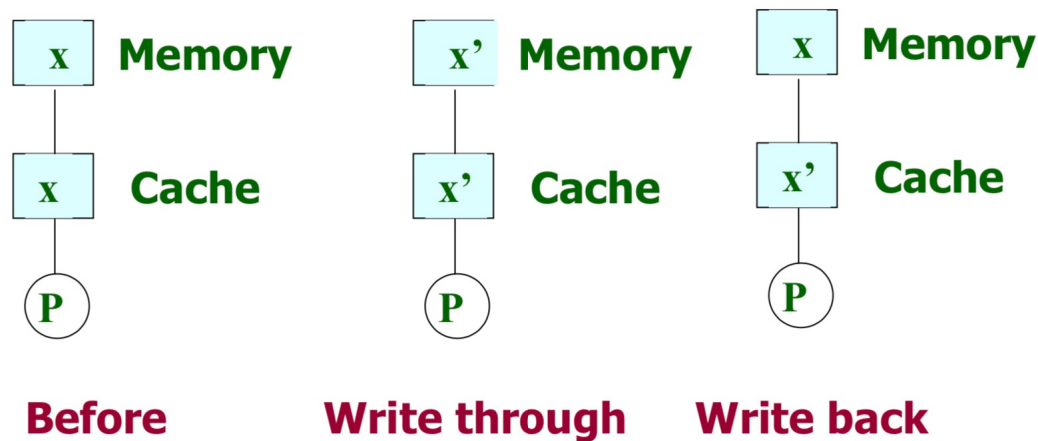
- A problem might happen in a multiprocessor system



What happen if processor 1 updates the shared data!!

Cache Coherence (Cont'd)

- Write-Through vs Write-Back (quick review)
 - **Write-Through policy**
 - The memory is updated every time the cache is updated
 - **Write-back policy**
 - The memory is updated only when the cache block in the cache is being replaced.




Cache Coherence (Cont'd)

- Memory System is Coherent if..
 - **P** writes **X**, **P** reads **X**
 - X returns value written by **P**
 - **Q** writes **X**, <sufficient time period>, **P** reads **X**
 - X returns value written by **Q**
 - **Writes are serialized**
 - Written values seen in the same order by all processors
- **Coherence** defines **what** values can be returned by a read
- **Consistency** defines **when** a written value is returned by a read
 - If processor **P** writes to location **A** and then to **B**, then any processor **Q** that sees the new value in **B** must also see the new value in **A**

Cache Coherence (Cont'd)

- Coherent Caches provide:
 - **Migration**: movement of data
 - **Replication**: multiple copies of data



- **Cache coherence protocols**

- **Directory-based**

- Shared status of each cache line (a.k.a. cache block) kept in one location
 - Single directory for small core counts, or distributed directories for many cores



- **Snooping-based**

- Each core tracks the **state** of each cache line
 - Watch bus activity and carry out the appropriate coherency commands when necessary
 - Requires broadcast link (e.g., data bus)




Snooping Cache Protocols

- **Write-update Protocol (write broadcast)**
 - Upon write, make sure all copies are the same
 - Requires more bandwidth
- **Write-invalidate protocol (2-state invalidate protocol)**
 - Upon write, make sure there is only one valid copy
 - Most common protocol

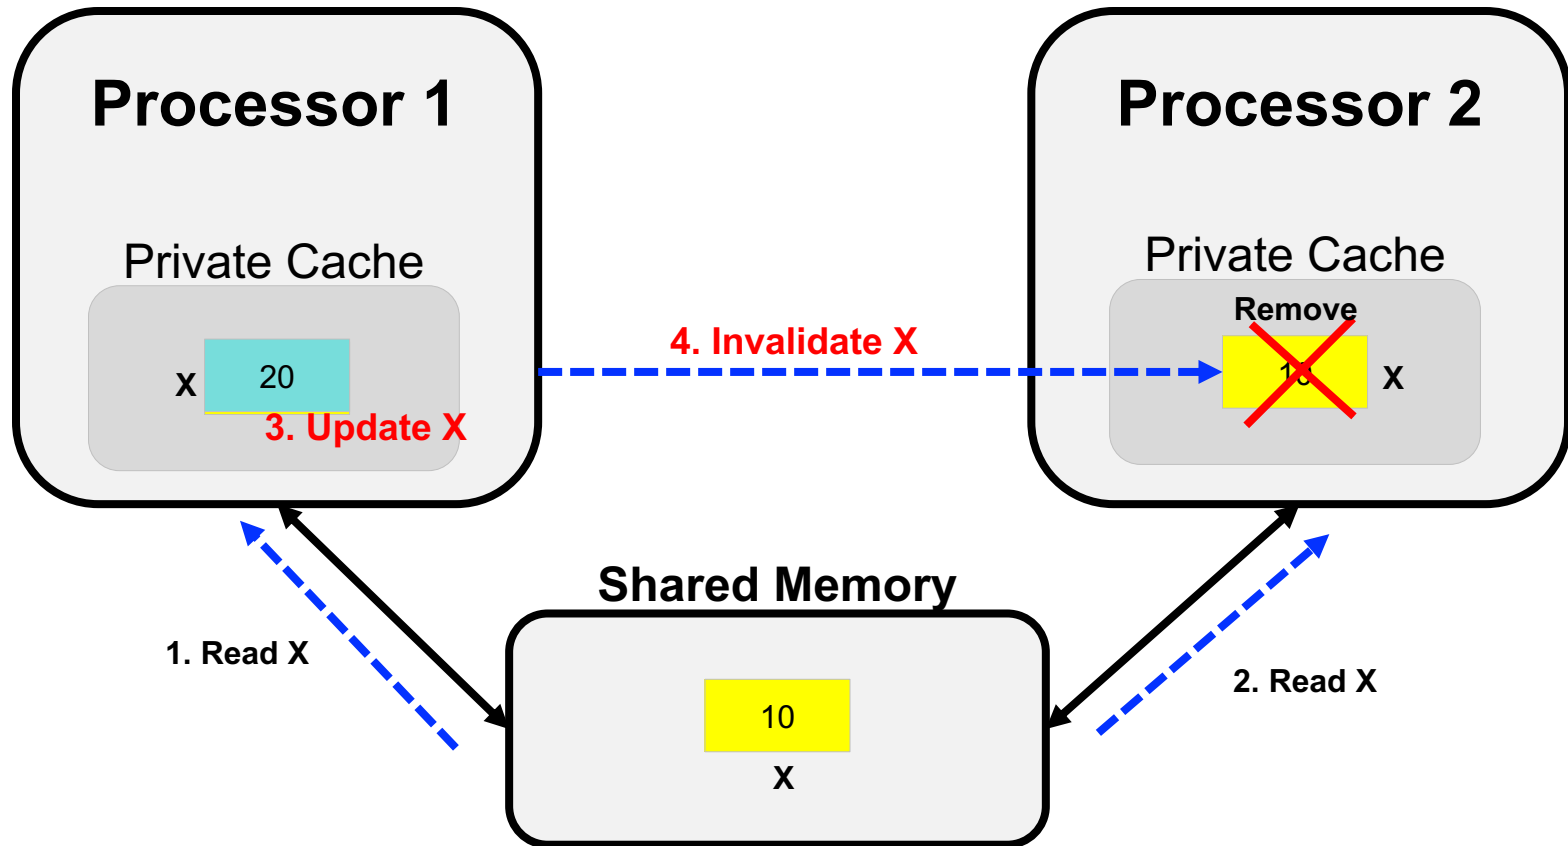
Snooping Cache Protocols (Cont'd)

▪ Write-update Protocol



Snooping Cache Protocols (Cont'd)

▪ Write-invalidate Protocol




Snooping Cache Protocols (Cont'd)

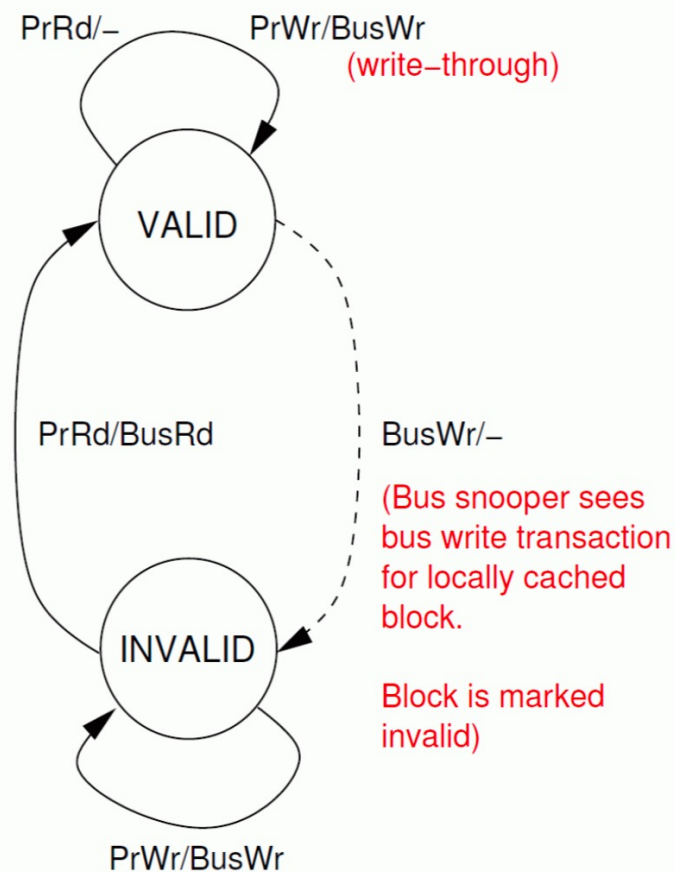
■ Write-invalidate Protocol (Cont'd)

= 2-state invalidate protocol

- **Each cache controller receives**
 - Memory requests from a local processor
 - memory **requests** from bus (snooping)
- **In response, the controller may**
 - update the state of a particular cache line
 - initiate a new bus transaction



State transition diagram for a cache line



PrRd : Read request from processor
PrWr : Write request from processor
BusRd: Read request from bus (other processors)
BusWr: Write request from bus (other processors)

Snooping Cache Protocols (Cont)

■ MSI (Modified-Shared-Invalid) Protocol

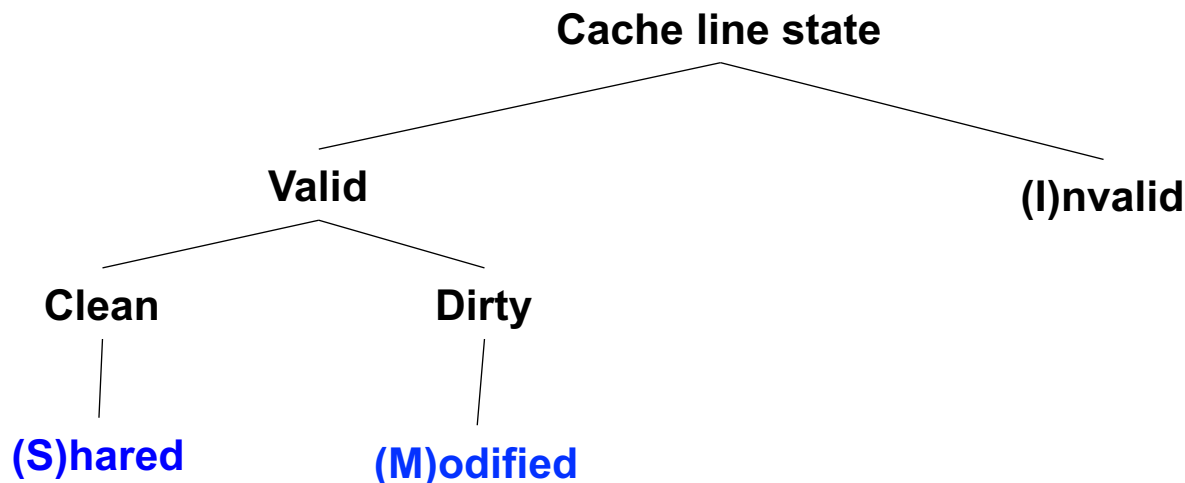

○ Cache line state

• (S)hared

- › The block is present in an unmodified state in this cache
- › The block is up-to-date in main memory
- › Zero or more up-to-date copies exist in other caches

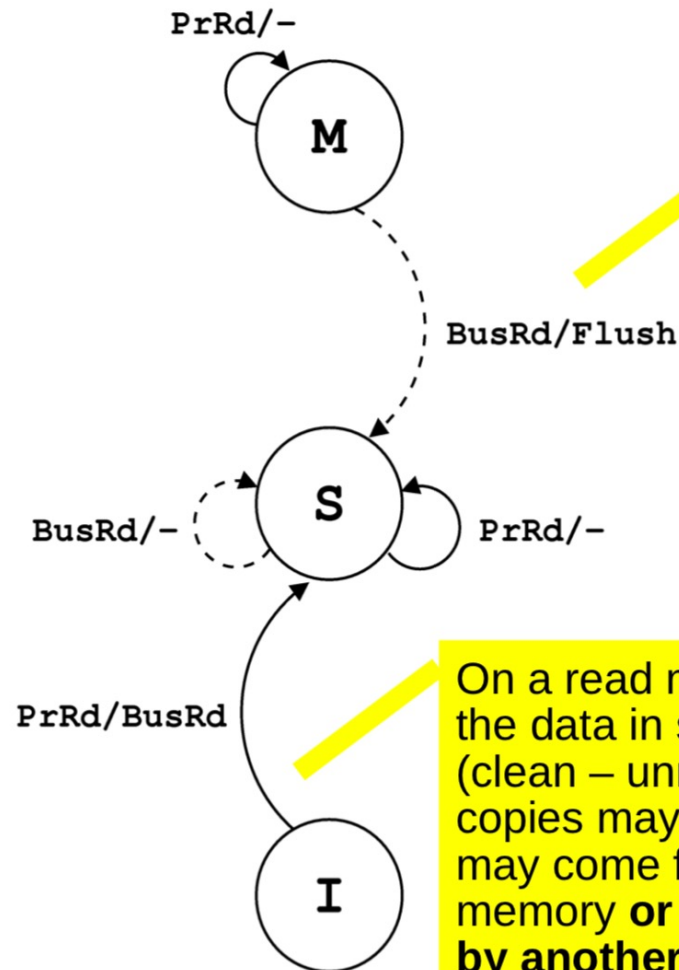

• (M)odified

- › Only this cache has a valid copy, the copy in memory and other cache is stale



Snooping Cache Protocols (Cont)

■ MSI (Modified-Shared-Invalid) Protocol




If another cache needs the block that we have in state M (the block is dirty/modified), we must flush the block to memory and provide the data to the requesting cache (over the bus)

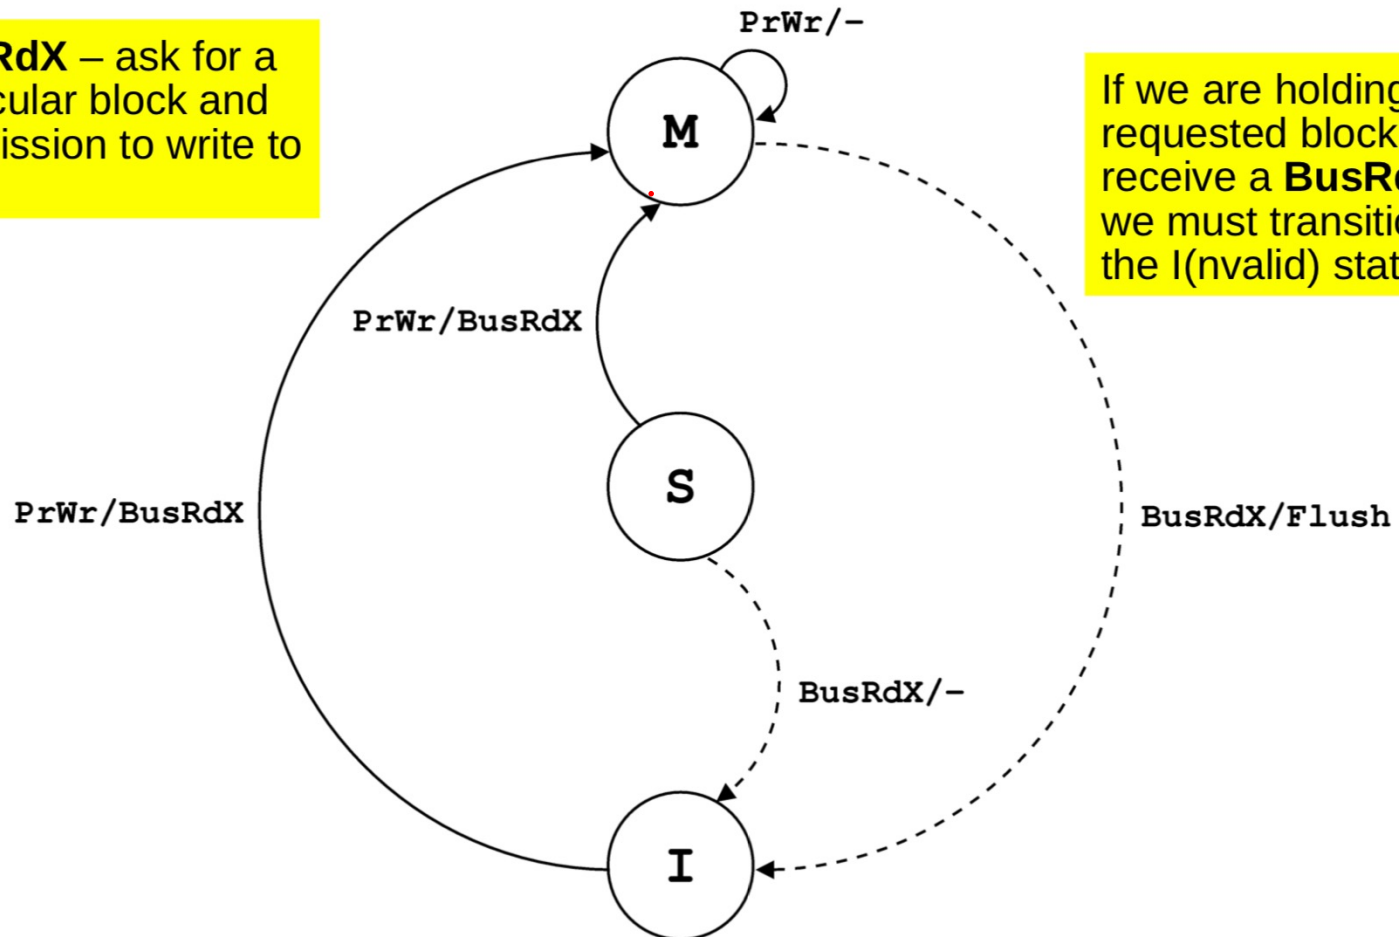
On a read miss, we get the data in state S (clean – unmodified, other copies may exist). Data may come from main memory **or be provided by another cache**

Snooping Cache Protocols (Cont)

■ MSI (Modified-Shared-Invalid) Protocol



BusRdX – ask for a particular block and permission to write to it.


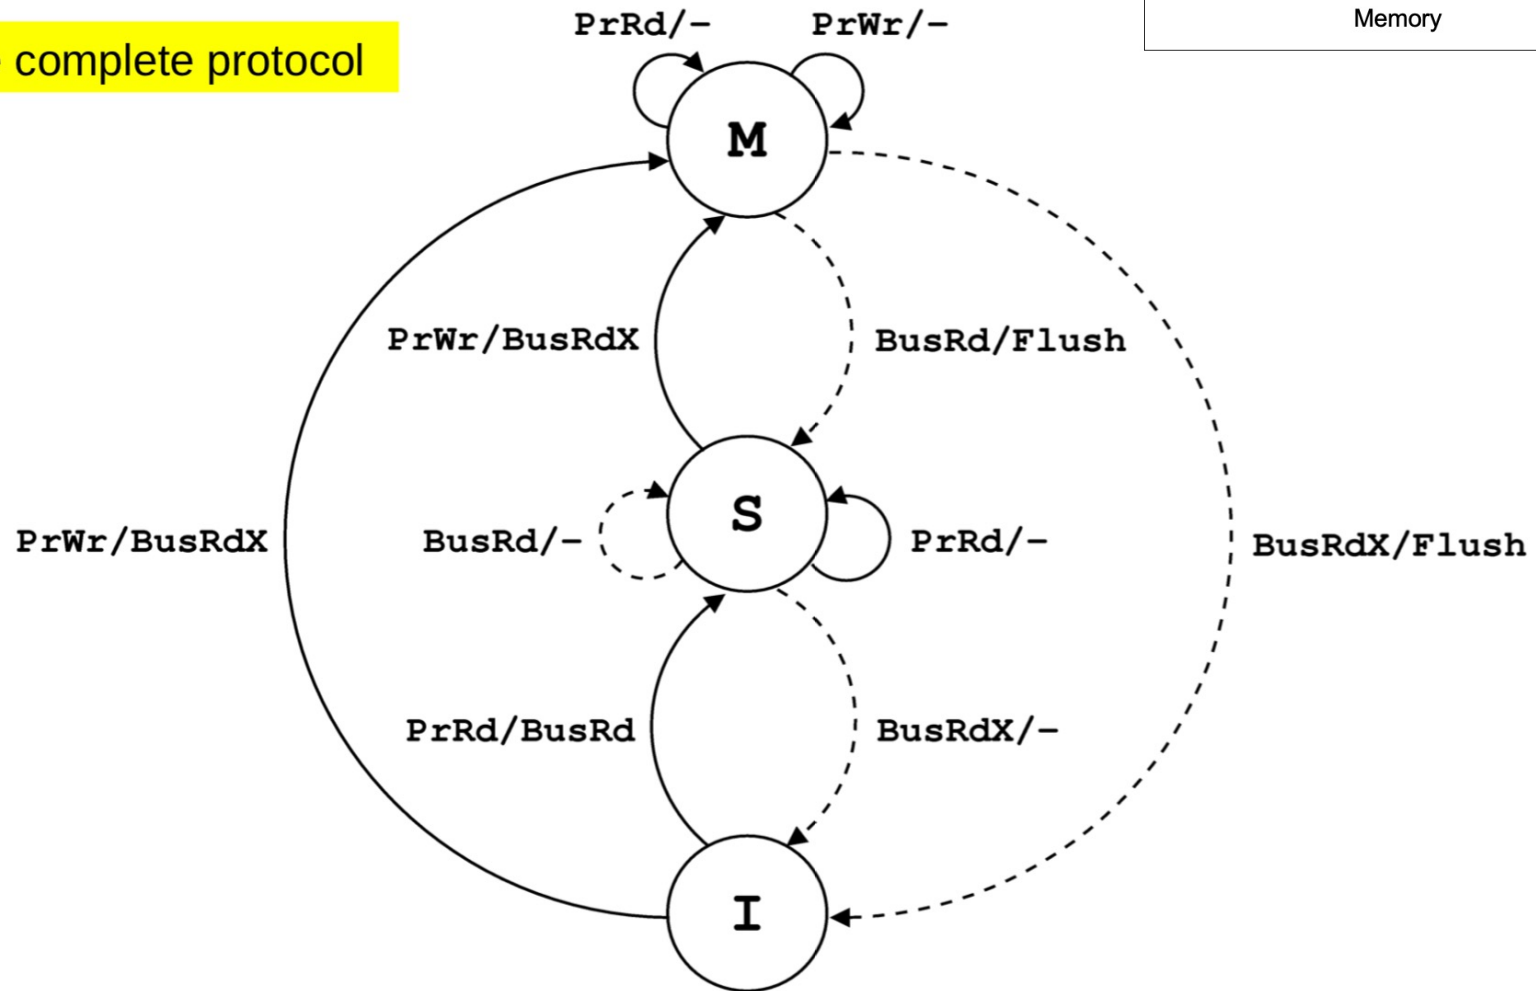


If we are holding the requested block and receive a **BusRdX**, we must transition to the I(nvalid) state

Snooping Cache Protocols (Cont)

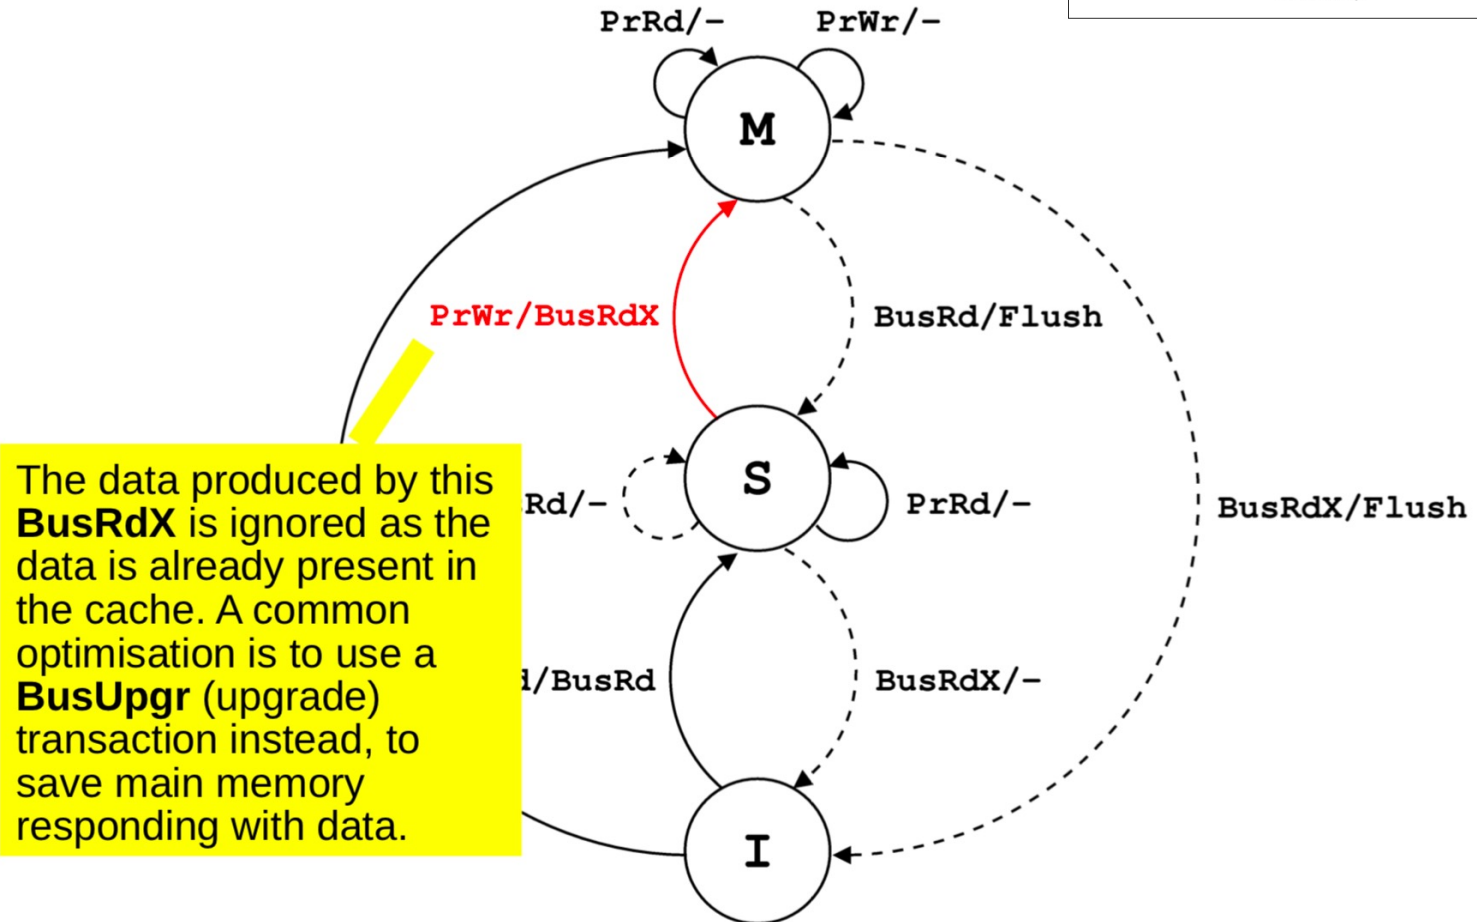

■ MSI (Modified-Shared-Invalid) Protocol

The complete protocol



Snooping Cache Protocols (Cont)

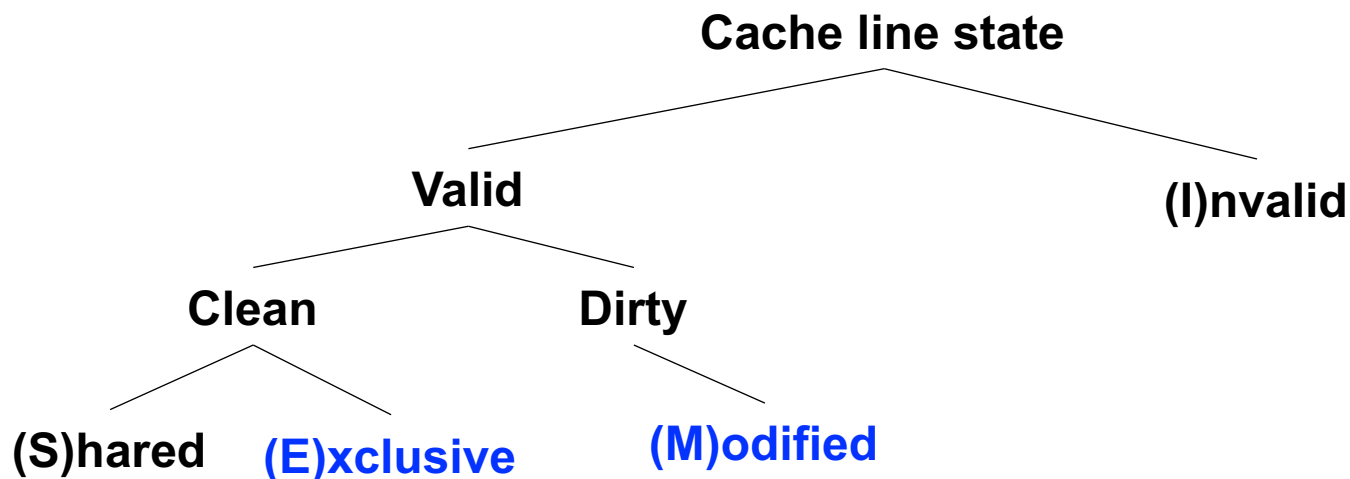
■ MSI (Modified-Shared-Invalid) Protocol



Snooping Cache Protocols (Cont'd)


▪ MESI Protocol

- Let's consider what happens if we read a block and then subsequently modify it
 - This will require two bus transactions in the MSI protocol
 - But if we know that we have the only copy of the block, the transaction (BusUpgr) required to transition from state S to M is really unnecessary
 - › We could safely, and silently, transition from S to M
 - This will be a common event in parallel programs




Snooping Cache Protocols (Cont

■ MESI Protocol (Cont'd)



If we are in state E and need to write, we require no bus transaction to take place



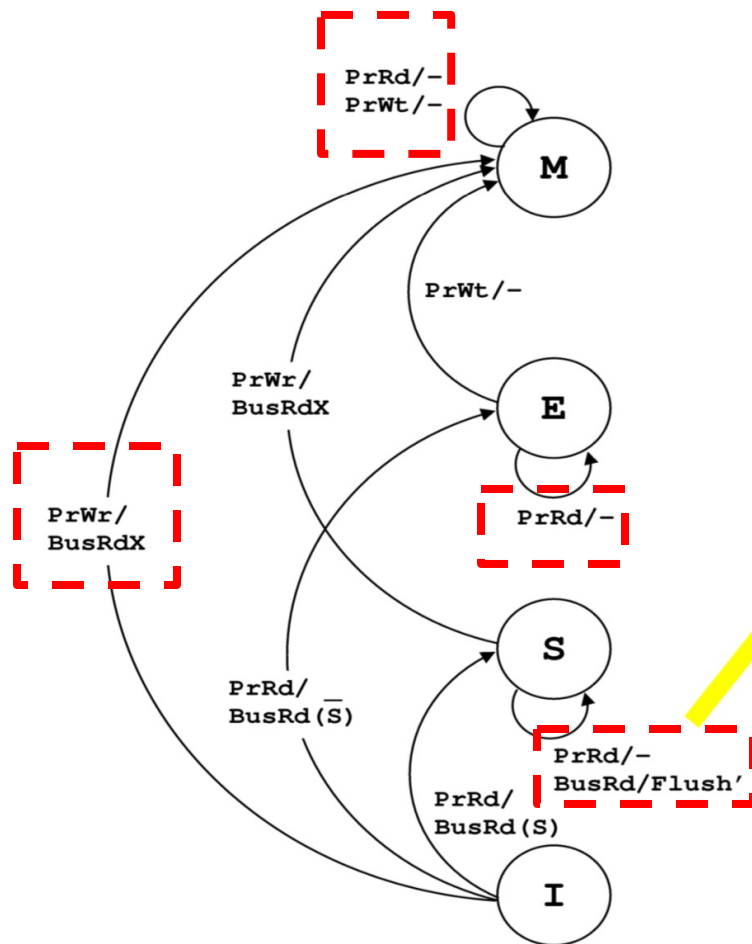

The shared signal (S) is used to determine if any caches currently hold the requested data on a PrRd.

BusRd(S) means the bus read transaction caused the shared signal to be asserted (another cache has a copy of the data).
BusRd(Not S) means no cache has the data.

This signal is used to decide if we transition from state I to S or E

Snooping Cache Protocols (Cont

■ MESI Protocol (Cont'd)



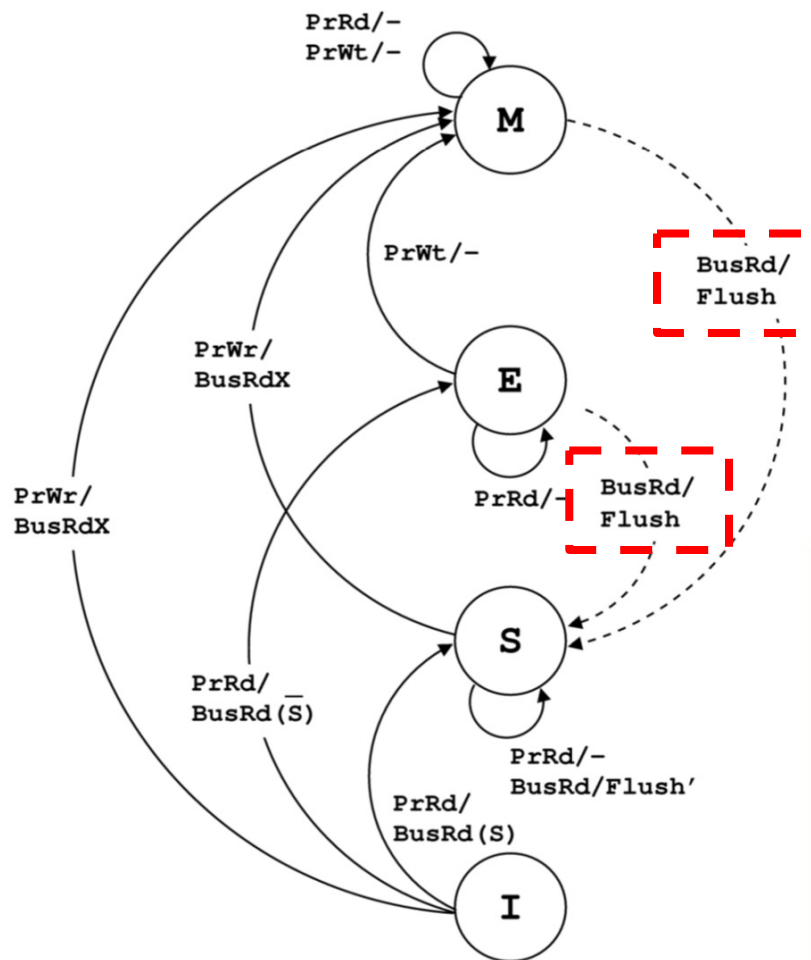

BusRd/Flush'

To enable cache-to-cache sharing we must be sure to select only one cache to provide the required block. Flush' is used to indicate that only **one** cache will flush the data (in order to supply it to the requesting cache). This is easy in a bus-based system of course as only one transaction can take place at a time.

Of course a cache will normally be able to provide the block faster than main memory.

Snooping Cache Protocols (Cont

■ MESI Protocol (Cont'd)

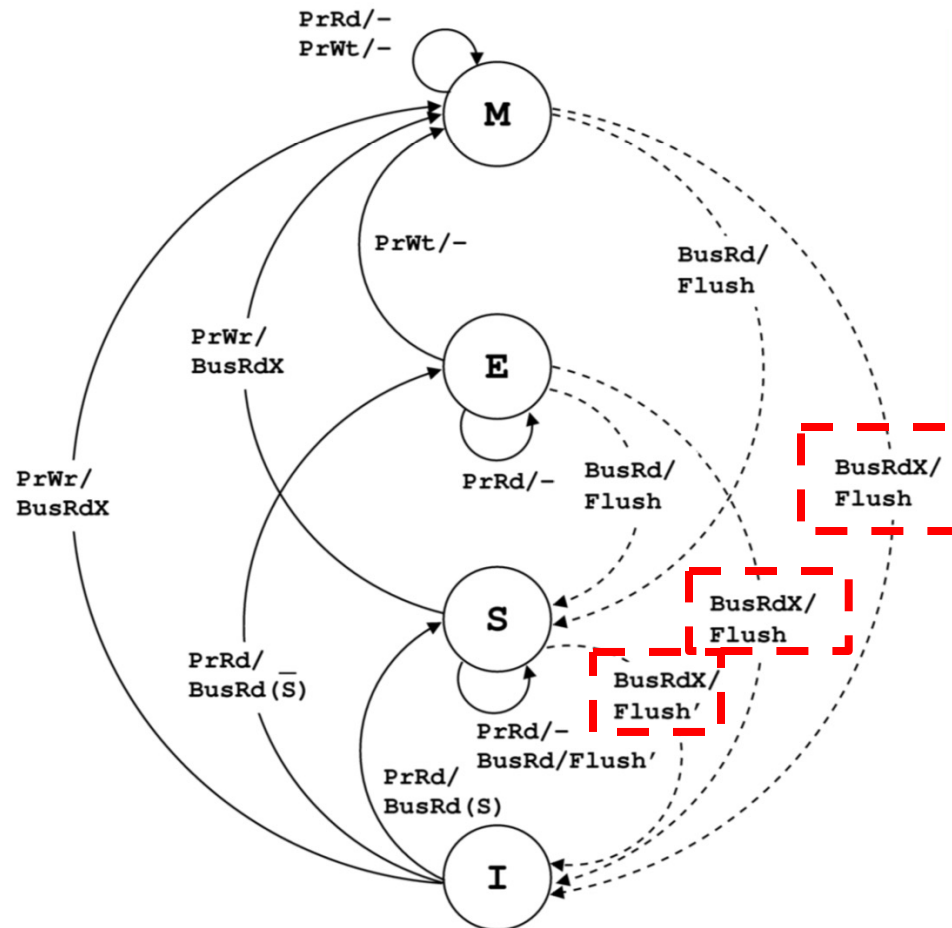



BusRd/Flush

If we are in state M or E and other cache requests the block, we provide the data (we have the only copy) and move to state S (clean & zero or more copies)

Snooping Cache Protocols (Cont

■ MESI Protocol (Cont'd)



BusRdX/Flush

If we are in states M, E or S and receive a BusRdX we simply flush our data and move to state I

Directory-Based Coherence

- Scalability bottleneck: snooping bandwidth
 - Snooping broadcast are a major drain of cache bus bandwidth
 - Practical scalability limit: 8 cores
- Solution: add a **directory**
 - Directory keeps
 - The state of every cache line
 - Which caches have a copy of the block
 - Directory also must be distributed

